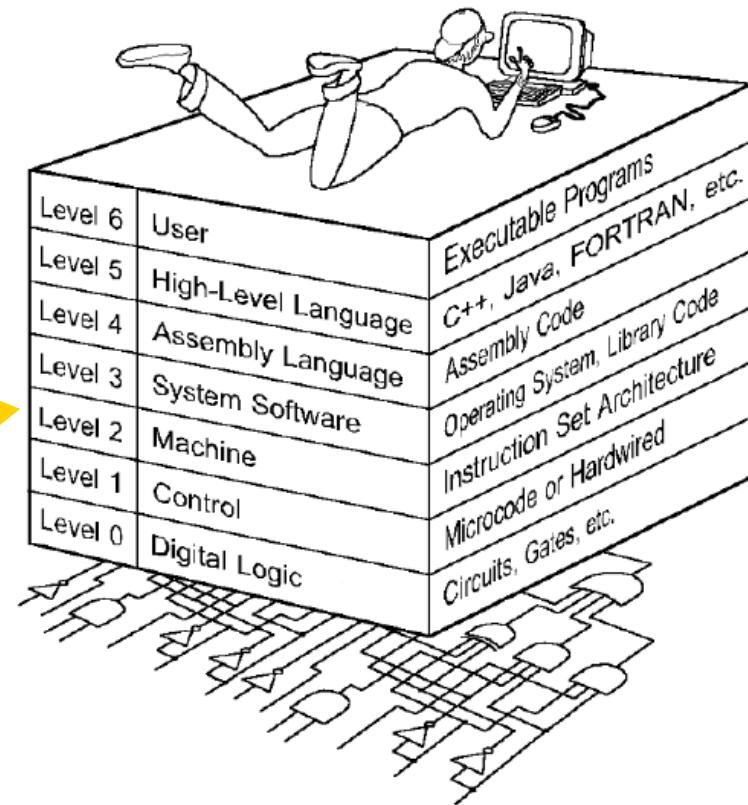
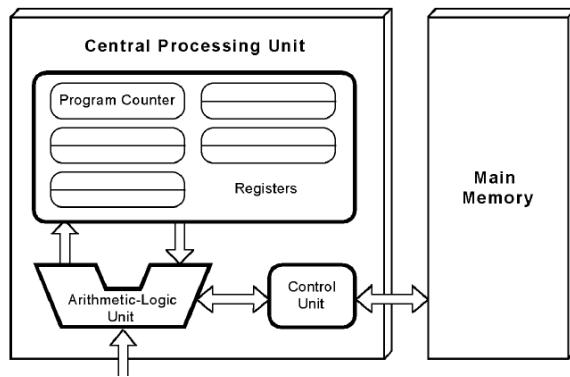


Fetch-Decode-  
Execution  
CYCLE



# Chapter 5 – A Closer Look at Instruction Set Architecture (ISA)

## LO-3:

- Describe and explain the **organization** of the classical **von Neumann computer** and its **major functional units**.
- Describe the functioning of a **single cycle CPU** and its internal operations.

>>> Quiz-3 and Test-3 (Chs-4 & 5)

# Chapter 5 Objectives

---

- Understand the **factors** involved in **ISA design**.
- Gain familiarity with **memory addressing modes**.
- Understand the concepts of instruction-level  
**pipelining** and its affect upon execution  
**performance**.

# Outline

---

- Instruction formats and types
- Operand types
- Addressing (memory access)
- Instruction pipelining
- Real-world examples of ISAs

An Instruction:

OpCode Operand1 Addr. Operand2 Addr. Operand3 Addr.

## 5.2 Instruction Formats

Instruction sets are **differentiated** by the following:

- Number of bits per instruction.
  - Whether short, long, or variable.
- Types of operations.
- Stack-based or register-based.
- Number of explicit operands per instruction.
- Operand location.
  - Number of addressable registers.
  - Memory byte- or word- addressable.
  - Memory Addressing modes.
- Type and size of operands.

Primarily, based on the operator and operand types, size, location

## 5.2 Instruction Formats

---

Instruction set architectures are **measured** according to:

- Main memory **space** occupied by a program.
- Instruction **complexity**.
- Instruction **length** (in bits).
- **Total number of instructions** in the instruction set.

## 5.2 Instruction Formats

---

- Byte ordering, or *endianness*, is another major architectural consideration.
- If we have a two-byte integer, the integer may be stored so that the least significant byte is followed by the most significant byte or vice versa.

e.g. IA x86,  MIPS, ARM\*  
Windows  
\* configurable

In *little endian* machines, the least significant byte is followed by the most significant byte.

e.g. RISC,  IBM  
Motorolla  
Unix, Linux, Mac

*Big endian* machines store the most significant byte first (at the lower address).

## 5.2 Instruction Formats

---

- As an example, suppose we have the hexadecimal number 0x12345678.
- The big endian and small endian arrangements of the bytes are shown below.

Address →	00	01	10	11
natural Big Endian	12	34	56	78
Little Endian	78	56	34	12

## 5.2 Instruction Formats

- A larger example: A computer uses 32-bit integers. The values 0xABCD1234, 0x00FE4321, and 0x10 would be stored sequentially in memory, starting at address 0x200 as below.

32 bit word boundary

Address	Big Endian	Little Endian
0x200	AB	34
0x201	CD	12
0x202	12	CD
0x203	34	AB
0x204	00	21
0x205	FE	43
0x206	43	FE
0x207	21	00
0x208	00	10
0x209	00	00
0x20A	00	00
0x20B	10	00

## 5.2 Instruction Formats

ASM

```
.data  
dw1 DWORD 12345678h  
dw2 DWORD 'AB', '123', 123h  
;dw3 DWORD 'ABCDE' ; error A2084: constant value too Large  
by3 BYTE 'ABCDE', 0FFh, 'A', 0Dh, 0Ah, 0  
w1 WORD 123h, 'AB', 'A'
```

For simplicity, the hexadecimal constants are used as initializer. The memory repre

Memory 1	
Address:	0x00CF51A6
0x00CF51A6	78 56 34 12 42 41 00 00   33 32 31 00 23 01 00 00 xv4.BA..321.#...
0x00CF51B6	41 42 43 44 45 ff 41 0d 0a 00 23 01 42 41 41 00 ABCDEjA...#.BAA.

### □ Big endian:

- Is more natural.
- The sign of the number can be determined by looking at the byte at address offset 0.
- Strings and integers are stored in the same order.

### □ Little endian:

- Makes it easier to place values on non-word boundaries.
- Conversion from a 16-bit integer address to a 32-bit integer address does not require any arithmetic.

### Tradeoffs:

- simplicity, speed, cost, ease of use

### Where are operands?

Memory-memory: 2+ in memory.  
Register-memory: 1+ in a register.  
Load-store: no operand in memory.

## 5.2 Instruction Formats

---

- In a **stack** architecture, instructions and operands are implicitly taken from the stack.
  - A stack cannot be accessed randomly!
- In an **accumulator** architecture, one operand of a binary operation is implicitly in the accumulator.
  - Other operand is in memory, creating lots of bus traffic.
- In a **general purpose register (GPR)** architecture, most used today, registers can be used instead of memory.
  - Faster than accumulator architecture.
  - Efficient implementation for compilers.
  - Results in longer instructions.

>>> Java

>>> ~MARIE

>>> IA x86

## 5.2 Instruction Formats

---

- Let's see how to evaluate an **infix expression** using different instruction formats.
- With a **three-address ISA**, (e.g., **mainframes**), the infix expression, e.g. RISC, ARM

$$Z = X \times Y + W \times U$$

might look like this:

MULT R1,X,Y  
MULT R2,W,U  
ADD Z,R1,R2

<== A commonly used notation in arithmetical and logical formulae and statements.

## 5.2 Instruction Formats

---

- In a two-address ISA, (e.g., Intel, Motorola), the infix expression,

$$Z = X \times Y + W \times U$$

then Atmel (used in Arduino), now Microchip

might look like this:

LOAD	R1,X
MULT	R1,Y
LOAD	R2,W
MULT	R2,U
ADD	R1,R2
STORE	Z,R1

Note: Two-address ISAs usually require one operand to be a register.

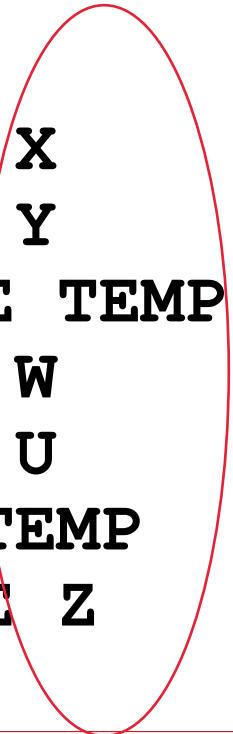
## 5.2 Instruction Formats

---

- In a one-address ISA, like MARIE, the infix expression,

$$Z = X \times Y + W \times U$$

looks like this:



```
LOAD X
MULT Y
STORE TEMP
LOAD W
MULT U
ADD TEMP
STORE Z
```

---

Memory references (X, Y, ...) come from the Symbol Table<sup>18</sup>

## 5.2 Instruction Formats

---

- In a stack ISA, the **postfix** expression,

**Z = X Y × W U × +**

might look like this:

**PUSH X**  
**PUSH Y**  
**MULT**  
**PUSH W**  
**PUSH U**  
**MULT**  
**ADD**  
**POP Z**

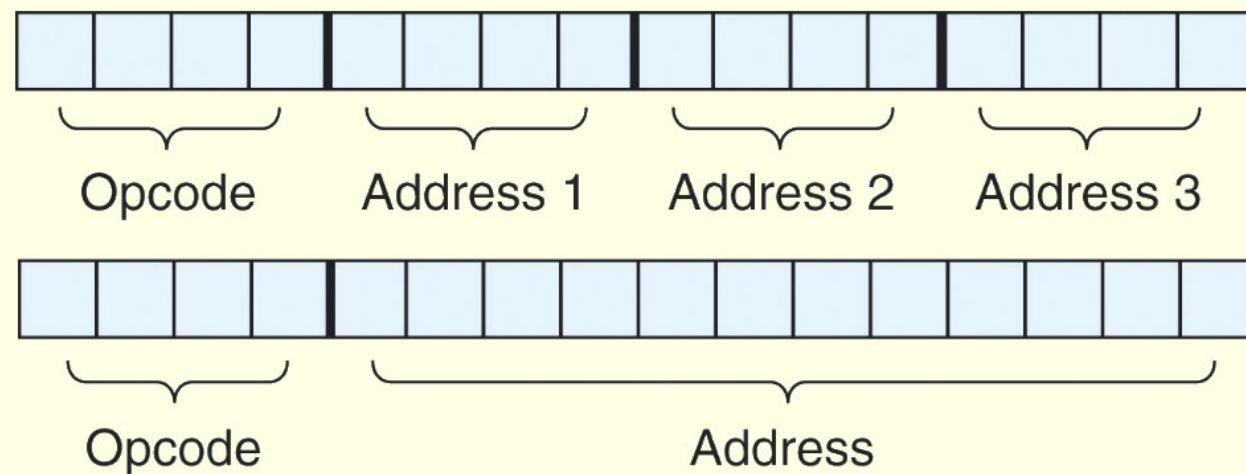
In sum, the ISA with lesser\* number of operands results in bigger/longer program that will have more number of instructions to be executed.

\* implicit registers may be involved.

Expanding opcodes are used to save program space wastage when fixed-length instructions are used.

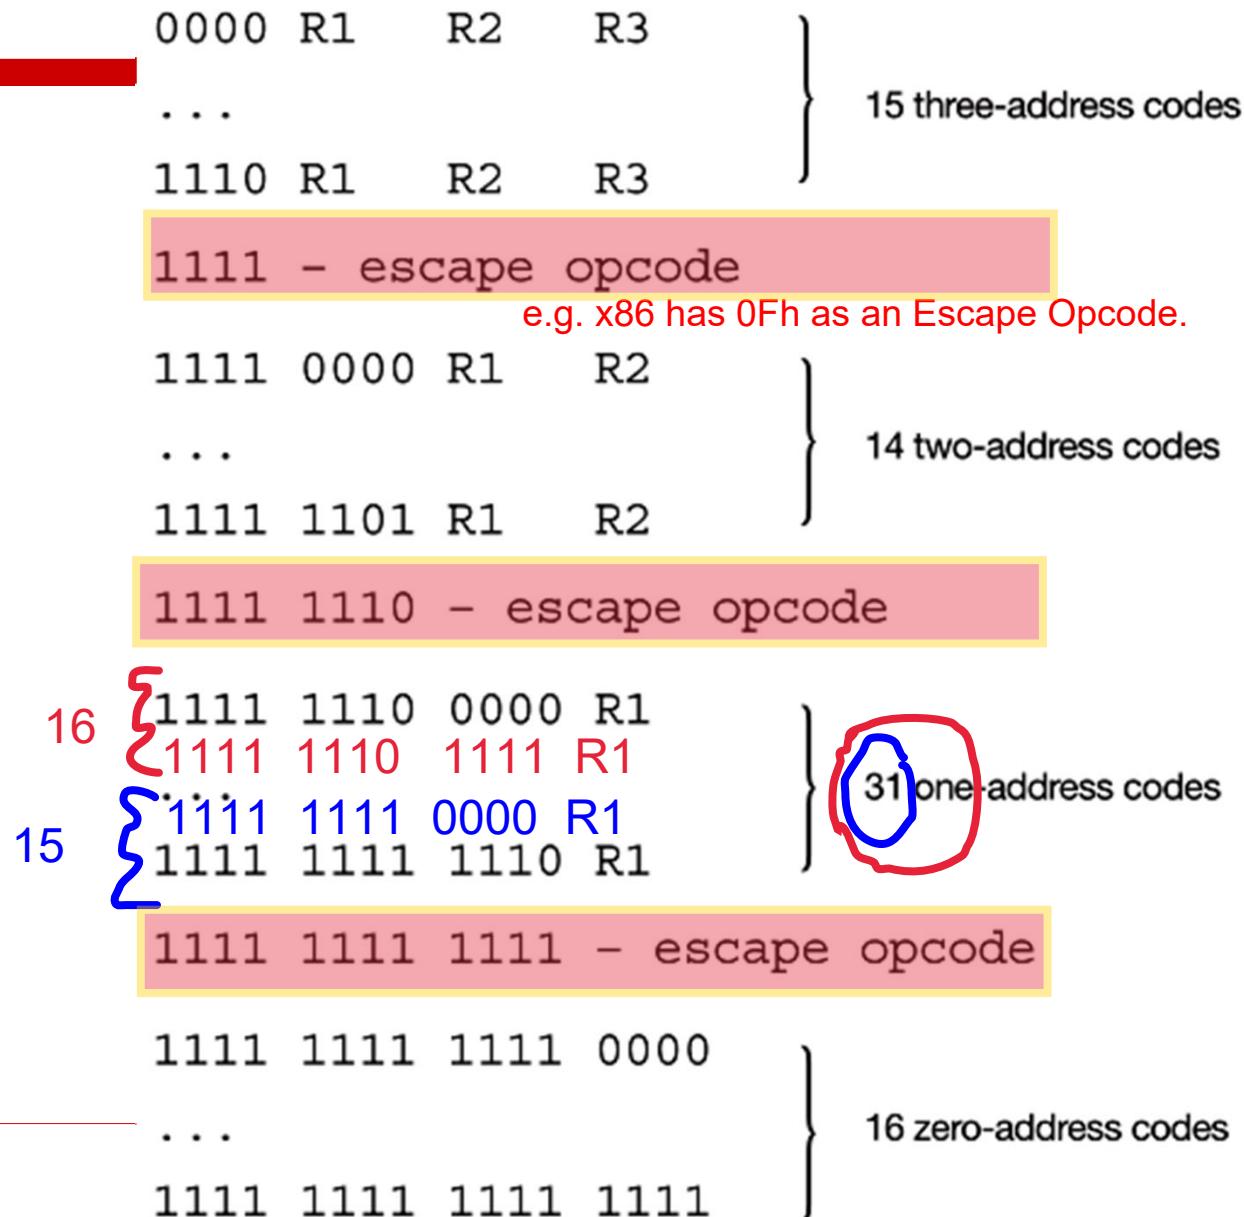
## 5.2 Instruction Formats

- A system has 16 registers and 4K of memory.
- We need 4 bits to access one of the registers. We also need 12 bits for a memory address.
- If the system is to have 16-bit instructions, we have two choices for our instructions:



## 5.2 Instruction Formats

- ☐ If we allow the length of the opcode to **vary**, we could create a very **rich** instruction set:



## 5.3 Instruction types

Instructions fall into several broad categories that you should be familiar with:

- Data movement.
- Arithmetic.
- Boolean.
- Bit manipulation.
- I/O.
- Control transfer.
- Special purpose.

Can you think of some examples of each of these?

Instruction
Load X
Store X
Add X
Subt X
Input
Output
Halt
Skipcond
Jump X

Instruction
JnS X
Clear
AddI X

JumpI X

LoadI X

StoreI X

## 5.4 Addressing

- where an operand is located?
- operand's actual location is its effective address.

- *Immediate addressing* is where the data is part of #  
the instruction.
- *Direct addressing* is where the address of the data X  
is given in the instruction.
- *Register addressing* is where the data is located in R1  
a register.
- *Indirect addressing* gives the address of the \*x  
address of the data in the instruction. e.g. pointers, IVT
- *Register indirect addressing* uses a register to store \*r1  
the address of the data. e.g. MAR in MARIE

## 5.4 Addressing

---

- *Indexed addressing* uses a **register** (implicitly or explicitly) as an **offset**, which is **added** to the **address in the operand to determine the effective address** of the data. e.g. arrays
- *Base addressing* is similar except that a **base register** is used instead of an **index register**.
- The difference between these two is that an **index register** holds an **offset addition** to the address given in the instruction, a **base register** holds a base address where the address field represents a **displacement from this base**. e.g. JVM, segmentation

## 5.4 Addressing

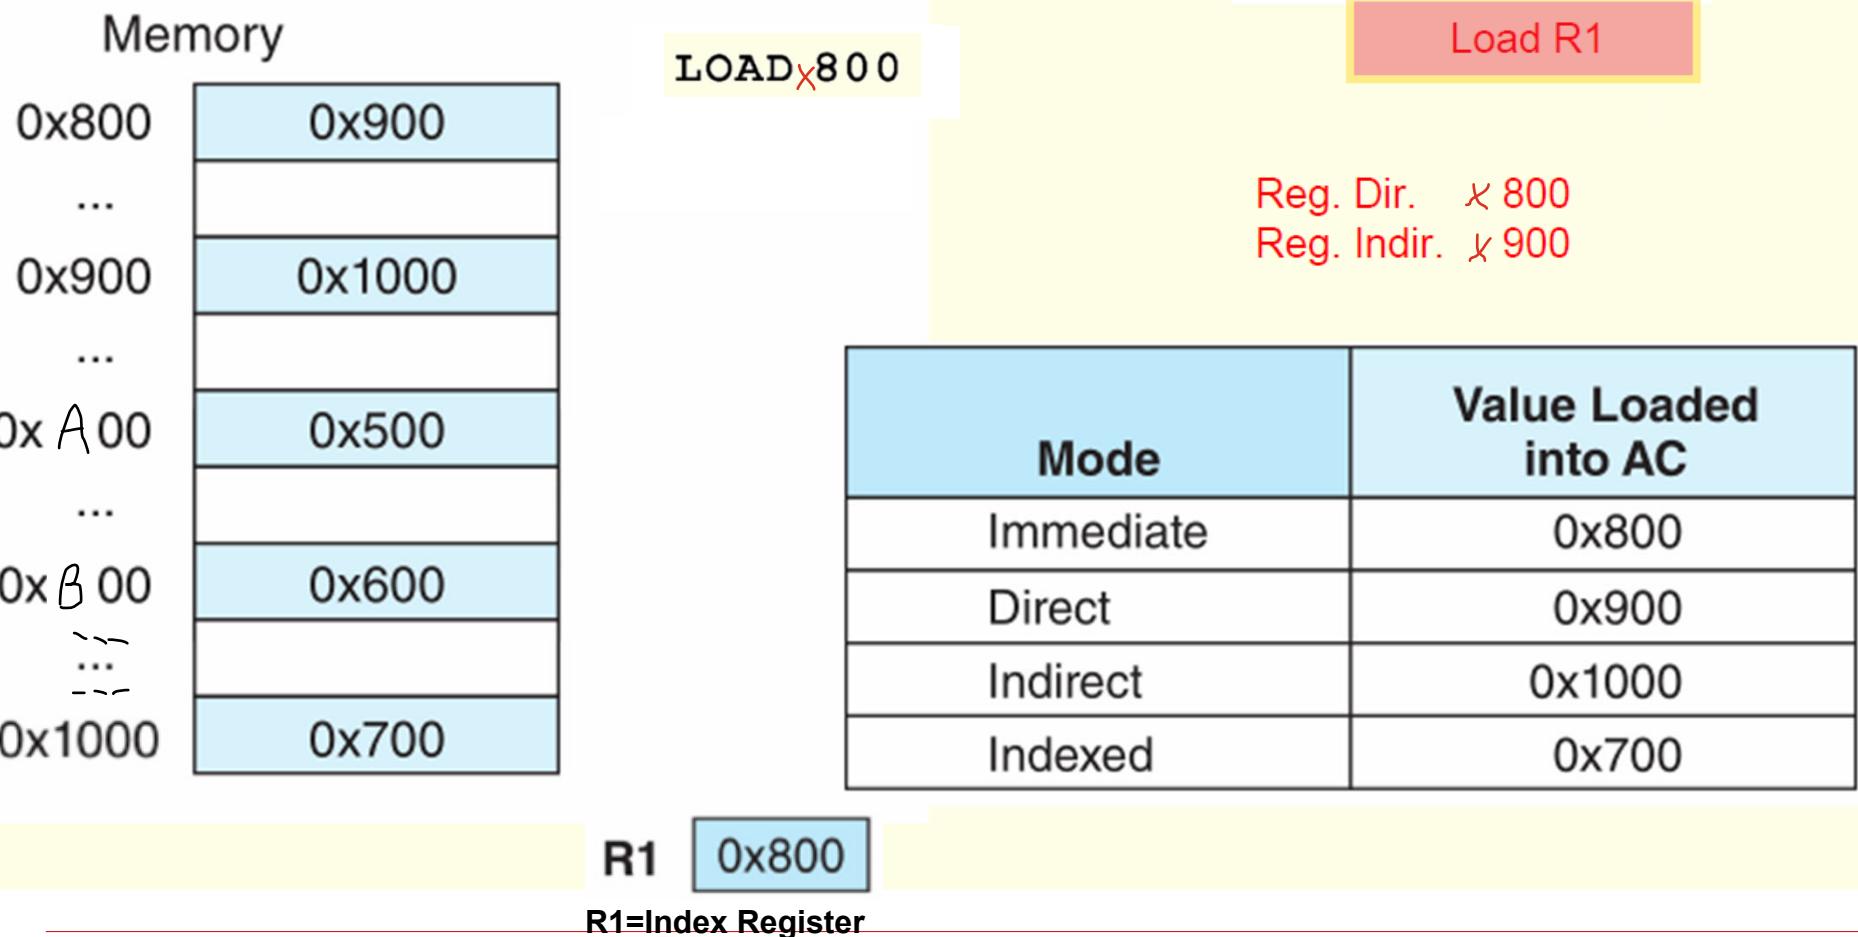
---

- In *stack addressing* the operand is assumed to be on top of the stack.
- There are many variations to these addressing modes including:
  - Indirect indexed.
  - Base/offset.
  - Self-relative
  - Auto increment - decrement.
- We won't cover these in detail.

---

Let's look at an example of the principal addressing modes.

# 5.4 Addressing

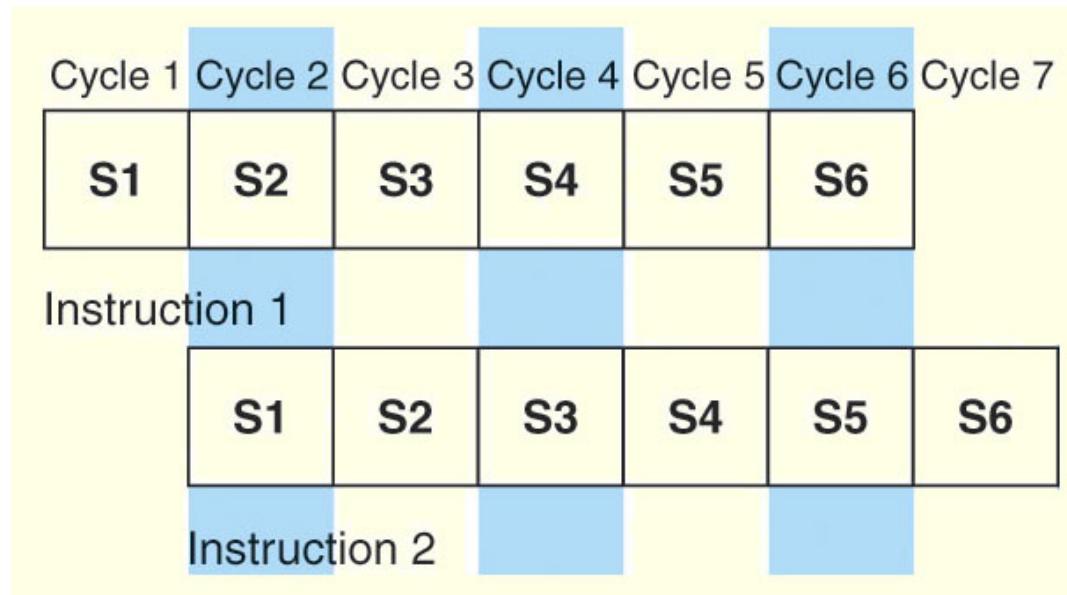


## 5.5 Instruction Pipelining

==> provides instruction level parallelism (ILP)

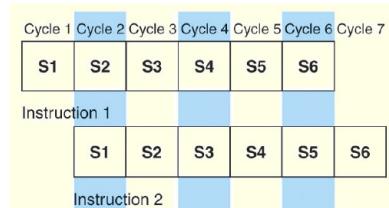
- For every clock cycle, one small step is carried out, and the stages are **overlapped**.

6-stage pipeline



- S1. Fetch instruction.
- S2. Decode opcode.
- S3. Calculate effective address of operands.

- S4. Fetch operands.
- S5. Execute.
- S6. Store result.



## 5.5 Instruction Pipelining

- The **theoretical** speedup offered by a pipeline can be determined as follows:

Let  $t_p$ , be the **time per stage**. **Each instruction represents a task,  $T$ , in the pipeline.**

The first task (instruction) **requires  $k \times t_p$**  time to complete in a  **$k$ -stage** pipeline. The remaining ( $n - 1$ ) tasks emerge from the pipeline one per cycle. So the total time to complete the remaining tasks is  $(n - 1)t_p$ .

Thus, to complete  $n$  tasks using a  $k$ -stage pipeline requires:

$$(k \times t_p) + (n - 1)t_p = (k + n - 1)t_p.$$

From e.g. above (copied to the left):

- $t_p = 1$  clock cycle per stage
- $k=6$  stage-pipeline per instruction  $T$
- $k \times t_p = 6$  cycles per instruction
- remaining  $(n-1)$   $T$  take  $t_p=1$  clock cycle per task !

$\Rightarrow 7$  clock cycles for 2 ins./Tasks

## 5.5 Instruction Pipelining

- If we take the time required to complete  $n$  tasks without a pipeline and divide it by the time it takes to complete  $n$  tasks using a pipeline, we find:

$$\text{Speedup } S = \frac{n t_n}{(k + n - 1) t_p}$$

here, let's assume:  
 $t_n$ =time per task  
=  $k t_p$

- If we take the limit as  $n$  approaches infinity,  $(k + n - 1)$  approaches  $n$ , which results in a theoretical speedup of:

$$\text{Speedup } S = \frac{k t_p}{t_p} = k$$

$$= \frac{k n t_p}{(k+n-1) t_p}$$

Assumptions in the equation:  
1) architecture supports fetching instructions and data in parallel.  
2) pipeline can be kept filled at all times. Pipeline hazards arise that cause pipeline conflicts and stalls.

## 5.5 Instruction Pipelining

---

- An instruction pipeline may **stall**, or be **flushed** for any of the following reasons:
  - Resource conflicts.
  - Data dependencies.
  - Conditional branching.
- **Measures** can be taken at the **software level** as well as at the **hardware level** to reduce the effects of these hazards, but they **cannot** be totally **eliminated**.
  - Pipelining, performance, capability:
  - multi-tasking:
    - concurrency (multi-threading)
    - simultaneously (xCores, multi-processing)

## 5.6 Real-World Examples of ISAs: Intel

---

- Pentium-I introduced two 5-stage **pipelines**. Pentium IV had a 24-stage, Itanium (IA-64) had a 10-stage pipeline.
- The original **8086 provided 17 address modes**, Pentium supported them for backward compatibility.
- The Itanium, having a RISC core, supports only one: register indirect addressing, with optional post increment.

## 5.6 Real-World Examples of ISAs: **MIPS**

---

- ***Microprocessor Without Interlocked Pipeline Stages.***
- Little endian and word-addressable with three-address, fixed-length instructions.
- R2000/R3000 5-stage, R4000/R4400 8-stage pipelines
  - R10000 instructions: 5-stage for integer, 7-stage for FP, 6-stage for LOAD/STORE.
  - only load/store can access memory.
  - uses only base addressing mode.

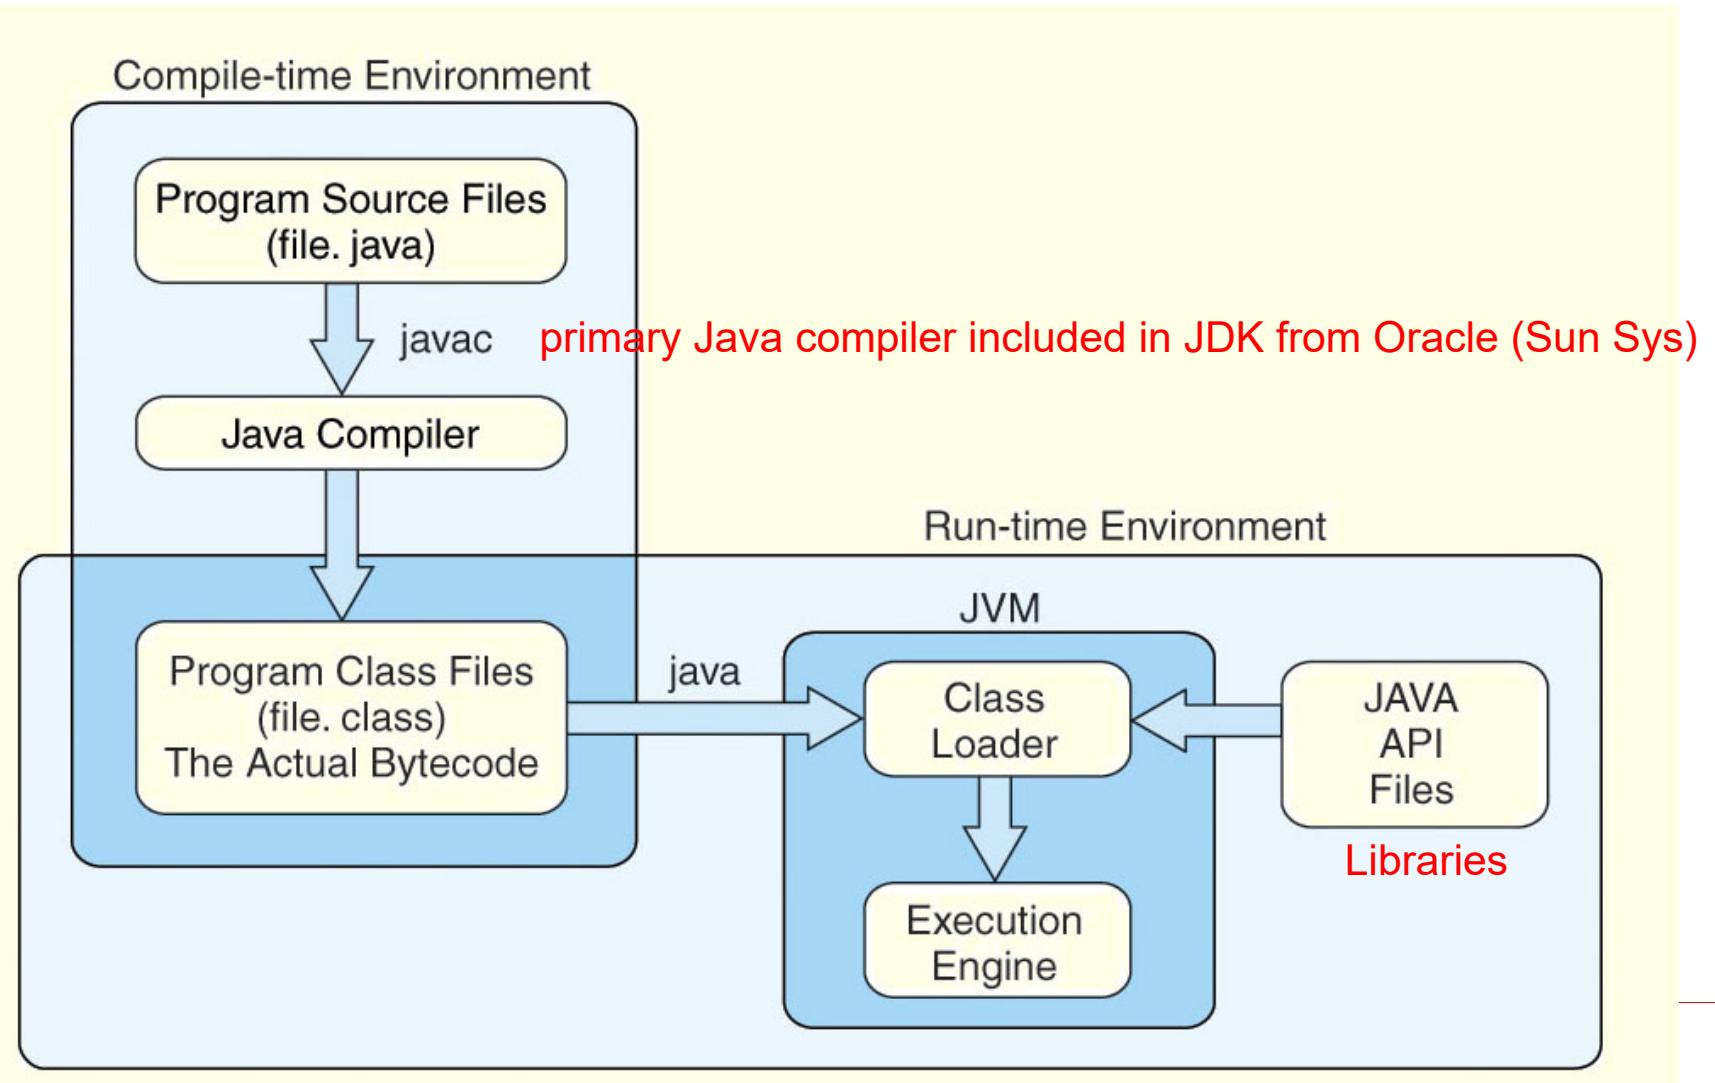
\* MIPS is a family of RISC ISA developed by MIPS Computer Systems (now MIPS Technologies) from Stanford Uni's work. Interlocks were stalling the pipeline when hazard arises and defeats the purpose of pipelining.

## 5.6 Real-World Examples of ISAs: Java4Any!

---

- interpreted language that **runs in a software machine** called the *Java Virtual Machine (JVM)*.
- A JVM is **written in a native language** for a wide array of processors, including MIPS and Intel.
- **has a stack-based ISA** all of its own, called **bytecode**.
- zero address instructions. The JVM has four registers that provide access to five regions of main memory.
- All references to memory are offsets from these registers. Java uses no pointers or absolute memory

## 5.6 Real-World Examples of ISAs: Java



## 5.6 Real-World Examples of ISAs: **ARM**

---

- Most widely used 32-bit instruction architecture:
  - 95%+ smartphones, 80%+ cameras, 40%+ TVs
- Founded in 1990, by Apple and others, ARM (Adv. RISC Machine) is a British firm, ARM Holdings, does not manufacture, but licenses to manufacture.
- Fixed-length, three-operand instructions, load/store arch., simple addressing modes, 3-stage pipeline.
- ARM8+ implementations has 13-stage integer pipeline
- 37 registers, simultaneously load/store any of 16 GPR
- Most instructions execute in a single cycle, if no pipeline hazards or memory accesses.

# ISA's Capability: Summary

---

- Instruction:
  - length, numbers,
    - format, operator/opcode expanding, operand
    - number, location, type, size
  - execution/CPU: stack/Accumulator/GPR, #addr.
  - memory addressing
    - immi, dir, indir, reg, regind, index, base, stack
  - instruction complexity: k-stage pipelining, ILP speedup
- program: length, density
- Adv Arch.
  - Intel: addr modes, pipelining, --> MIPS
  - JVM: stack 0-addr/operand, xPlatform, Referenced
  - ARM: load/store 16 GPRs to 37, fixed 3 ins, pipelined.