



Embedded Systems: Introduction

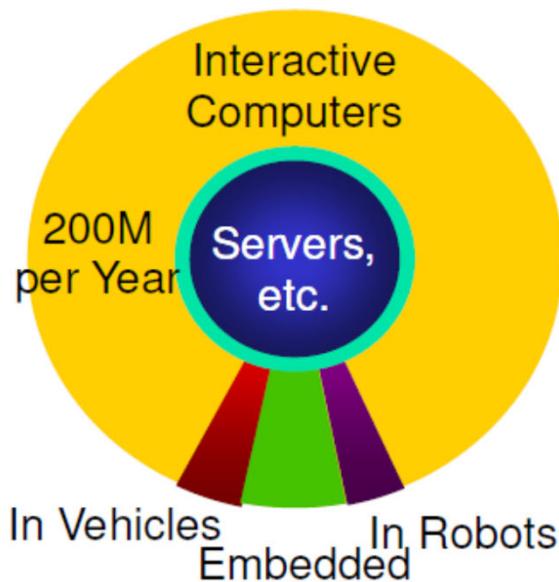
Edited: Waqas Majeed
© Craig Chin



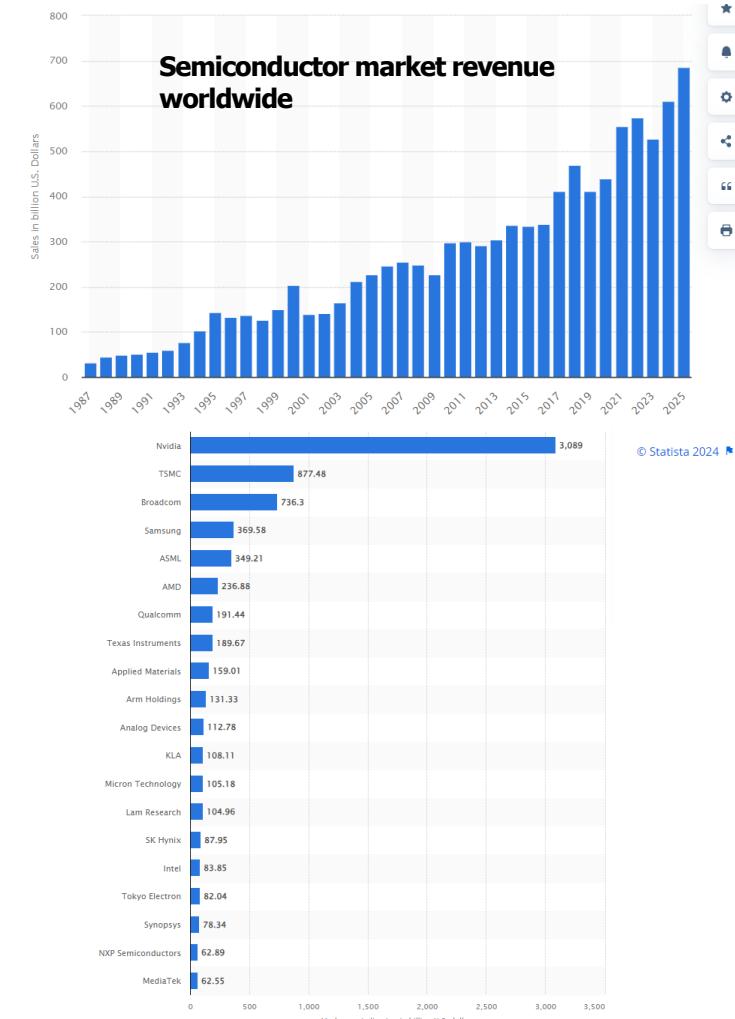
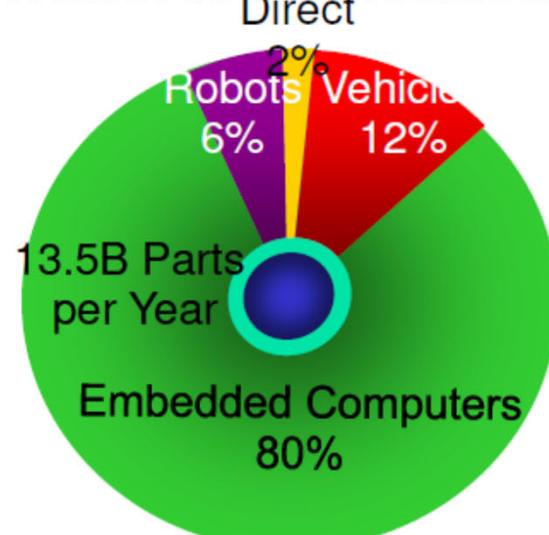
Why Embedded Systems?

Estimated 98% of 20 Billion CPUs produced in 2015 used for embedded apps

Where Has CS Focused?



Where Are the Processors?



Look for the CPUs...the Opportunities Will Follow!

Source: DARPA/Intel (Tennenhouse)



Embedded Systems





Computers Comparison

Embedded Systems

- Special purpose HW and embedded OS.
- Few applications that are known at design-time.
- Not programmable by end user.
- Fixed run-time requirements (additional computing power not useful).
- e.g. MCU, PLC, ...
- Package
- Criteria:
 - cost
 - power consumption
 - Predictability (speed)
 - Performance (Constrained)

General Purpose Computers

- Generic HW and general purpose OS
- Broad class of applications.
- Programmable by end user.
- Faster is better.
- Criteria:
 - cost
 - average speed
- Examples: PCs, Workstations, Servers
- Modular



Embedded Systems Comparison

ASIC	uC	FPGA
Speed: High		
Cost (prototyping): High		
Cost (mass production): Low		
Power Consumption: Low		
Customizability: Low		
Usability: Hard		

ASIC

- Large production volumes

Microcontroller (MCU)

- CPU with RAM, ROM, I/O ports and other peripheral devices integrated as one, to perform specific tasks.
- Assembly, GP Programming Languages

Field Programmable Gate Array (FPGA)

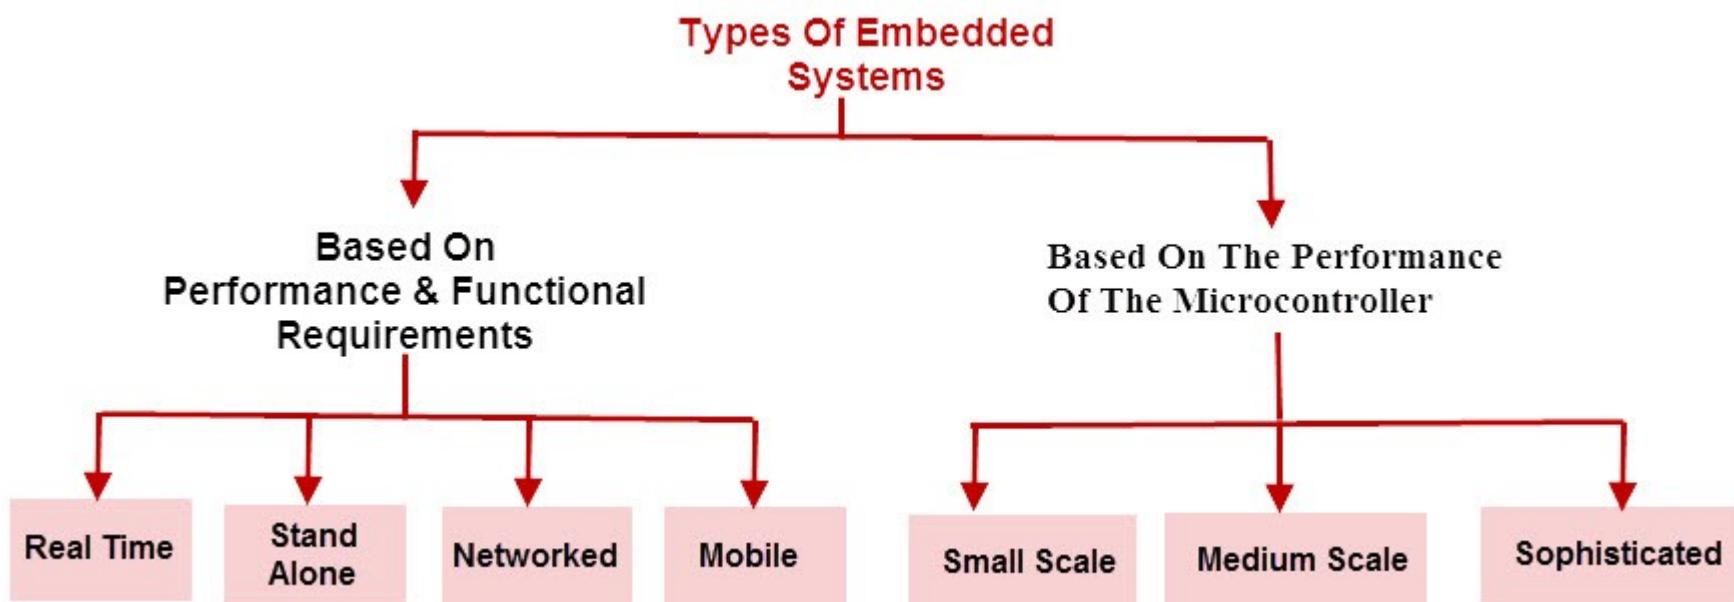
- GP IC for embedded applications
- Co-designing SW and HW
- An integrated circuit that contains millions of logic gates that can be electrically re-configured, through programmable interconnects to perform the required task.
- HDL (VHDL, Verilog) Programming
- Small scale use, prototyping

SoC

- CPU+FPU+iGPU+Cache/s all in one
- e.g. Raspberry Pi SBC:
 - DRAM+SDC Flash Storage+xDisplay+GPIO+USBs
 - Comm (WiFi, BT, Ethernet)



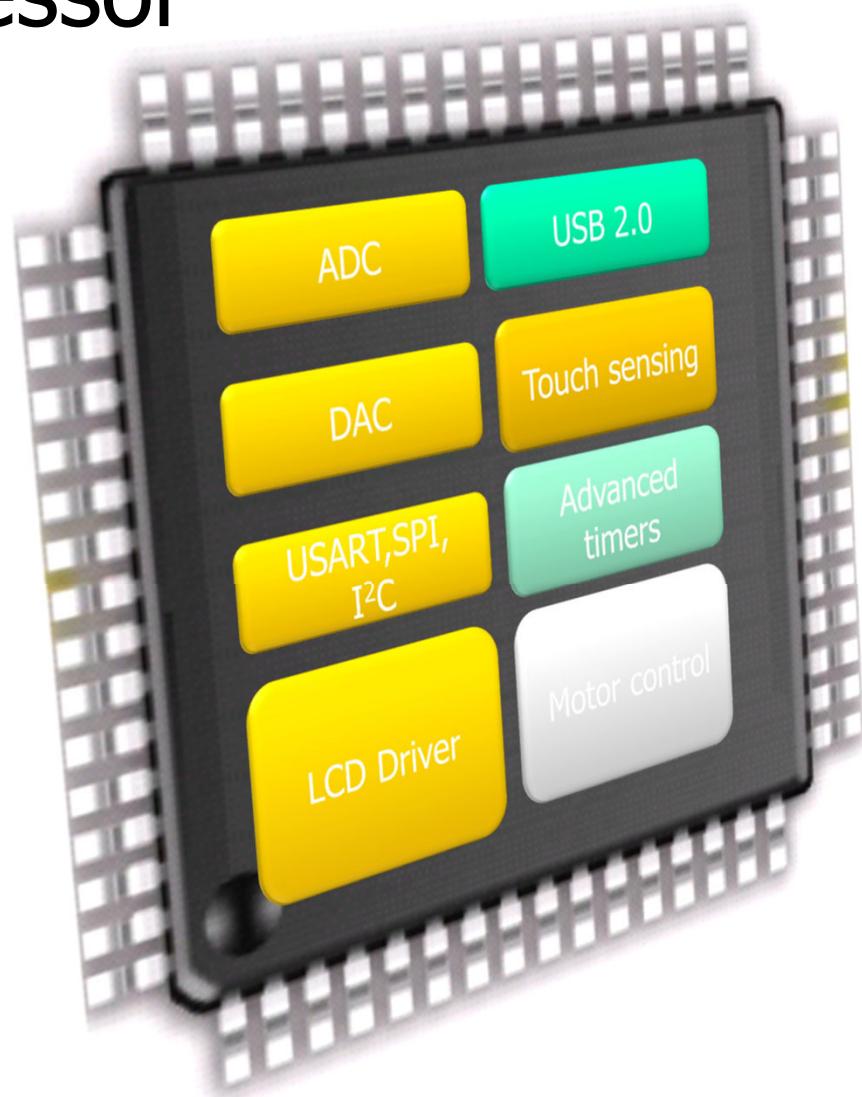
Classification of Embedded Systems





Why ARM processor

- As of 2005, **98%** of the more than one billion mobile phones sold each year used ARM processors
- As of 2009, ARM processors accounted for approximately **90%** of all embedded 32-bit RISC processors
- In 2010 alone, **6.1 billion** ARM-based processors, representing **95%** of smartphones, **35%** of digital televisions and set-top boxes and **10%** of mobile computers
- As of 2016, over 80 billion ARM processors have been produced





Laboratory Main Component

- We will be working with a state-of-the-art ARM processor:
- STM32L476VGT6 microcontroller
- On-board USB interface
- Mbed-enabled
- LCD Display 4 x 24 seg
- 7 LEDs
- Pushbutton + Joystick
- Stereo Audio DAC
- 9 axis Inertial Motion Sensor (IMU)
- 128Mbit flash memory



<http://www.st.com/en/evaluation-tools/32l476gdiscovery.html>

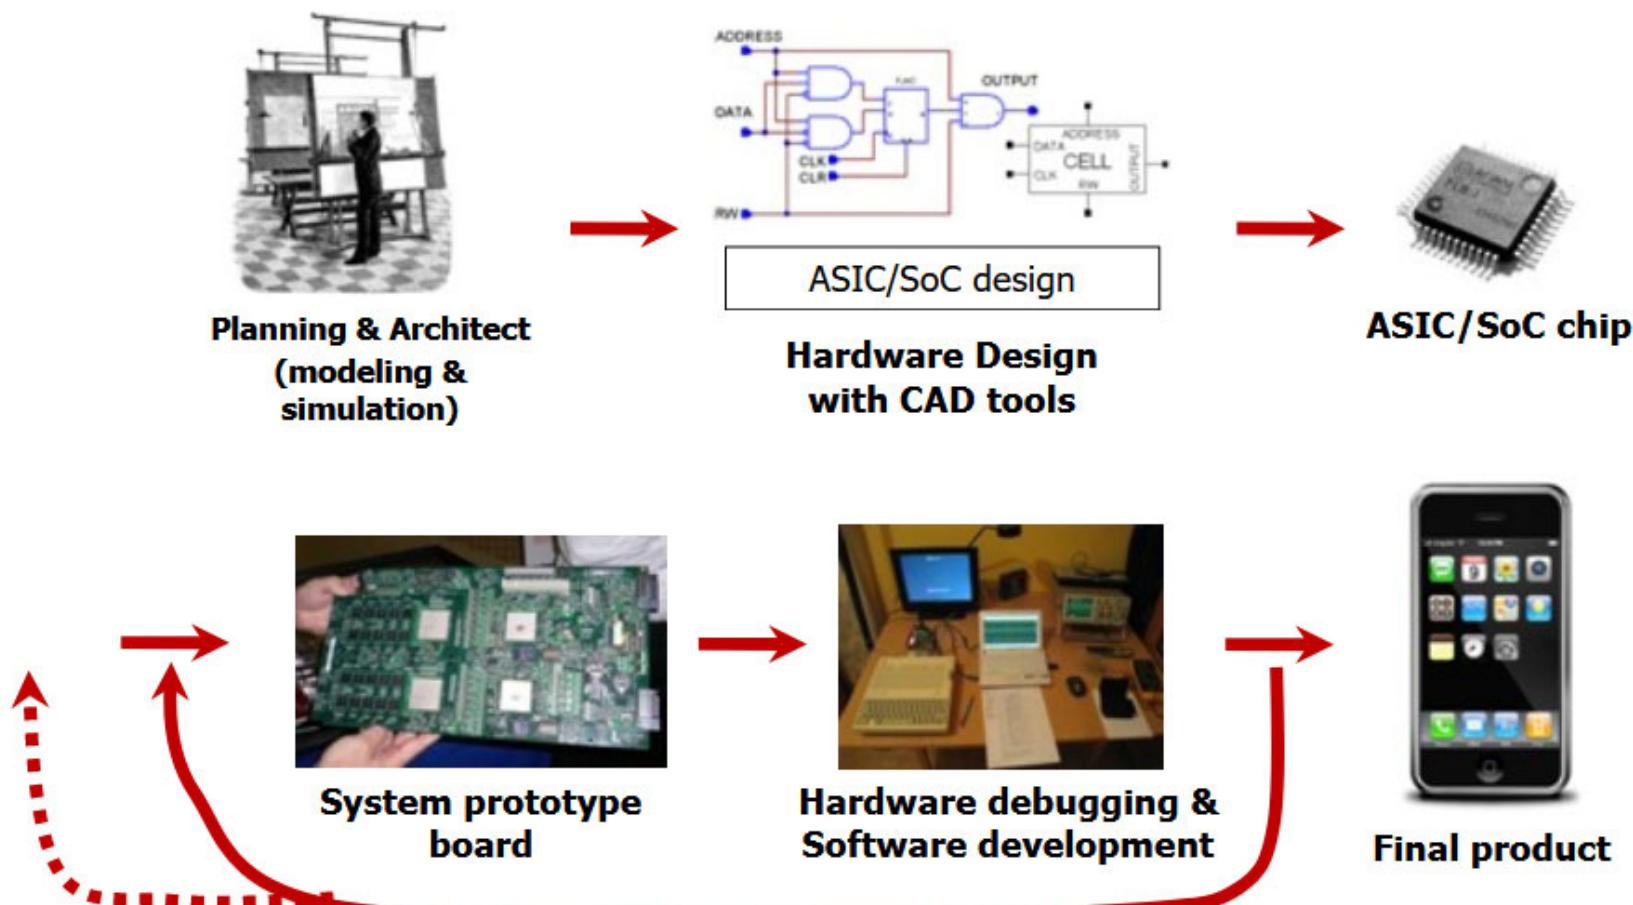


Why do we learn Assembly?

- Hardware/processor specific code
 - Processor booting code
 - Device drivers
 - A test-and-set atomic assembly instruction can be used to implement locks, mutex and semaphores.
- Cost-sensitive applications
 - Embedded devices, where the size of code is limited, washing machine controller, automobile controllers
- **The best applications are written by those who've mastered assembly language or fully understand the low-level implementation of the high-level language statements they're choosing.**



Embedded System Design Flow





Chapter 1

Introduction to Computer Systems and Computer Programs

Edited: Waqas Majeed
©Yifeng Zhu



Structure and Function

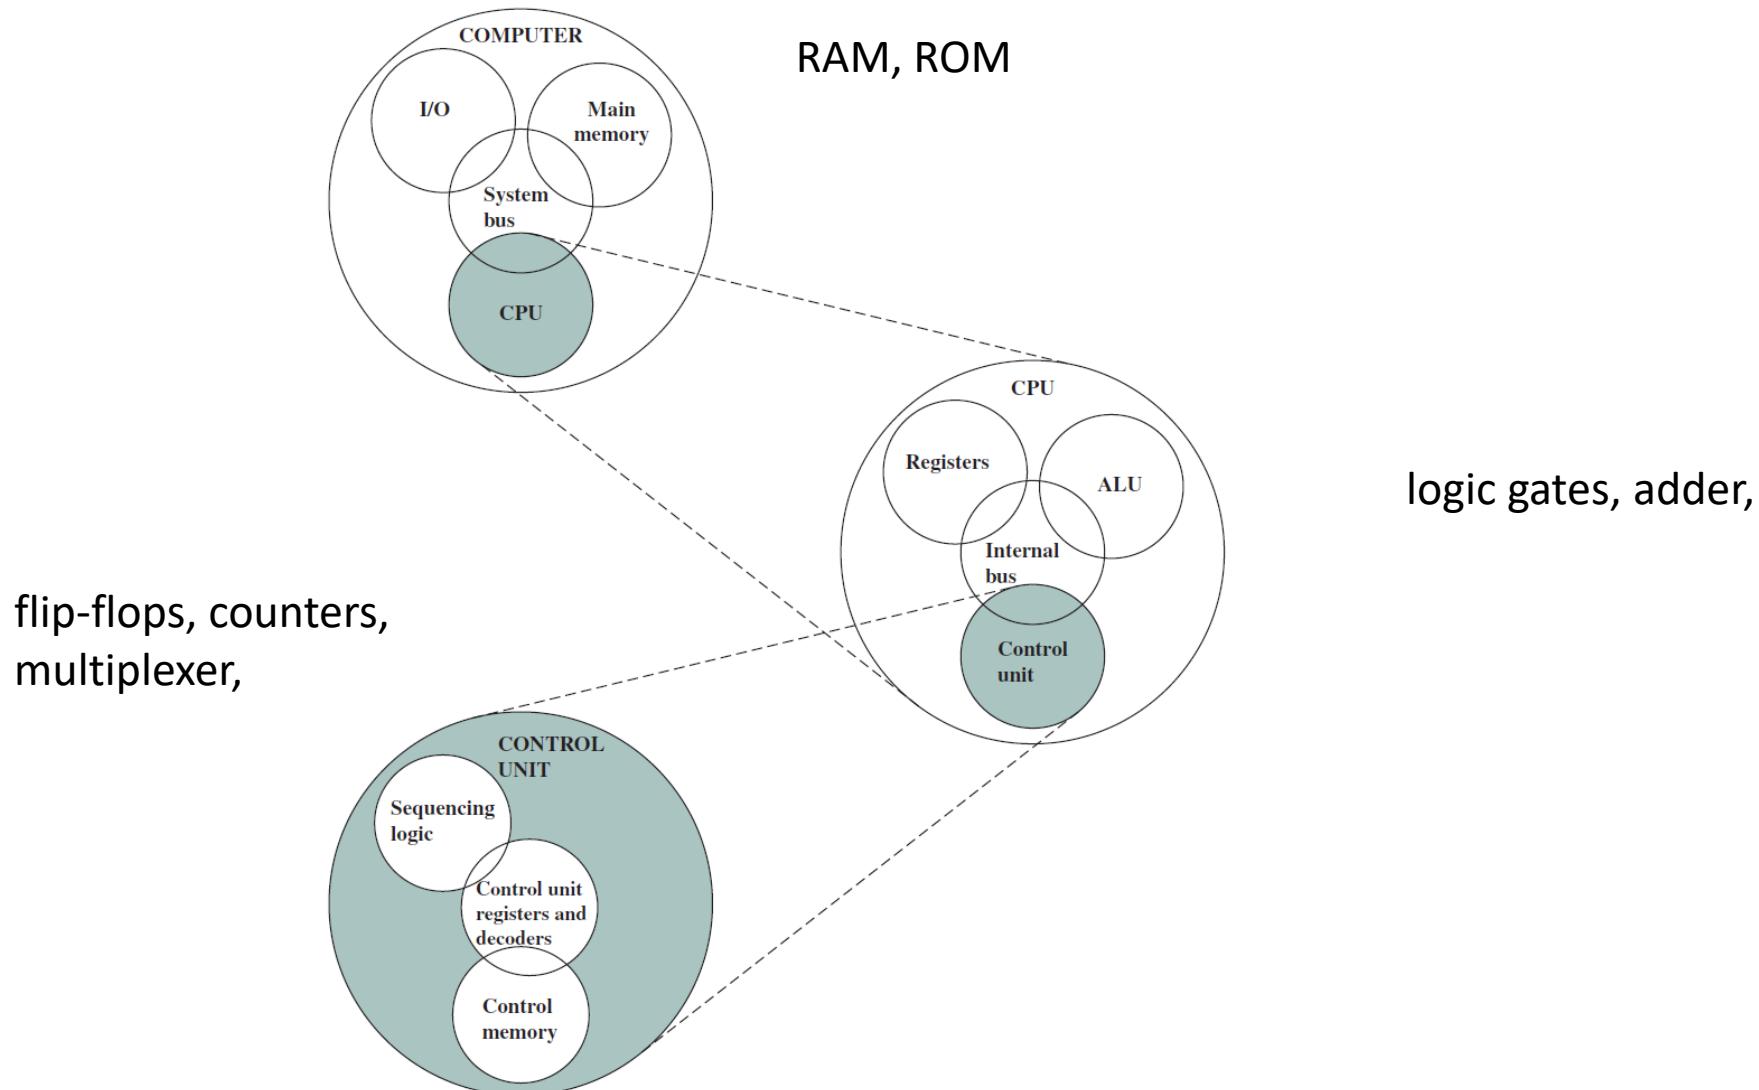
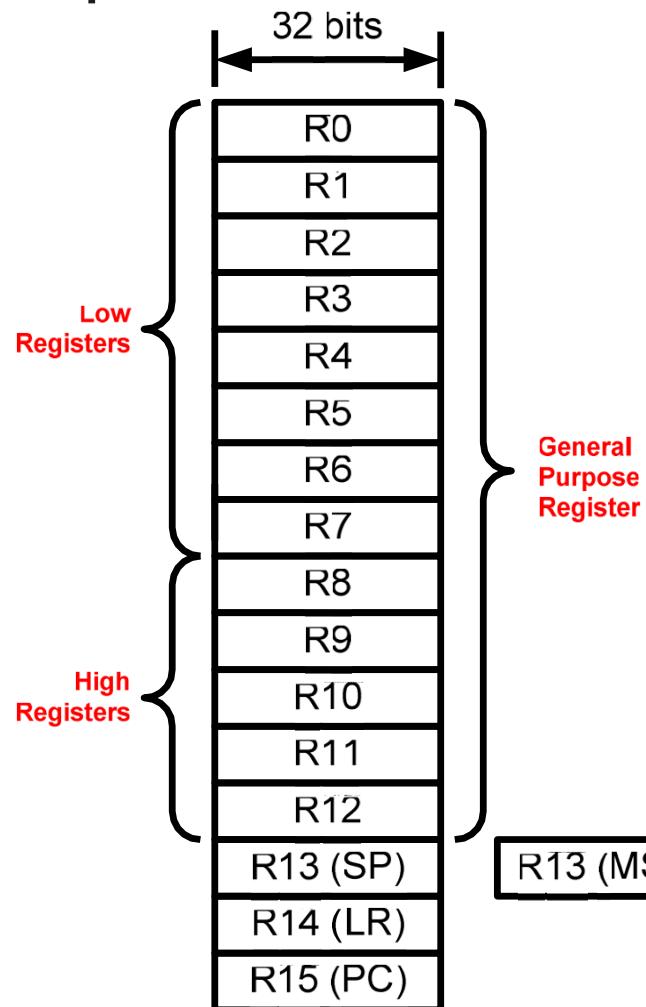


Figure 1.1 The Computer: Top-Level Structure



ARM Processor Registers



- Fastest way to read and write
- Registers are within the processor chip
- A register stores 32-bit value
- STM32L has
 - **R0-R12**: 13 general-purpose registers
 - **R13**: Stack pointer (Shadow of MSP or PSP)
 - **R14**: Link register (LR)
 - **R15**: Program counter (PC)
 - Special registers (xPSR, BASEPRI, PRIMASK, etc)

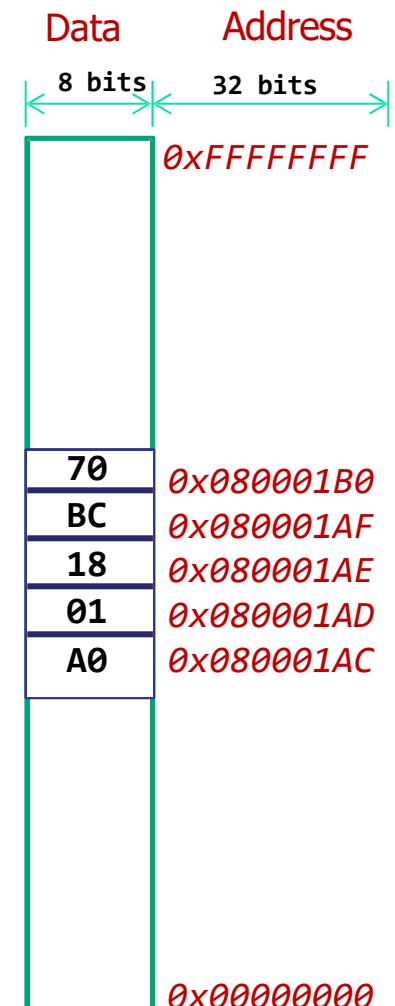
Thread mode
SPSEL-control
0 or 1





Memory

- Memory is arranged as a series of “locations”
 - Each location has a unique “address”
 - Each location holds a byte (**byte-addressable**)
 - e.g. the memory location at address **0x080001B0** contains the byte value **0x70**, i.e., 112
- The number of locations in memory is limited
 - e.g. 4 GB of RAM
 - 1 Gigabyte (GB) = 2^{30} bytes
 - 2^{32} locations → 4,294,967,296 locations!
- Values stored at each location can represent either **program data** or **program instructions**
 - e.g. the value **0x70** might be the code used to tell the processor to add two values together

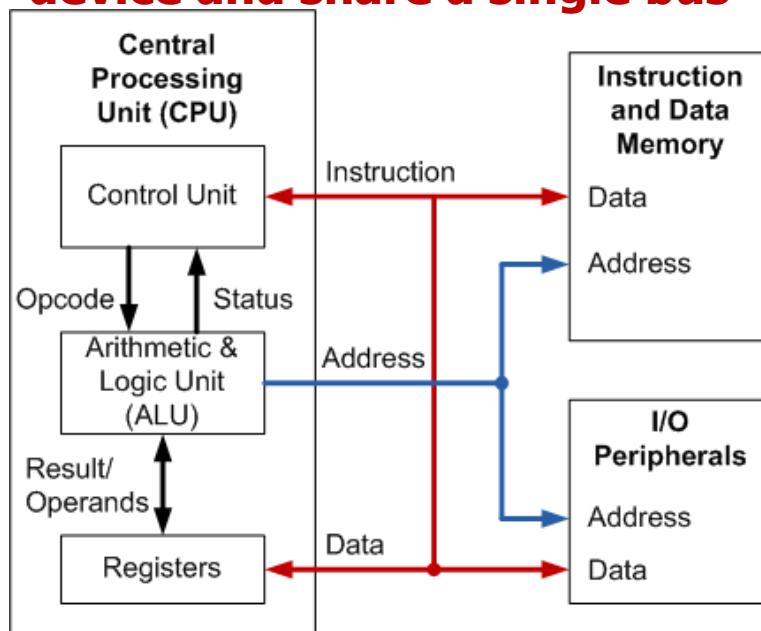




Computer Architecture

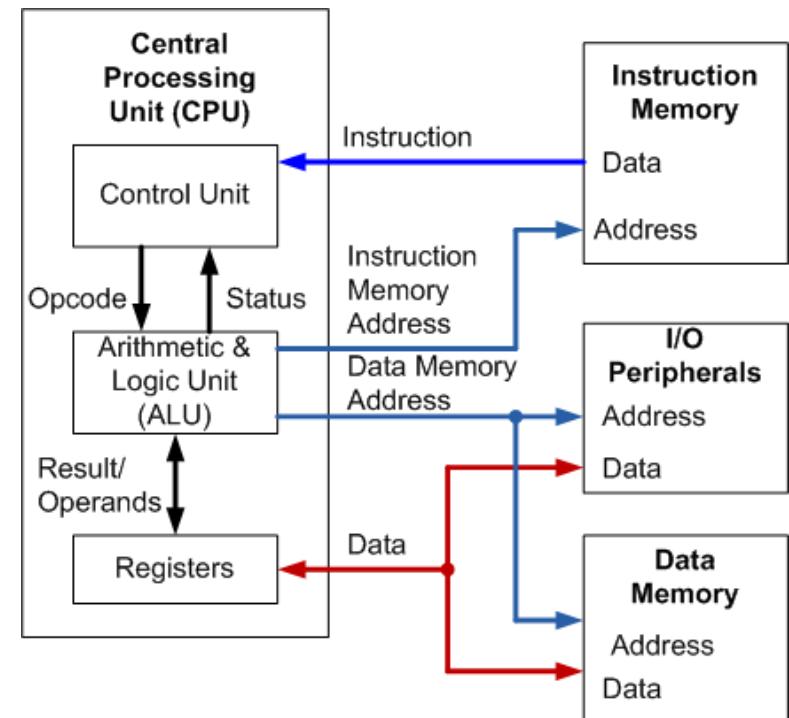
Von Neumann

Instructions and data are stored in the same memory device and share a single bus



Harvard

Data and instructions are stored in separate memory devices. Each has its own buses.





How Programs are Built

C Program

```
int main(void){  
    int i;  
    int total = 0;  
    for (i = 0; i < 10; i++) {  
        total += i;  
    }  
    while(1); // Dead loop  
}
```

Compile

Assembly Program

```
MOVS r1, #0  
MOVS r0, #0  
B check  
loop ADD r1, r1, r0  
ADDS r0, r0, #1  
check CMP r0, #10  
BLT loop  
self B self
```

Assemble

Machine Program

```
0010000100000000  
0010000000000000  
1110000000000001  
0100010000000001  
0001110001000000  
0010100000001010  
1101110011111011  
1011111000000000  
1110011111111110
```

High-level language

- Level of abstraction closer to problem domain
- Provides for productivity and portability

Assembly language

- Textual representation of machine instructions

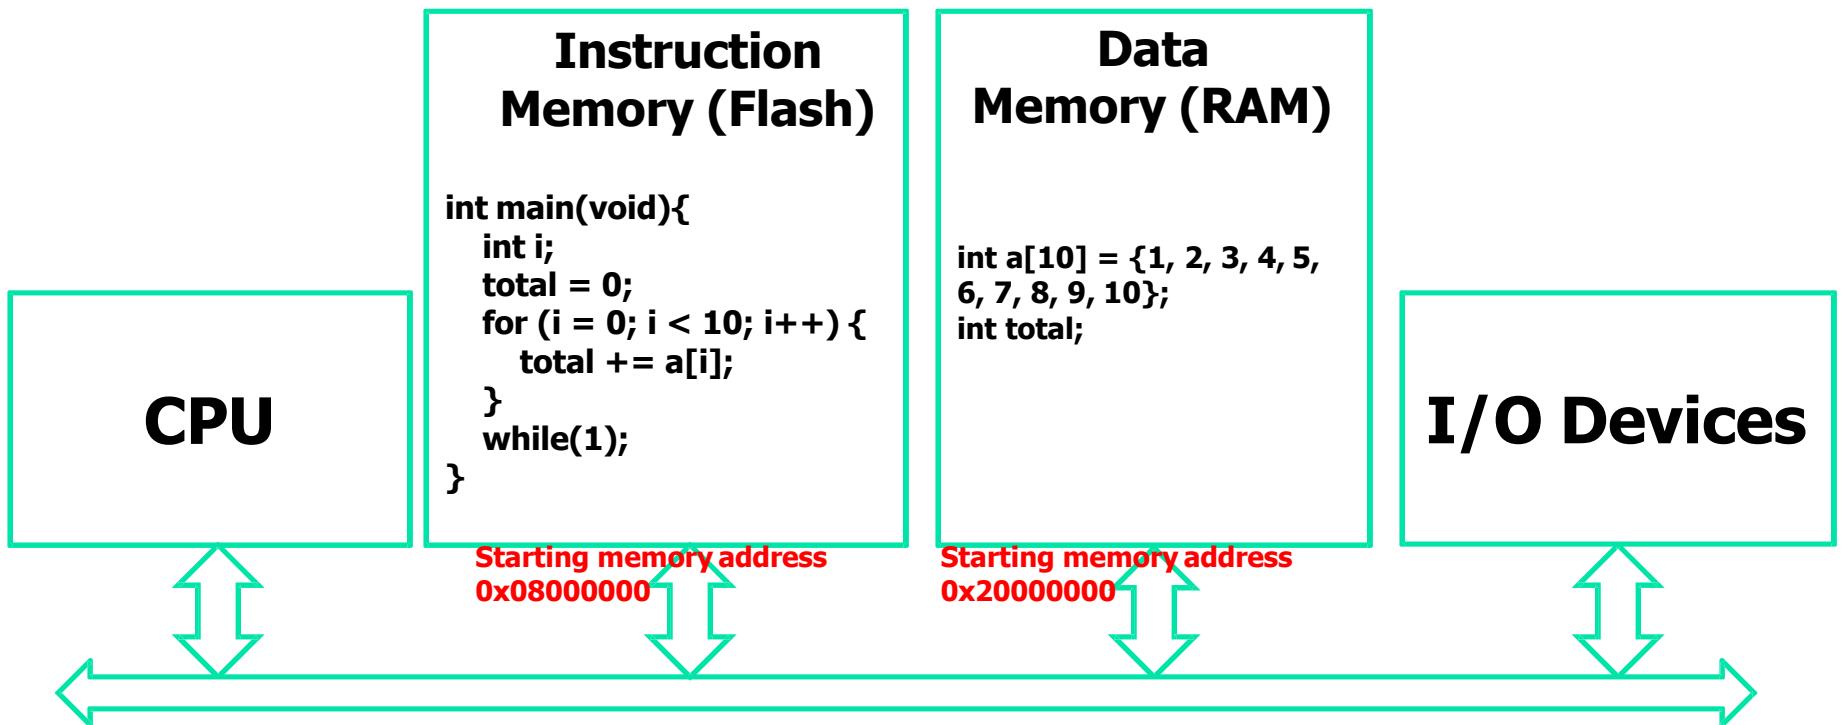
Hardware representation

- Binary digits (bits)
- Encoded instructions and data



Example:

Calculate the Sum of an Array

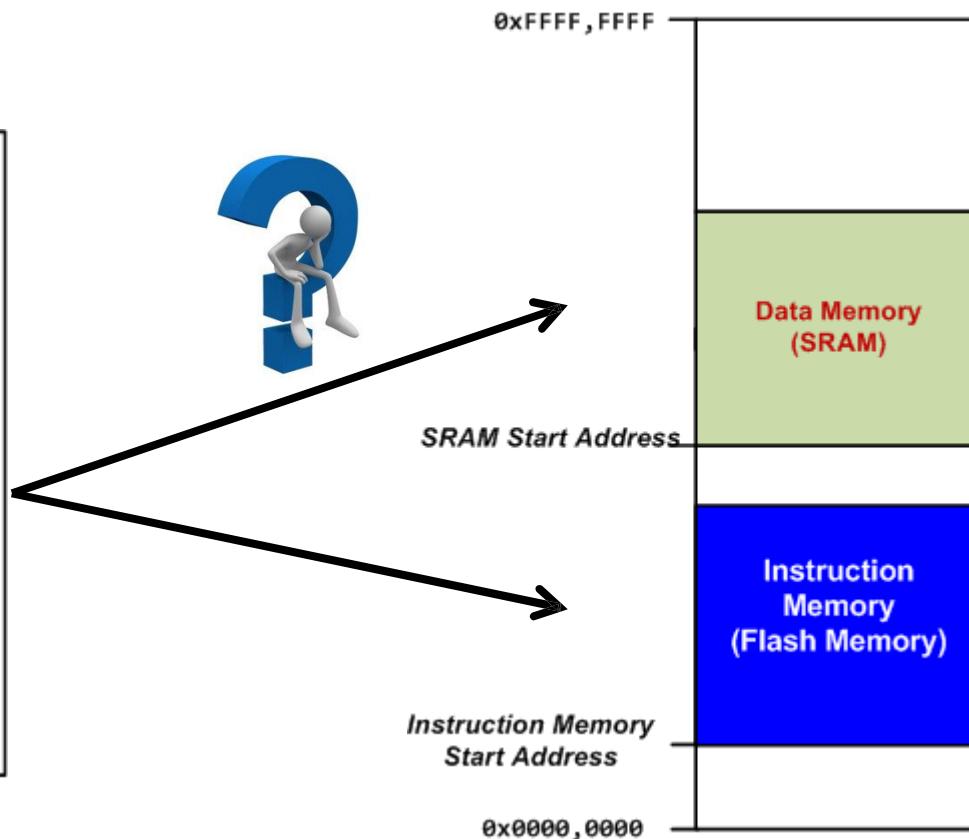




Loading Code and Data into Memory

```
int counter;  
  
int a[5] = {1, 2, 3, 4, 5};  
  
int main(void){  
    int i;  
  
    int b[5];  
  
    counter = 0;  
    for (i = 0; i < 5; i++){  
        b[i] = a[i];  
        counter++;  
    }  
    while(1);  
}
```

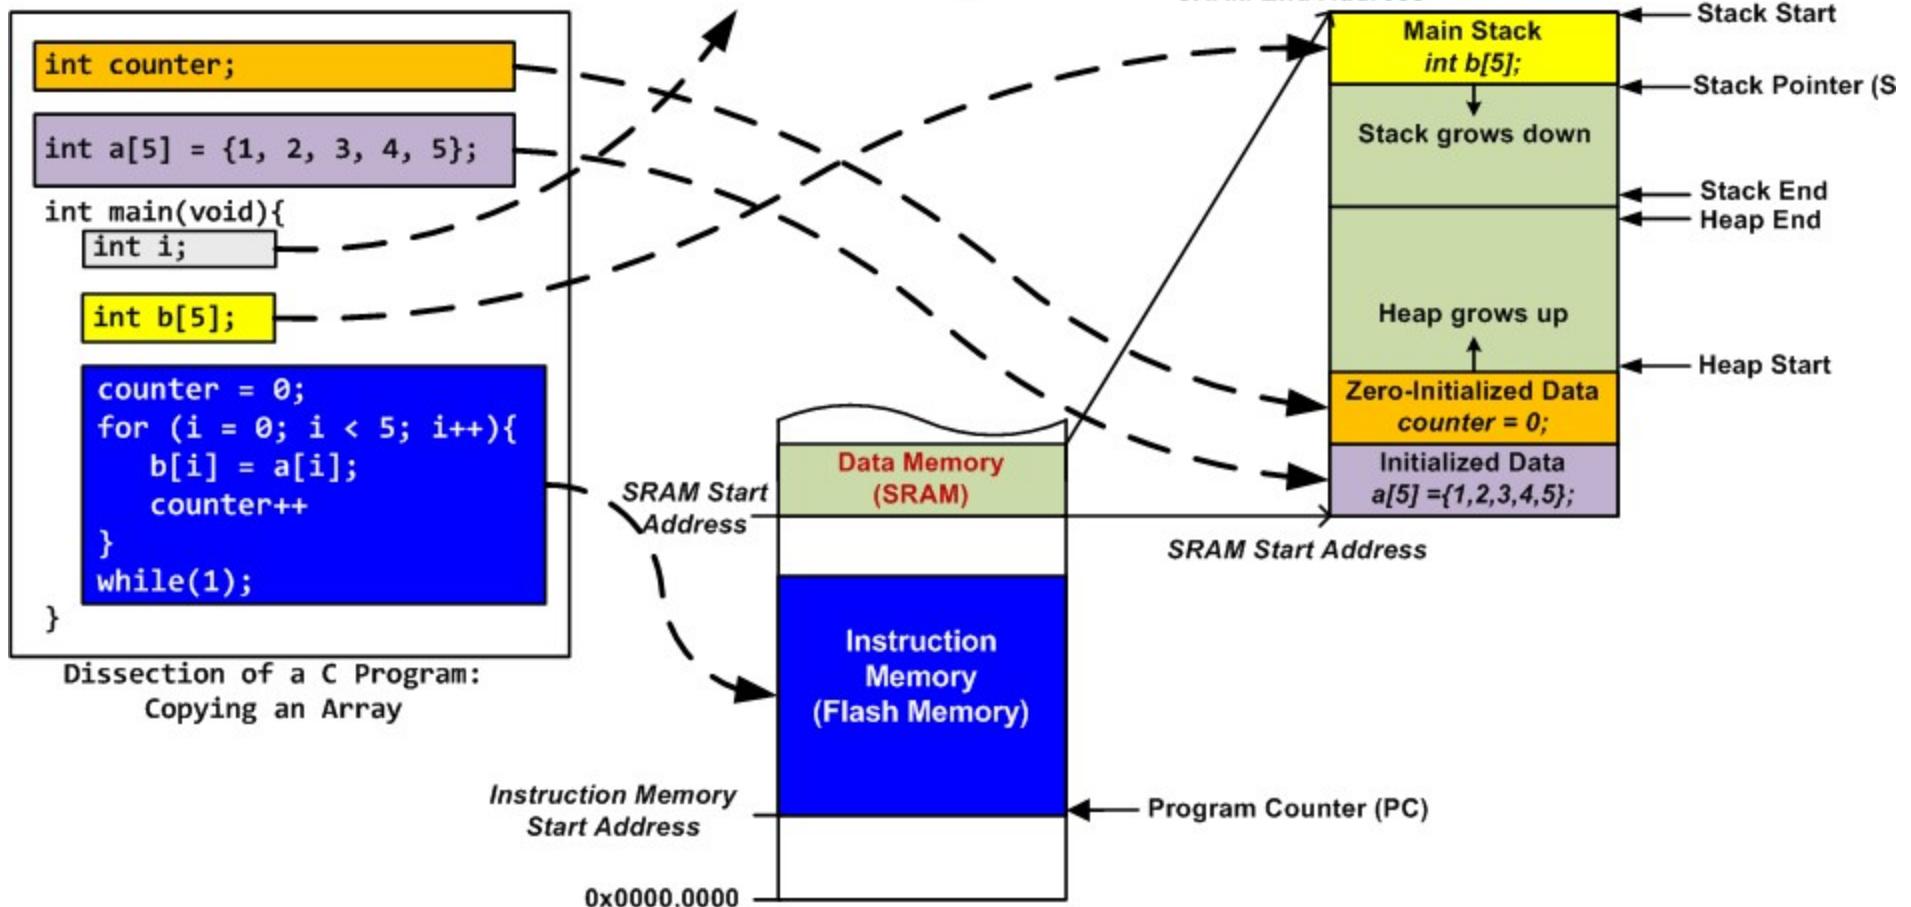
C Program:
Copying an Array





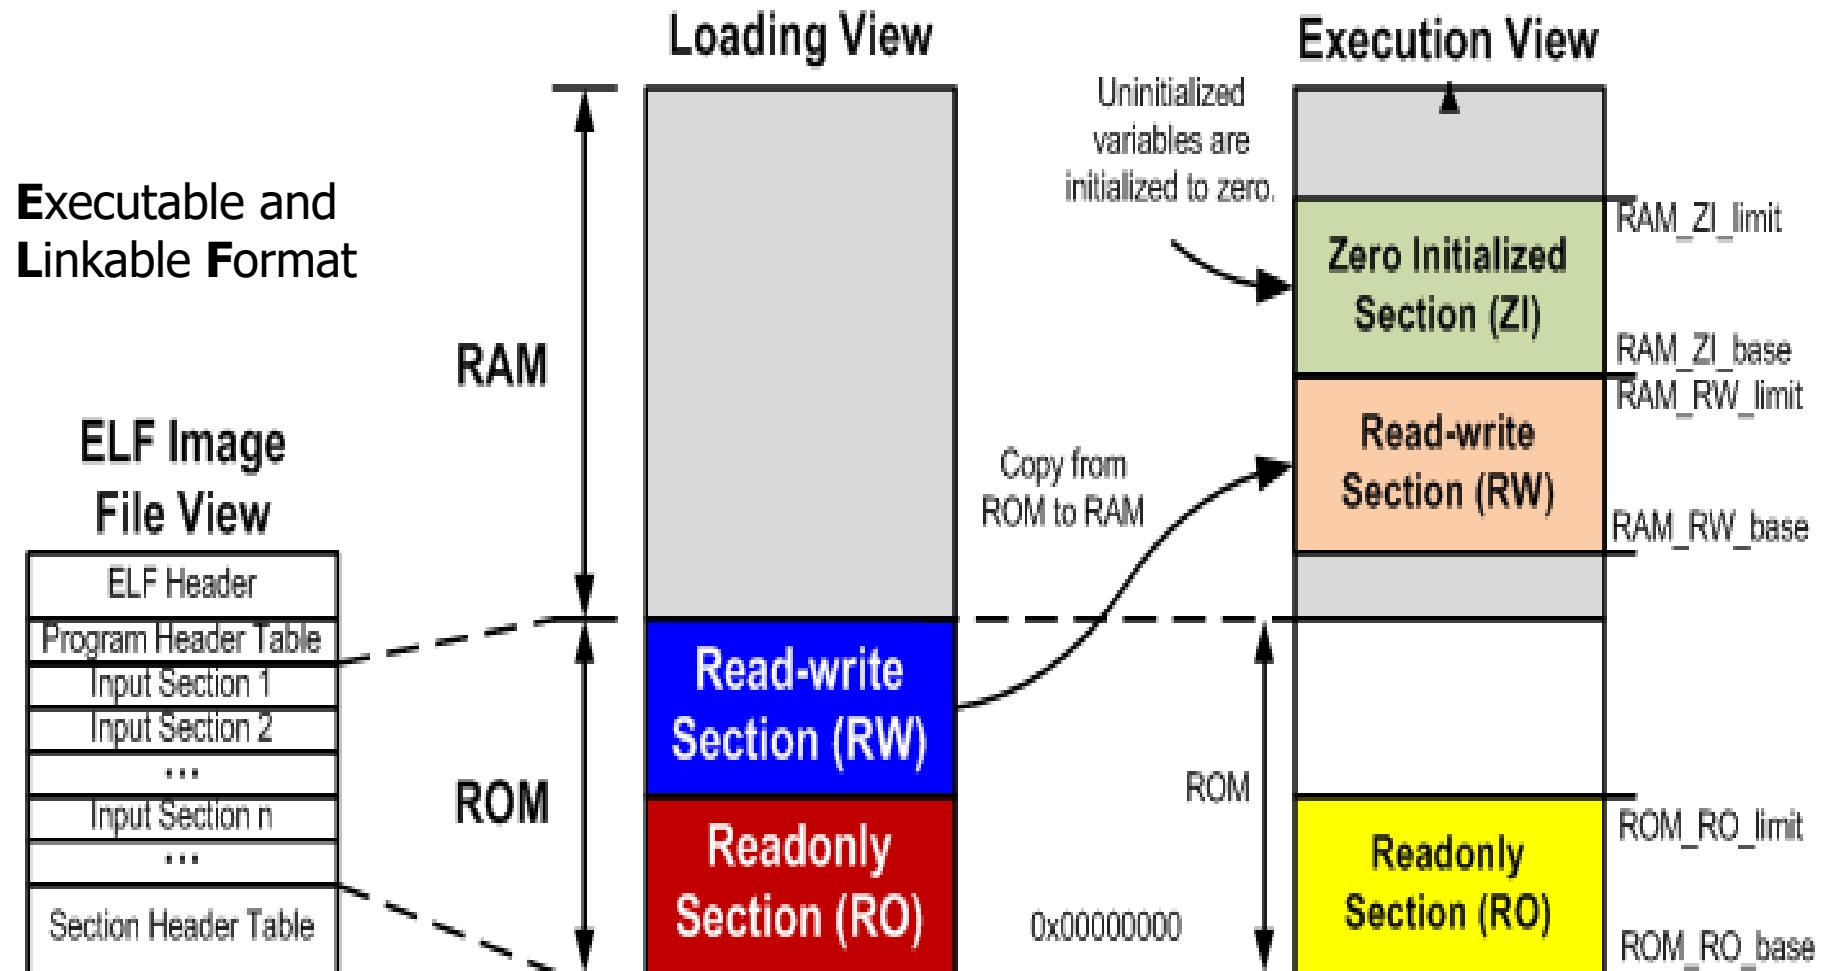
Loading Code and Data into Memory

To improve performance, some variables are not stored in memory.
Variable i will be stored in a register.





View of a Binary Program





STM32L476 Discovery – HMI

Integrated ST-Link/V2-1 (for programming and debugging)

LCD 96 segments

Motion Mems (9-axis)

push buttons and joystick,
2 color LEDs

Quad SPI NOR Flash
16 MB

USB OTG connector

from st.com



STM32L476 Discovery - Audio and connector

APC connector (for Apple connector)

MFX to auto-measure power consumption

Direct access to all MCU I/Os

Audio Codec and 3.5 mm connector

Microphone Mems



from st.com



STM32L4

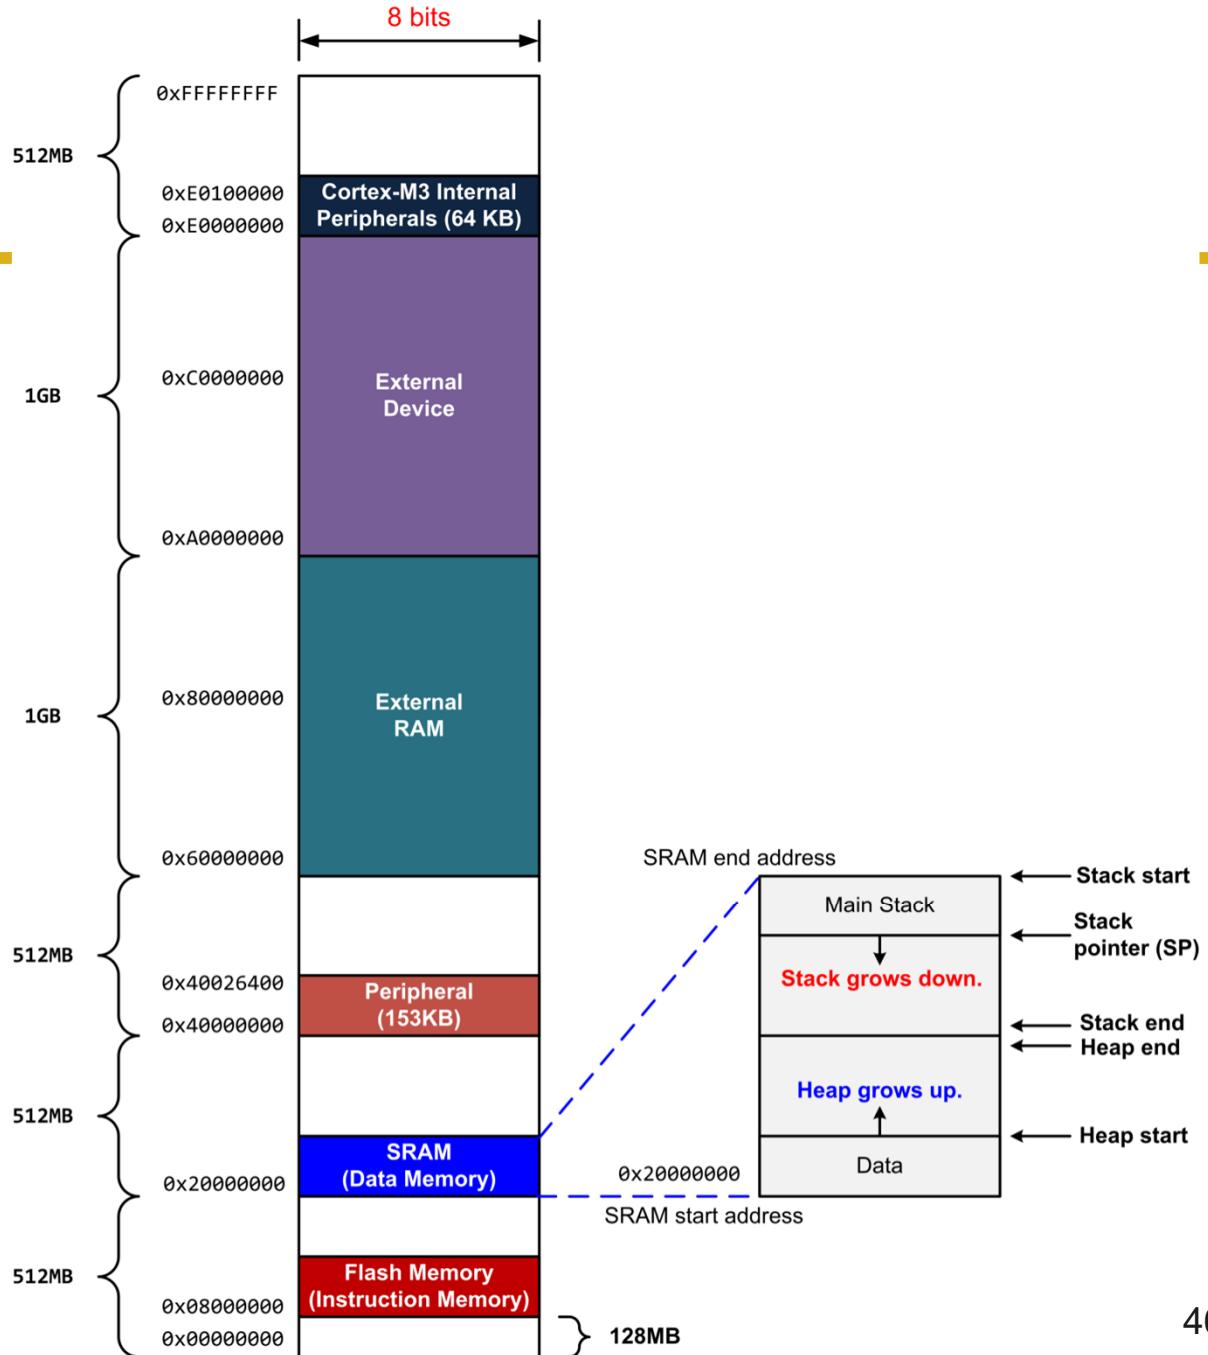


Parallel Interface FSMC 8/16-bit (TFT-LCD, SRAM, NOR, NAND)	Cortex-M4 80 MHz FPU MPU ETM DMA ART Accelerator™ Up to 1-Mbyte Flash with ECC Dual Bank 128-Kbyte RAM	Connectivity USB OTG, 1x SD/SDIO/MMC, 3 x SPI, 3 x I ² C, 1x CAN, 1 x Quad SPI, 5 x USART + 1 x ULP
Display LCD driver 8 x 40		Digital AES (256-bit), TRNG, 2 x SAI, DFSDM (8 channels)
Timers 17 timers including: 2 x 16-bit advanced motor control timers 2 x ULP timers 7 x 16-bit-timers 2 x 32-bit timers		Analog 3 x 16-bit ADC, 2 x DAC, 2 x comparators, 2 x op amps 1 x temperature sensor
I/Os Up to 114 I/Os Touch-sensing controller		

from st.com



Memory Map



Embedded Systems with ARM Cortex-M Microcontrollers in Assembly Language and C



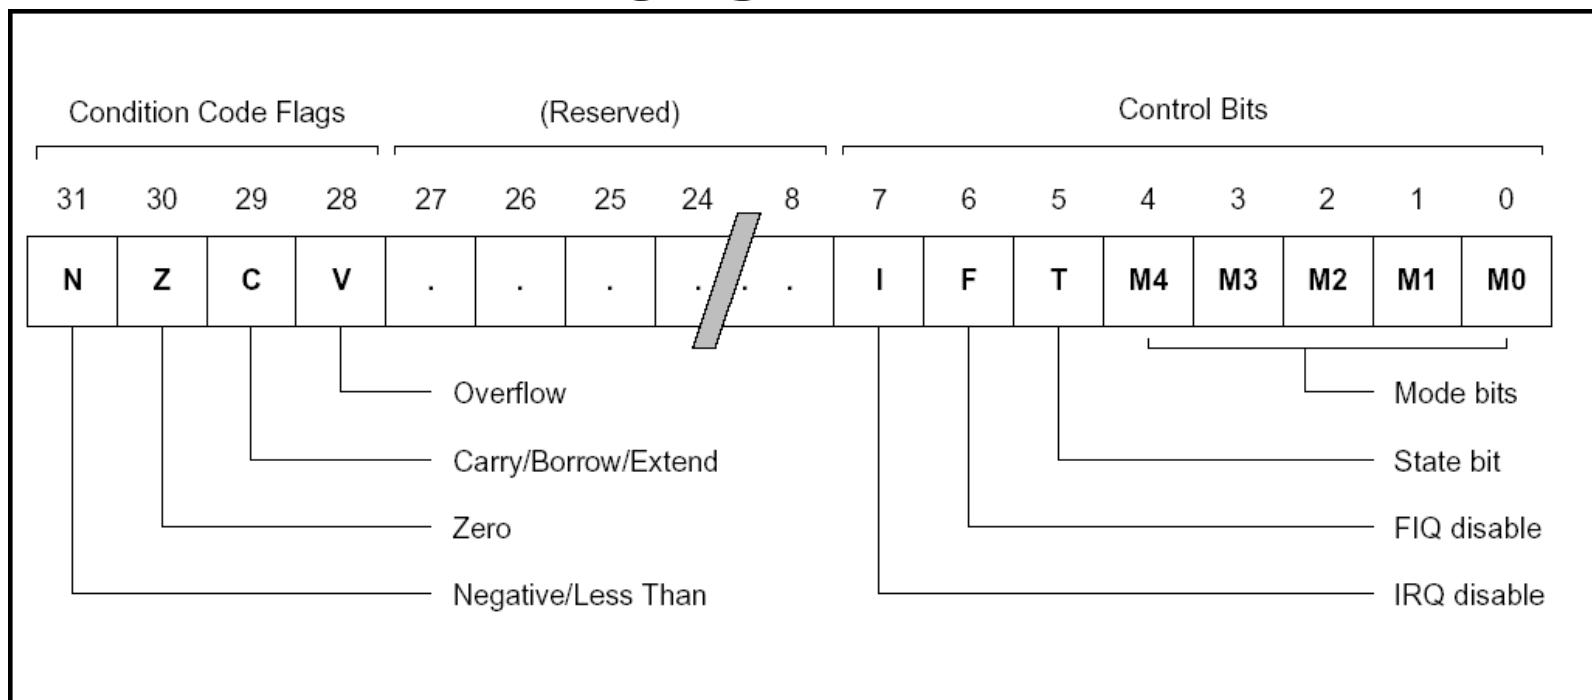
Chapter 2 Data Representation

Edited: Waqas Majeed
©Yifeng Zhu



ARM Current Program Status Register

CPSR





ARM CPSR Format

- N (Negative)
- Z (Zero)
- C (Carry)
- V (oVerflow)
- M0-M4 mode – control processor mode (see manual)
- T – control instruction set
 - T = 1 – instruction stream is 16-bit Thumb instructions
 - T = 0 – instruction stream is 32-bit ARM instructions
- I, F – interrupt enables



Signed Integer Representation Overview

- Three ways to represent signed binary integers:
 - Signed magnitude
 - $\text{value} = (-1)^{\text{sign}} \times \text{Magnitude}$
 - One's complement ($\tilde{\alpha}$)
 - $\alpha + \tilde{\alpha} = 2^n - 1$
 - Two's complement ($\bar{\alpha}$)
 - $\alpha + \bar{\alpha} = 2^n$

	Sign-and-Magnitude	One's Complement	Two's Complement
Range	$[-2^{n-1} + 1, 2^{n-1} - 1]$	$[-2^{n-1} + 1, 2^{n-1} - 1]$	$[-2^{n-1}, 2^{n-1} - 1]$
Zero	Two zeroes (± 0)	Two zeroes (± 0)	One zero
Unique Numbers	$2^n - 1$	$2^n - 1$	2^n



Signed Integer Representation

Method 3: Two's Complement

Assume a four-bit system:

Expression	Result	Carry ?	Overflow?	Correct Result?
$0100 + 0010$	0110	No	No	Yes
$0100 + 0110$	1010	No	Yes	Yes(unsigned), No(signed)
$1100 + 1110$	1010	Yes	No	No(unsigned), Yes(signed)
$1100 + 1010$	0110	Yes	Yes	No



Why use Two's Complement

Two's complement simplifies hardware

Operation	Are signed and unsigned operations the same?
Addition	Yes
Subtraction	Yes
Multiplication	Yes if the product is required to keep the same number of bits as operands
Division	No



Condition Codes

Bit	Name	Meaning after add or sub
N	negative	result is negative
Z	zero	result is zero
V	overflow	signed overflow
C	carry	unsigned overflow

C is set upon an **unsigned** addition if the answer is wrong

C is cleared upon an **unsigned** subtract if the answer is wrong

V is set upon a **signed** addition or subtraction if the answer is wrong

Why do we care about these bits?

Embedded Systems with ARM Cortex-M Microcontrollers in Assembly Language and C

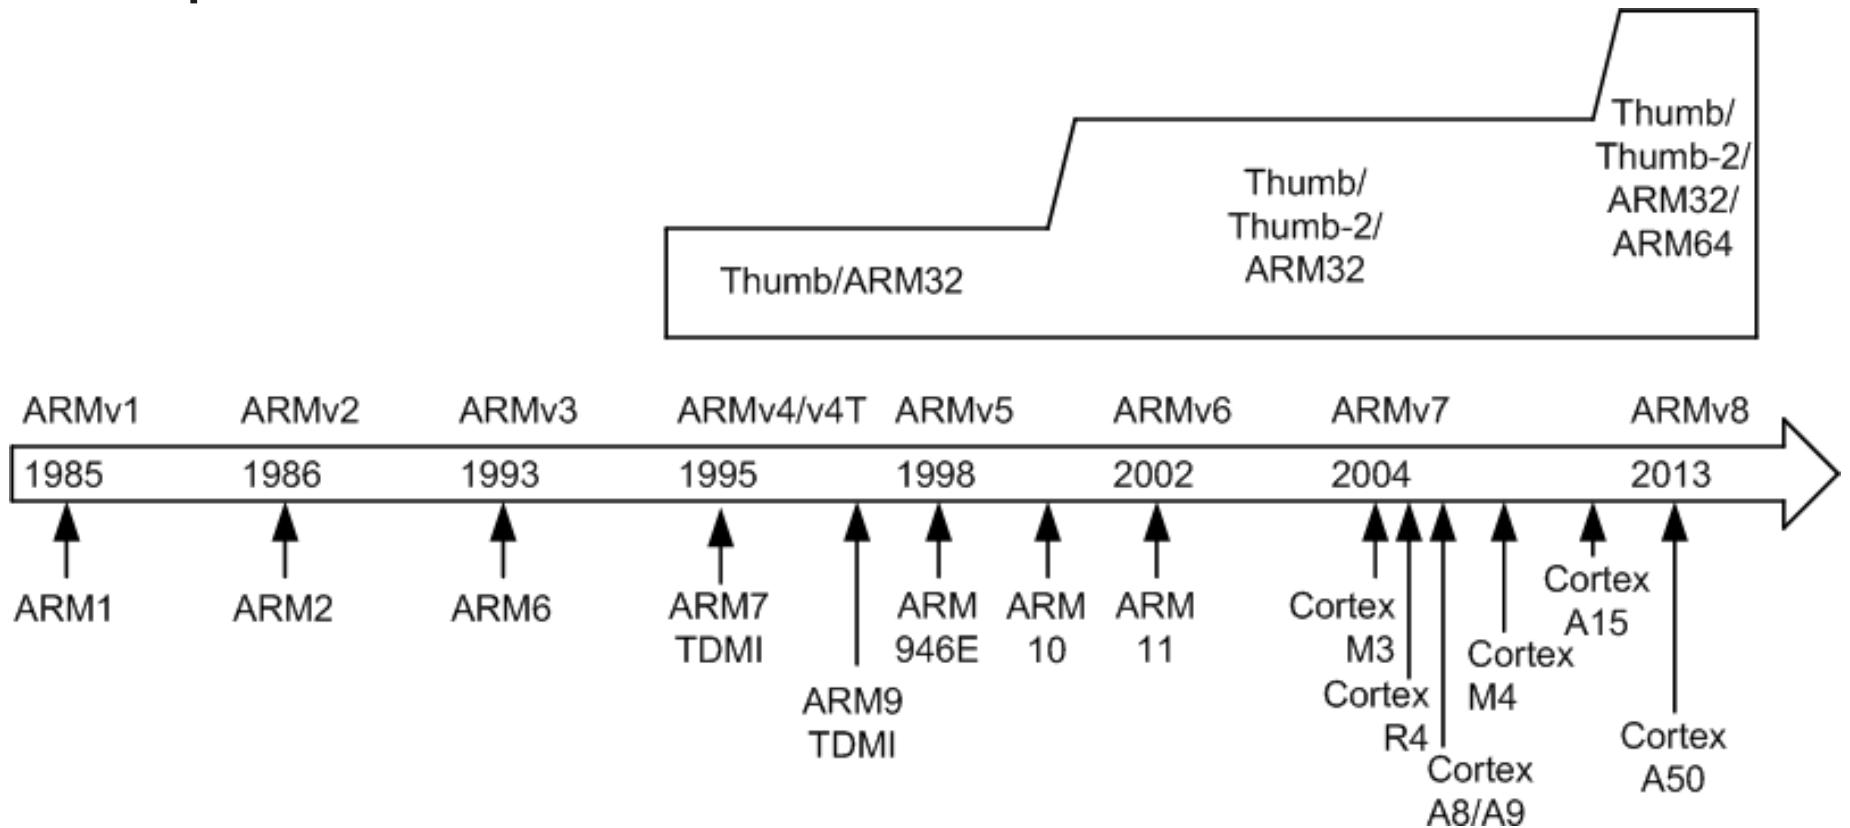


Chapter 3 ARM Instruction Set Architecture

Edited: Waqas Majeed
©Yifeng Zhu



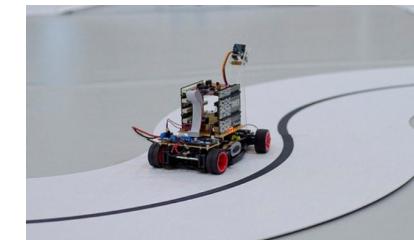
History





ARM Cortex Processors

- ARM Cortex-A family:
 - Applications processors
 - Support OS and high-performance applications
 - Such as Smartphones, Smart TV
- ARM Cortex-R family:
 - Real-time processors with high performance and high reliability
 - Support real-time processing and mission-critical control
- ARM Cortex-M family:
 - Microcontroller
 - Cost-sensitive, support SoC



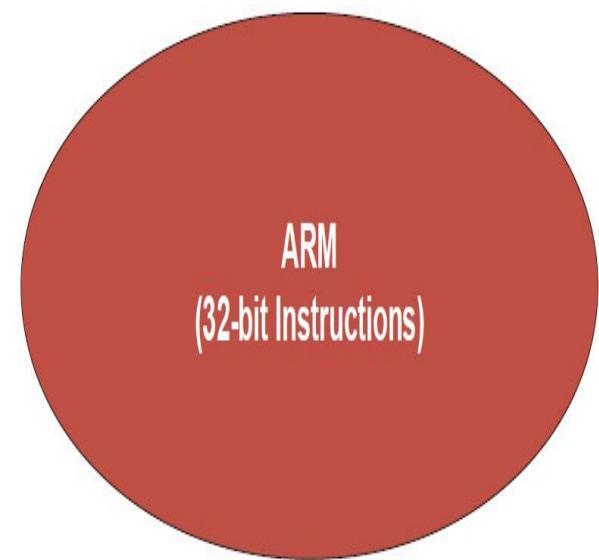
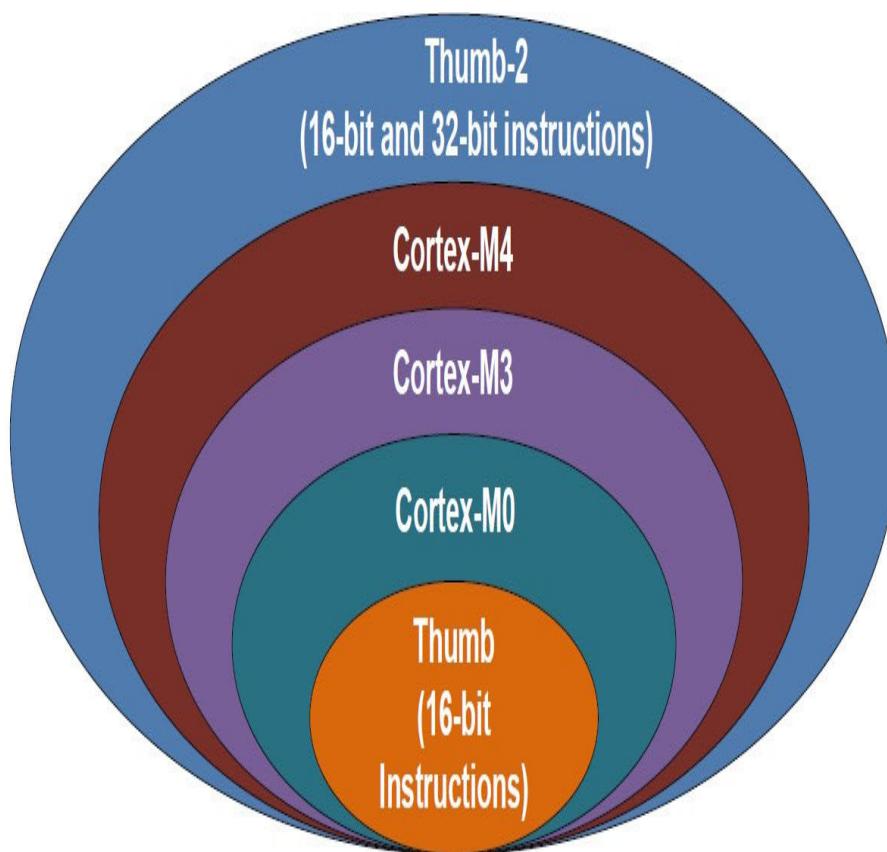


ARM Assembly Instruction Sets

- ARM processors support 4 different assembly instruction sets
 1. Thumb – 16-bit length, improve code density
 2. ARM32 – 32-bit length, more flexible and faster than Thumb, lower code density
 3. Thumb-2 – Mixture of 16-bit and 32-bit instructions, compromise between Thumb and ARM32
 4. ARM64 – 64-bit length, used for desktops and servers



Instruction Sets





Instruction Sets

PKH	QADD	QADD16	QADD8	QASX	QDADD	QDSUB	QSAX	QSUB
QSUB16	QSUB8	SADD16	SADD8	SASX	SEL	SHADD16	SHADD8	SHASX
SHSAX	SHSUB16	SHSUB8	SMLABB	SMLABT	SMLATB	SMLATT	SMLAD	SMLALBB
CORTEX-M0/M1								
ADC	ADD	ADR	AND	ASR	B	CLZ	SMLALBT	SMLALTB
BFC	BFI	BIC	CDP	CLREX	CBNZ	CBZ	CMN	SMLALDT
CMP				DBG	EOR	LDC	SMLAWB	SMLAWT
LDMIA	BLKPT	BLX	ADC	ADD	ADR		SMLSD	SMLSID
LDRBT	BX	CPS	AND	ASR	B		SMMLA	SMMLS
LDREXH	DMB		BL	BIC			SMMUL	SMUAD
LDRSBT	DSB		CMN	CMP	EOR		SMULBB	SMULBT
MCR	ISB		LDR	LDRB	LDM		SMULTB	SMULTT
MCRR	MRS		LDRH	LDRSB	LDRSH		SMULWB	SMULWT
MRC	MSR		LSL	LSR	MOV		SMUSD	SSAT16
NOP	NOP	REV	MUL	MVN	ORR		SSAX	SSUB16
PLDW	REV16	REVSH	POP	PUSH	ROR		SSUB8	SXTAB
RBIT	SEV	SXTB	RSB	SBC	STM		SXTAB16	SXTAH
ROR	SXTH	UXTB	STR	STRB	STRH		SXTB16	UADD16
SBFX	UXTH	WFE	SUB	SVC	TST		UADD8	UASX
SMULL	WFI	YIELD					UHADD16	UHADD8
STMDB							UHASX	UHSAX
STRD	STREX	STREXB	STREXH				UHSUB16	UHSUB8
SUB	SXTB	SXTH	TBB				UMAAL	UQADD16
UBFX	UDIV	UMLAL	UMULL	USAT			UQADD8	UQASX
WFE	WFI	YIELD	IT				UQSAX	UQSUB16
Cortex-M3								
USAX	USUB16	USUB8	UXTAB	UXTAB16	UXTAH	UXTB16	Cortex-M4	
VABS	VADD	VCMP	VCMPE	VCVT	VCVTR	VDIV	VLD	VLDR
VMLA	VMLS	VMOV	VMRS	VMSR	VMUL	VNEG	VNMLA	VNMLS
VNMUL	VPOP	VPUSH	VSQRT	VSTM	VSTR	VSUB	Cortex-M4F	
from arm.com								

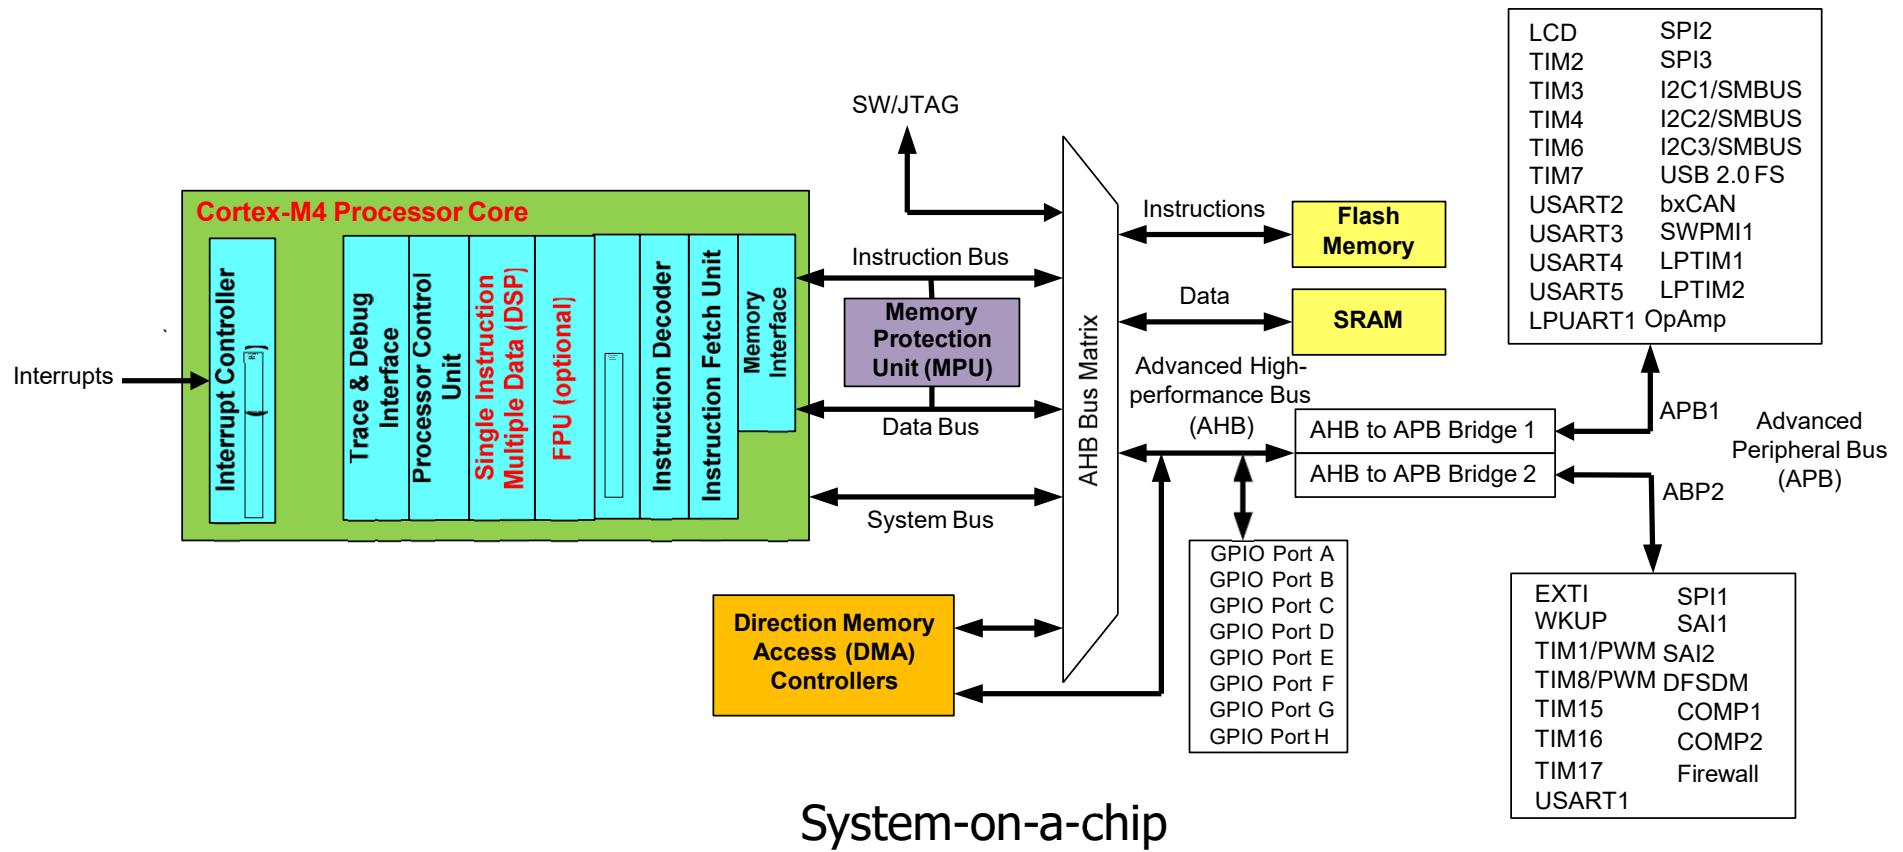


ARM Instruction Set Architecture (ISA)

- ARM processors have a load-store instruction set architecture
 - ALU cannot directly use data stored in memory as operands.
 - ALU source /destination operands must be general-purpose registers.
 - In order to access/modify data in memory, load and store instructions must be used.



ARM Cortex-M4 Organization (STM32L4)





Assembly Instructions Supported

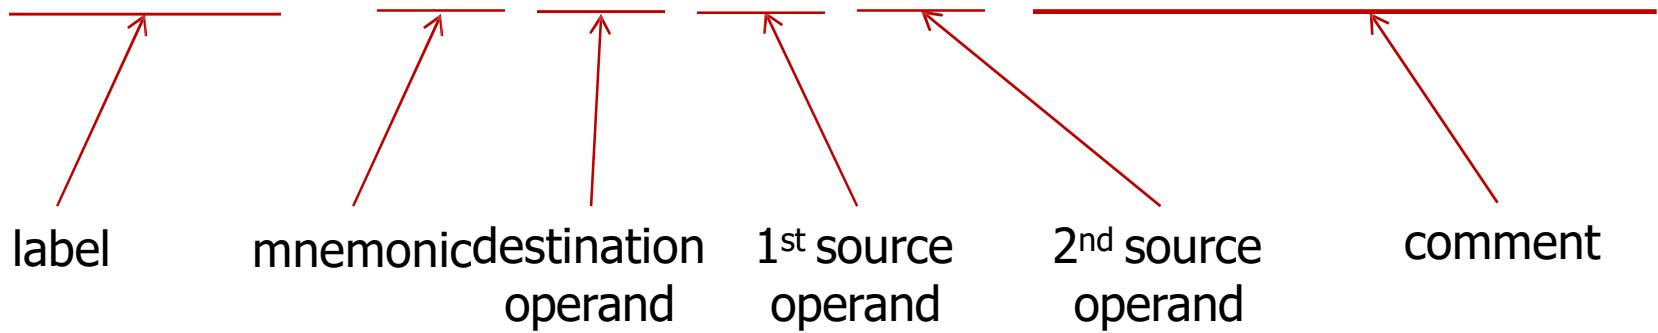
- Arithmetic and logic
 - Add, Subtract, Multiply, Divide, Shift, Rotate
- Data movement
 - Load, Store, Move
- Compare and branch
 - Compare, Test, If-then, Branch, compare and branch on zero
- Miscellaneous
 - Breakpoints, wait for events, interrupt enable/disable, data memory barrier, data synchronization barrier



ARM Instruction Format

label mnemonic operand1, operand2, operand3 ; comments

target ADD r0, r2, r3 ; r0 = r2 + r3





Assembly Directives

- Directives are **NOT** instructions. Instead, they are used to provide key information to compiler.

AREA	Make a new block of data or code
ENTRY	Declare an entry point where the program execution starts
ALIGN	Align data or code to a particular memory boundary
DCB	Allocate one or more bytes (8 bits) of data
DCW	Allocate one or more half-words (16 bits) of data
DCD	Allocate one or more words (32 bits) of data
SPACE	Allocate a zeroed block of memory with a particular size
FILL	Allocate a block of memory and fill with a given value
EQU	Give a symbol name to a numeric constant
RN	Give a symbol name to a register
EXPORT	Declare a symbol and make it referable by other source files
IMPORT	Provide a symbol defined outside the current source file
INCLUDE/GET	Include a separate source file within the current source file
PROC	Declare the start of a procedure
ENDP	Designate the end of a procedure
END	Designate the end of a source file



Directive: Data Allocation

- Data allocation directives are used to reserve space in data memory for variables and set their initial content.

Directive	Description	Memory Space
DCB	Define Constant Byte	Reserve 8-bit values
DCW	Define Constant Half-word	Reserve 16-bit values
DCD	Define Constant Word	Reserve 32-bit values
DCQ	Define Constant	Reserve 64-bit values
DCFS	Define single-precision floating-point numbers	Reserve 32-bit values
DCFD	Define double-precision floating-point numbers	Reserve 64-bit values
SPACE	Defined Zeroed Bytes	Reserve a number of zeroed bytes
FILL	Defined Initialized Bytes	Reserve and fill each byte with a value

**Embedded Systems with ARM Cortex-M Microcontrollers in Assembly
Language and C**



Chapter 4
ARM Arithmetic and Logic Instructions

Edited: Waqas Majeed
©Yifeng Zhu



Overview: Arithmetic and Logic Instructions

- Shift
 - LSL (logic shift left), LSR (logic shift right), ASR (arithmetic shift right), ROR (rotate right), RRX (rotate right with extend)
- Logic
 - AND (bitwise and), ORR (bitwise or), EOR (bitwise exclusive or), ORN (bitwise or not), MVN (move not)
- Bit set/clear
 - BFC (bit field clear), BFI (bit field insert), BIC (bit clear), CLZ (count leading zeroes)
- Bit/byte reordering
 - RBIT (reverse bit order in a word), REV (reverse byte order in a word), REV16 (reverse byte order in each half-word independently), REVSH (reverse byte order in each half-word independently)
- Addition
 - ADD, ADC (add with carry)
- Subtraction
 - SUB, RSB (reverse subtract), SBC (subtract with carry)
- Multiplication
 - MUL (multiply), MLA (multiply-accumulate), MLS (multiply-subtract), SMULL (signed long multiply-accumulate), SMLAL (signed long multiply-accumulate), UMULL (unsigned long multiply-subtract), UMLAL (unsigned long multiply-subtract)
- Division
 - SDIV (signed), UDIV (unsigned)
- Saturation
 - SSAT (signed), USAT (unsigned)
- Sign extension
 - SXTB (signed), SXTH, UXTB, UXTH
- Bit field extract
 - SBFX (signed), UBFX (unsigned)
- Syntax
 - <Operation>{<cond>} {S} Rd, Rn, Operand2



Example: Add

- Unified Assembler Language (UAL) Syntax

```
ADD r1, r2, r3      ; r1 = r2 + r3
```

```
ADD r1, r2, #4      ; r1 = r2 + 4
```

- Traditional Thumb Syntax

```
ADD r1, r3          ; r1 = r1 + r3
```

```
ADD r1, #15         ; r1 = r1 + 15
```



Program Status Register (PSR)

- Cortex-M processors have 5 status flags
 - Negative (**N**)
 - Set if ALU result is negative.
 - Zero (**Z**)
 - Set if ALU result is zero.
 - Overflow (**V**)
 - Set if overflow occurs for signed addition/subtraction.
 - Carry (**C**)
 - Set (=1) if carry occurs for unsigned addition and set (=0) if borrow occurs for unsigned subtraction.
 - Saturation (**Q**)
 - Set if SSAT or USAT instruction causes saturation.



Program Status Register (PSR)

- Application PSR (**APSR**), Interrupt PSR (**IPSR**), Execution PSR (**EPSR**)

	31	30	29	28	27	26:2 5	24	23:2 0	19:1 6	15:1 0	9	8	7	6	5	4:0	
APSR	N	Z	C	V	Q				GE								
IPSR																Exception Number	
EPSR						ICI/I T	T						ICI/I T				

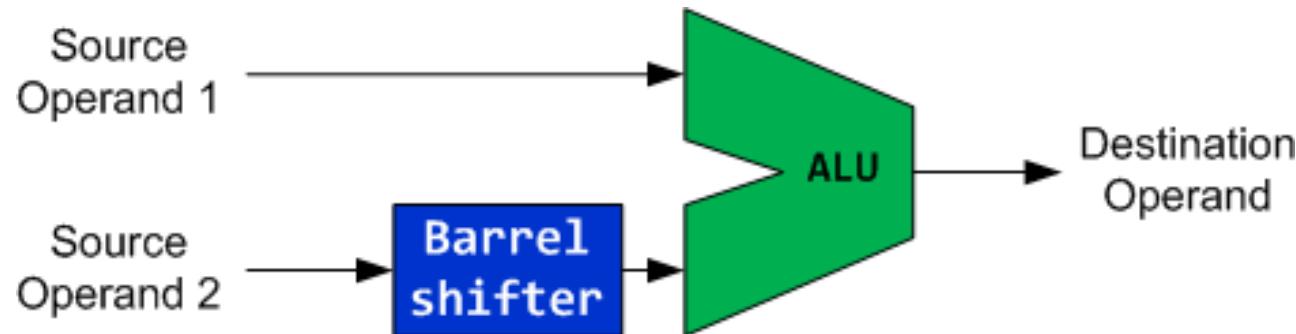
- Combine them together into one register (**PSR**)
- Use PSR in code

	31	30	29	28	27	26:2 5	24	23:2 0	19:1 6	15:1 0	9	8	7	6	5	4:0
PSR	N	Z	C	V	Q	ICI/I T	T		GE	ICI/I T						Exception Number

Note: GE flags are only available on Cortex-M4 and M7



Barrel Shifter



- The 2nd operand of ALU has a special hardware called **Barrel shifter** that enables quick shift and rotate operations on 2nd operand
- Example:

ADD r1, r0, r0, LSL #3 ; $r1 = r0 + r0 \ll 3 = 9 \times r0$



Shift and Rotate Instructions

Logical Shift Left (**LSL**)



Logical Shift Right (**LSR**)



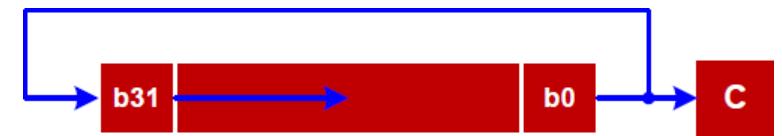
Rotate Right Extended (**RRX**)



Arithmetic Shift Right (**ASR**)



Rotate Right (**ROR**)



Why is there rotate right but no rotate left?

Rotate left can be replaced by a rotate right with a different rotate offset

Rotate left n bits = Rotate right 32 - n



Barrel Shifter

- Examples:

- ADD r1, r0, r0, **LSL #3**

; r1 = r0 + r0 << 3 = r0 + 8 × r0

- ADD r1, r0, r0, **LSR #3**

; r1 = r0 + r0 >> 3 = r0 + r0/8 (unsigned)

- ADD r1, r0, r0, **ASR #3**

; r1 = r0 + r0 >> 3 = r0 + r0/8 (signed)

- Use Barrel shifter to speed up the application

ADD r1, r0, r0, **LSL #3** <=> MOV r2, #9 ; r2 = 9

MUL r1, r0, r2 ; r1 = r0 * 9

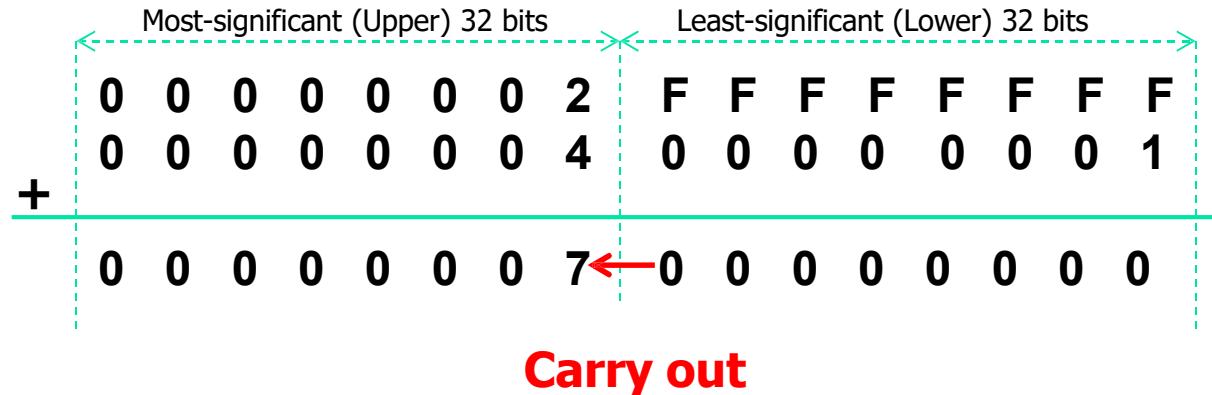


Arithmetic Operations

ADD {Rd,} Rn, Op2	Add. $Rd \leftarrow Rn + Op2$
ADC {Rd,} Rn, Op2	Add with carry. $Rd \leftarrow Rn + Op2 + \text{Carry}$
SUB {Rd,} Rn, Op2	Subtract. $Rd \leftarrow Rn - Op2$
SBC {Rd,} Rn, Op2	Subtract with carry. $Rd \leftarrow Rn - Op2 + \text{Carry} - 1$
RSB {Rd,} Rn, Op2	Reverse subtract. $Rd \leftarrow Op2 - Rn$
MUL {Rd,} Rn, Rm	Multiply. $Rd \leftarrow (Rn \times Rm)[31:0]$
MLA Rd, Rn, Rm, Ra	Multiply with accumulate. $Rd \leftarrow (Ra + (Rn \times Rm))[31:0]$
MLS Rd, Rn, Rm, Ra	Multiply and subtract, $Rd \leftarrow (Ra - (Rn \times Rm))[31:0]$
SDIV {Rd,} Rn, Rm	Signed divide. $Rd \leftarrow Rn / Rm$
UDIV {Rd,} Rn, Rm	Unsigned divide. $Rd \leftarrow Rn / Rm$
SSAT Rd, #n, Rm {,shift #s}	Signed saturate
USAT Rd, #n, Rm {,shift #s}	Unsigned saturate



Example: 64-bit Addition



- A register can only store 32 bits
- A 64-bit integer needs two registers
- Split 64-bit addition into two 32-bit additions



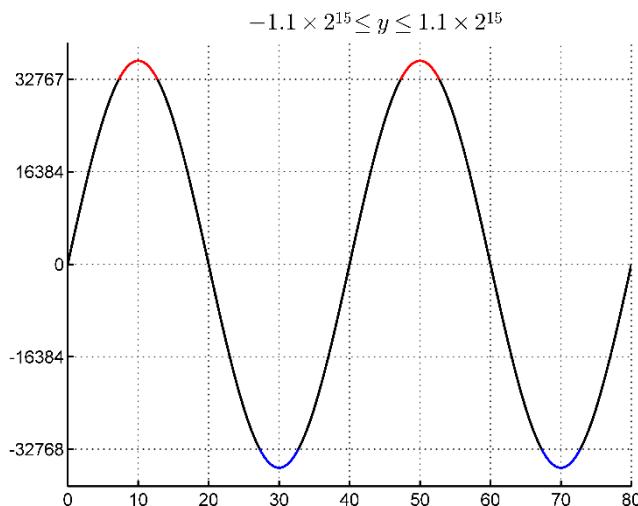
Saturating Instruction: **SSAT** and **USAT**

- Syntax:
 - op{cond} Rd, #n, Rm{}, shift}
- **SSAT** saturates a signed value to the signed range $-2^{n-1} \leq x \leq 2^{n-1}-1$.
$$SAT(x) = \begin{cases} 2^{n-1} - 1 & \text{if } x > 2^{n-1} - 1 \\ -2^{n-1} & \text{if } x < 2^{n-1} \\ x & \text{otherwise} \end{cases}$$
- **USAT** saturates a signed value to the unsigned range $0 \leq x \leq 2^n - 1$.
$$USAT(x) = \begin{cases} 2^{n-1} - 1 & \text{if } x > 2^{n-1} - 1 \\ 0 & \text{if } x \leq 0 \\ x & \text{otherwise} \end{cases}$$
- Examples:
 - SSAT r2, #11, r1 ; output range: $-2^{10} \leq r2 \leq 2^{10} - 1$
 - USAT r2, #11, r3 ; output range: $0 \leq r2 \leq 2^{11} - 1$



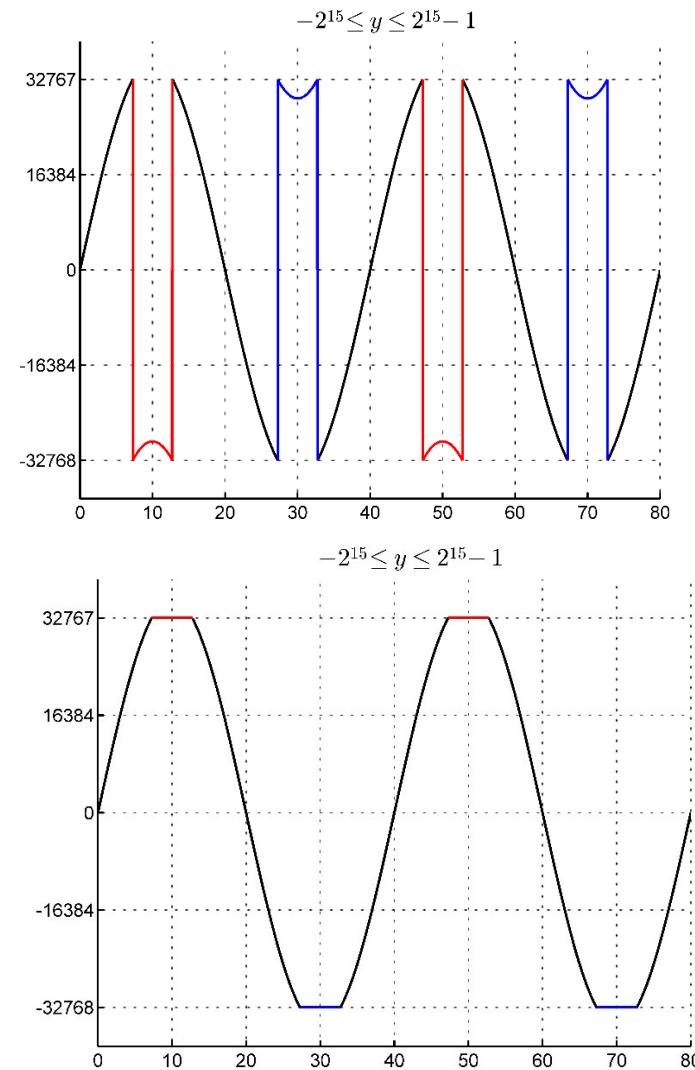
Example of Saturation

Assume data are limited to **16** bits



Without
saturation

With
saturation





Bitwise Logic

AND {Rd,} Rn, Op2	Bitwise logic AND. $Rd \leftarrow Rn \ \& \ \text{operand2}$
ORR {Rd,} Rn, Op2	Bitwise logic OR. $Rd \leftarrow Rn \ \ \text{operand2}$
EOR {Rd,} Rn, Op2	Bitwise logic exclusive OR. $Rd \leftarrow Rn \ ^\wedge \ \text{operand2}$
ORN {Rd,} Rn, Op2	Bitwise logic OR NOT. $Rd \leftarrow Rn \ \ (\text{NOT } \text{operand2})$
BIC {Rd,} Rn, Op2	Bit clear. $Rd \leftarrow Rn \ \& \ \text{NOT operand2}$
BFC Rd, #lsb, #width	Bit field clear. $Rd[(\text{width}+\text{lsb}-1):\text{lsb}] \leftarrow 0$
BFI Rd, Rn, #lsb, #width	Bit field insert. $Rd[(\text{width}+\text{lsb}-1):\text{lsb}] \leftarrow Rn[(\text{width}-1):0]$
MVN Rd, Op2	Move NOT, logically negate all bits. $Rd \leftarrow 0xFFFFFFFF \ EOR \text{Op2}$



Bit Operators (`&`, `|`, `~`) vs Boolean Operators (`&&`, `||`, `!`)

A && B	Boolean and	A & B	Bitwise and
A B	Boolean or	A B	Bitwise or
!B	Boolean not	~B	Bitwise not

- The Boolean operators perform word-wide operations, not bitwise
- For example
 - `"0x10 & 0x01" = 0x00`, but `"0x10 && 0x01" = 0x01`.
 - `"~0x01" = 0xFFFFFFF`E, but `"!0x01" = 0x00`.



Bitwise Logic Instructions

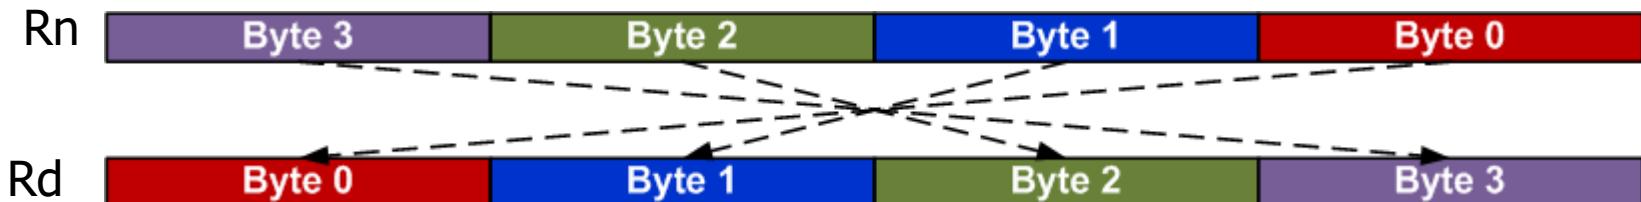
- Logic operations with **S** suffix (ANDS, ORRS, EORS, ORNS and MVNS) will update N, Z and C flags only.
- C flag is updated when 2nd operand of a logic operation
 - Is a constant with MSB = 1.
 - Is shifted resulting shift operation results in C = 1.



Reverse Order

RBIT Rd, Rn	Reverse bit order in a word. for ($i = 0; i < 32; i++$) $Rd[i] \leftarrow RN[31-i]$
REV Rd, Rn	Reverse byte order in a word. $Rd[31:24] \leftarrow Rn[7:0]$, $Rd[23:16] \leftarrow Rn[15:8]$, $Rd[15:8] \leftarrow Rn[23:16]$, $Rd[7:0] \leftarrow Rn[31:24]$
REV16 Rd, Rn	Reverse byte order in each half-word. $Rd[15:8] \leftarrow Rn[7:0]$, $Rd[7:0] \leftarrow Rn[15:8]$, $Rd[31:24] \leftarrow Rn[23:16]$, $Rd[23:16] \leftarrow Rn[31:24]$
REVSH Rd, Rn	Reverse byte order in bottom half-word and sign extend. $Rd[15:8] \leftarrow Rn[7:0]$, $Rd[7:0] \leftarrow Rn[15:8]$, $Rd[31:16] \leftarrow Rn[7] \&& 0xFFFF$

REV Rd, Rn



Example:

LDR R0, =0x12345678

REV R1, R0 ; R1 = 0x78563412



Sign and Zero Extension

SXTB {Rd,} Rm {,ROR #n}	Sign extend a byte. $Rd[31:0] \leftarrow \text{Sign Extend}((Rm \text{ ROR } n)[7:0])$
SXTH {Rd,} Rm {,ROR #n}	Sign extend a half-word. $Rd[31:0] \leftarrow \text{Sign Extend}((Rm \text{ ROR } n)[15:0])$
UXTB {Rd,} Rm {,ROR #n}	Zero extend a byte. $Rd[31:0] \leftarrow \text{Zero Extend}((Rm \text{ ROR } n)[7:0])$
UXTH {Rd,} Rm {,ROR #n}	Zero extend a half-word. $Rd[31:0] \leftarrow \text{Zero Extend}((Rm \text{ ROR } n)[15:0])$

```
LDR R0, =0x55AA8765
SXTB R1, R0      ; R1 = 0x00000065
SXTH R1, R0      ; R1 = 0xFFFF8765
UXTB R1, R0      ; R1 = 0x00000065
UXTH R1, R0      ; R1 = 0x00008765
```



Data Comparison

CMP Rn, Op2	Compare	Set NZCV flags on <i>Rn – Op2</i>
CMN Rn, Op2	Compare negative	Set NZCV flags on <i>Rn + Op2</i>
TST Rn, Op2	Test	Set NZCV flags on <i>Rn AND Op2</i>
TEQ Rn, Op2	Test Equivalence	Set NZCV flags on <i>Rn EOR Op2</i>



Data Movement between Registers

- 2 categories of instructions for moving data between registers
 - Move data between 2 general-purpose registers.
 - Move data between a general-purpose register and a special-purpose register.



Data Movement between Registers

MOV	Rd \leftarrow operand2
MVN	Rd \leftarrow NOT operand2
MRS Rd, spec_reg	Move from special register to general register
MSR spec_reg, Rm	Move from general register to special register

```
MOV r4, r5          ; Copy r5 to r4
MVN r4, r5          ; r4 = bitwise logical NOT of r5
MOV r1, r2, LSL #3  ; r1 = r2 << 3
MOV r0, PC          ; Copy PC (r15) to r0
MOV r1, SP          ; Copy SP (r14) to r1
```



Move Immediate Number to Register

MOVW Rd, #imm16	Move Wide, Rd \leftarrow #imm16
MOVT Rd, #imm16	Move Top, Rd \leftarrow #imm16 << 16
MOV Rd, #const	Move, Rd \leftarrow const

Example: Load a 32-bit number into a register

```
MOVW r0, #0x4321      ; r0 = 0x00004321
MOVT r0, #0x8765      ; r0 = 0x87654321
```

Order does matter!

- **MOVW** will zero the upper halfword
- **MOVT** won't zero the lower halfword

```
MOVT r0, #0x8765      ; r0 = 0x8765xxxx
MOVW r0, #0x4321      ; r0 = 0x00004321
```



Bit Field Extract

- These instructions extract adjacent bits from 1 register and place them in another

SBFX Rd, Rn, #lsb, #width	Signed Bit Field Extract $Rd[(width - 1):0] \leftarrow Rn[(width + lsb - 1):lsb]$ $Rd[31:width] \leftarrow \text{Replicate}(Rn[(width + lsb - 1)])$
UBFX Rd, Rn, #lsb, #width	Unsigned Bit Field Extract $Rd[(width - 1):0] \leftarrow Rn[(width + lsb - 1):lsb]$ $Rd[31:width] \leftarrow \text{Replicate}(0)$

```
; Assume r3 = 0x1234CDEF  
  
UBFX r4, r3, #4, #8      ; r4 = 0x000000DE (zero extension)  
SBFX r4, r3, #4, #8      ; r4 = 0xFFFFFDE (sign extension)
```



Chapter 5

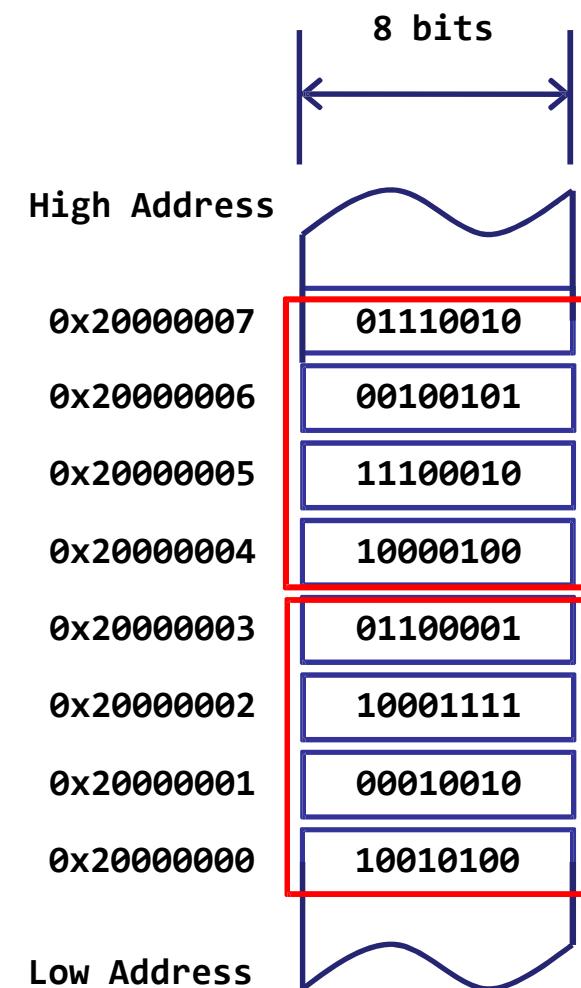
Memory Access



Logical View of Memory

- When we refer to memory locations by address, we can only do so in units of bytes, halfwords or words
- Words
 - $32 \text{ bits} = 4 \text{ bytes} = 1 \text{ word} = 2 \text{ halfwords}$
 - In diagram on the right, we have two words at addresses:
 - 0x20000000
 - 0x20000004
 - Can you store a word anywhere? **NO.**
 - A word can only be stored at an address that's divisible by 4.
 - Memory address of a word is the lowest address of all four bytes in that word.

Word-address mod 4 = 0





Single register data transfer

- You can load/store data of different sizes between memory and registers

LDR	Load Word
LDRB	Load Byte
LDRH	Load Halfword
LDRSB	Load Signed Byte
LDRSH	Load Signed Halfword

STR	Store Word
STRB	Store Lower Byte
STRH	Store Lower Halfword



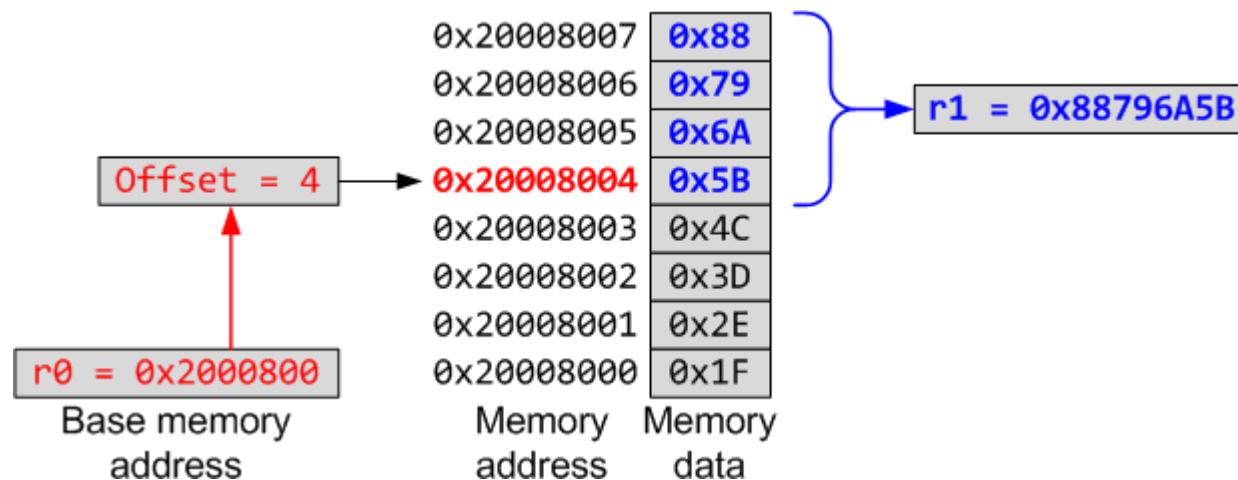
Address

- Address accessed by LDR/STR is specified by a base register **plus an offset**
- For word and unsigned byte accesses, offset can be
 - An unsigned 12-bit immediate value (i.e. 0 - 4095 bytes).
`LDR r0, [r1, #8]`
 - A register, optionally shifted by an immediate value
`LDR r0, [r1, r2] ; r1 = Memory.word[r1+r2]`
`LDR r0, [r1, r2, LSL#2] ; r1 = Memory.word[r1+4*r2]`
- This can be either added or subtracted from the base register:
 - `LDR r0, [r1, #-8]`
 - `LDR r0, [r1, -r2]`
 - `LDR r0, [r1, -r2, LSL#2]`
- For halfword and signed halfword / byte, offset can be:
 - An unsigned 8 bit immediate value (i.e. 0-255 bytes).
 - A register (unshifted).
- Choice of **pre-indexed** or **post-indexed** addressing



Pre-index

Pre-Index: LDR r1, [r0, #4]

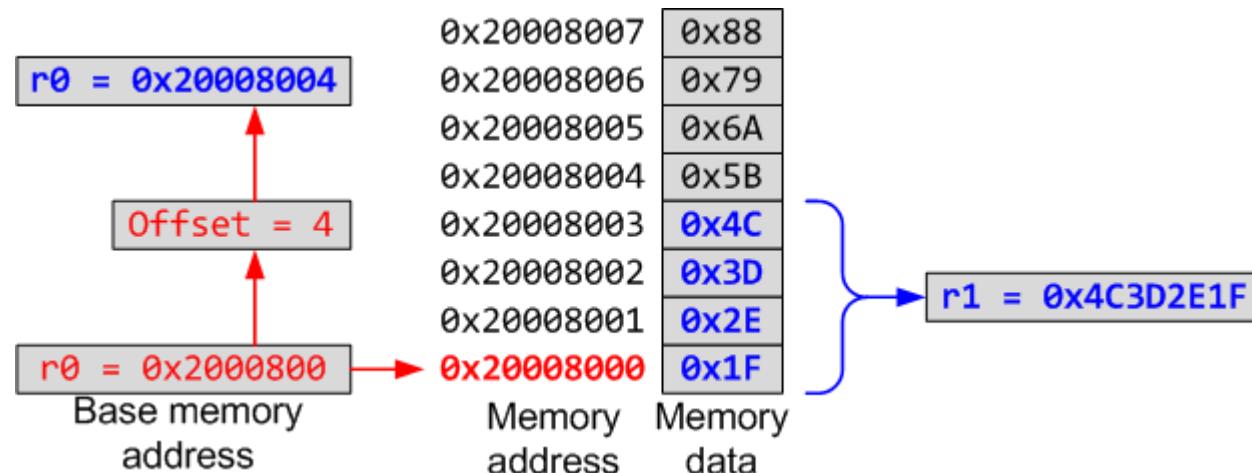


Offset range is -255 to +255



Post-index

Post-Index: LDR r1, [r0], #4



Offset range is -255 to +255



Summary of Pre-index and Post-index

Index Format	Example	Equivalent
Pre-index	LDR r1, [r0, #4]	$r1 \leftarrow \text{memory}[r0 + 4]$, $r0$ is unchanged
Pre-index with update	LDR r1, [r0, #4]!	$r1 \leftarrow \text{memory}[r0 + 4]$ $r0 \leftarrow r0 + 4$
Post-index	LDR r1, [r0], #4	$r1 \leftarrow \text{memory}[r0]$ $r0 \leftarrow r0 + 4$

Offset range is -255 to +255



Load/Store Multiple Registers

STMxx rn{!}, {register_list}

LDMxx rn{!}, {register_list}

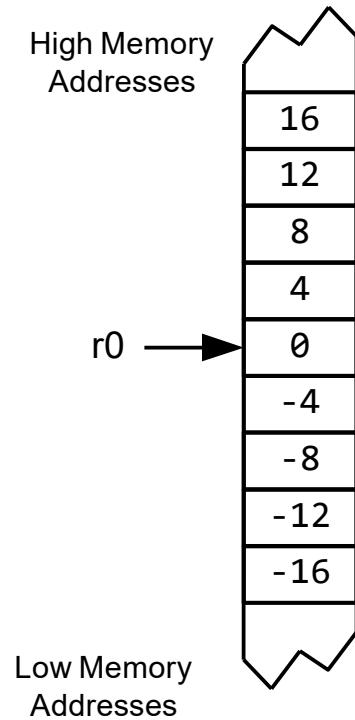
Addressing Modes	Description	Instructions
IA	Increment After	STMIA, LDMIA
IB	Increment Before	STMIB, LDMIB
DA	Decrement After	STMDA, LDMDA
DB	Decrement Before	STMDB, LDMDB

- **IA**: address is incremented by 4 after a word is loaded or stored.
- **IB**: address is incremented by 4 before a word is loaded or stored.
- **DA**: address is decremented by 4 after a word is loaded or stored.
- **DB**: address is decremented by 4 before a word is loaded or stored.



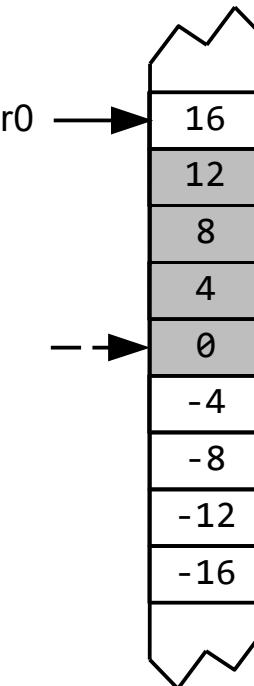
Load Multiple Registers

LDMxx r0!, {r3,r1,r7,r2}



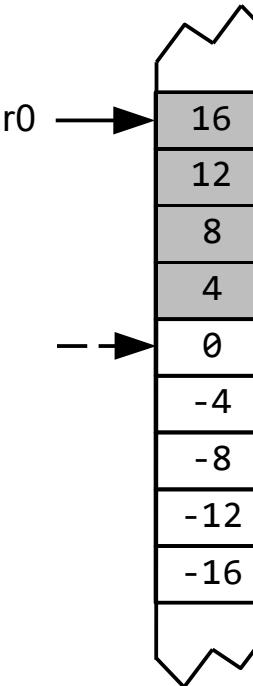
LDMIA

Increment After



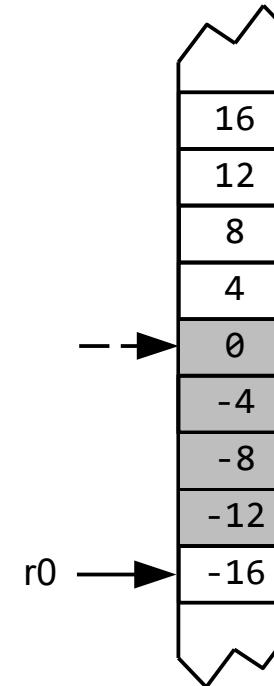
LDMIB

Increment Before



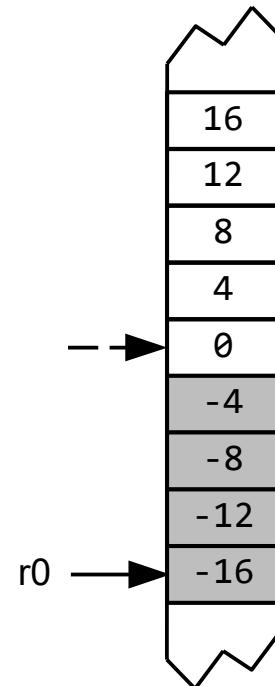
LDMDA

Decrement After



LDMDB

Decrement Before



**r1 = 0
r2 = 4
r3 = 8
r7 = 12**

**r1 = 4
r2 = 8
r3 = 12
r7 = 16**

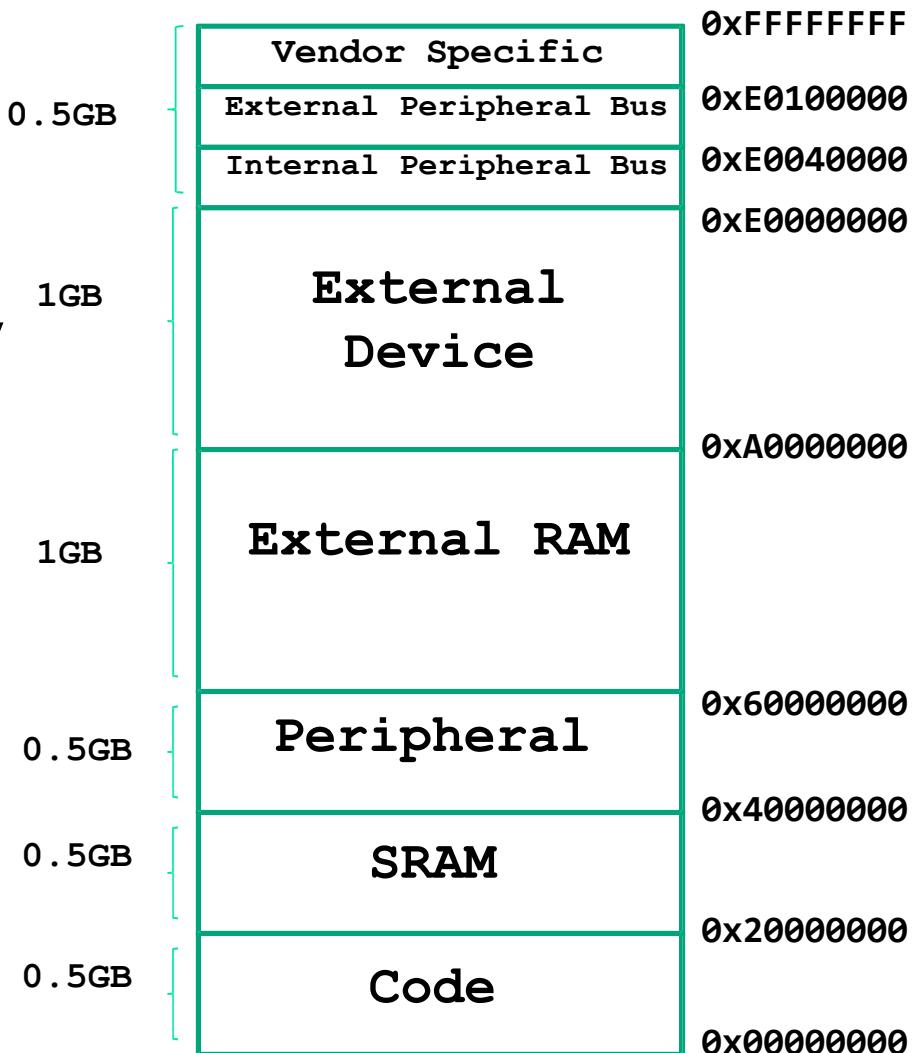
**r1 = -12
r2 = -8
r3 = -4
r7 = -0**

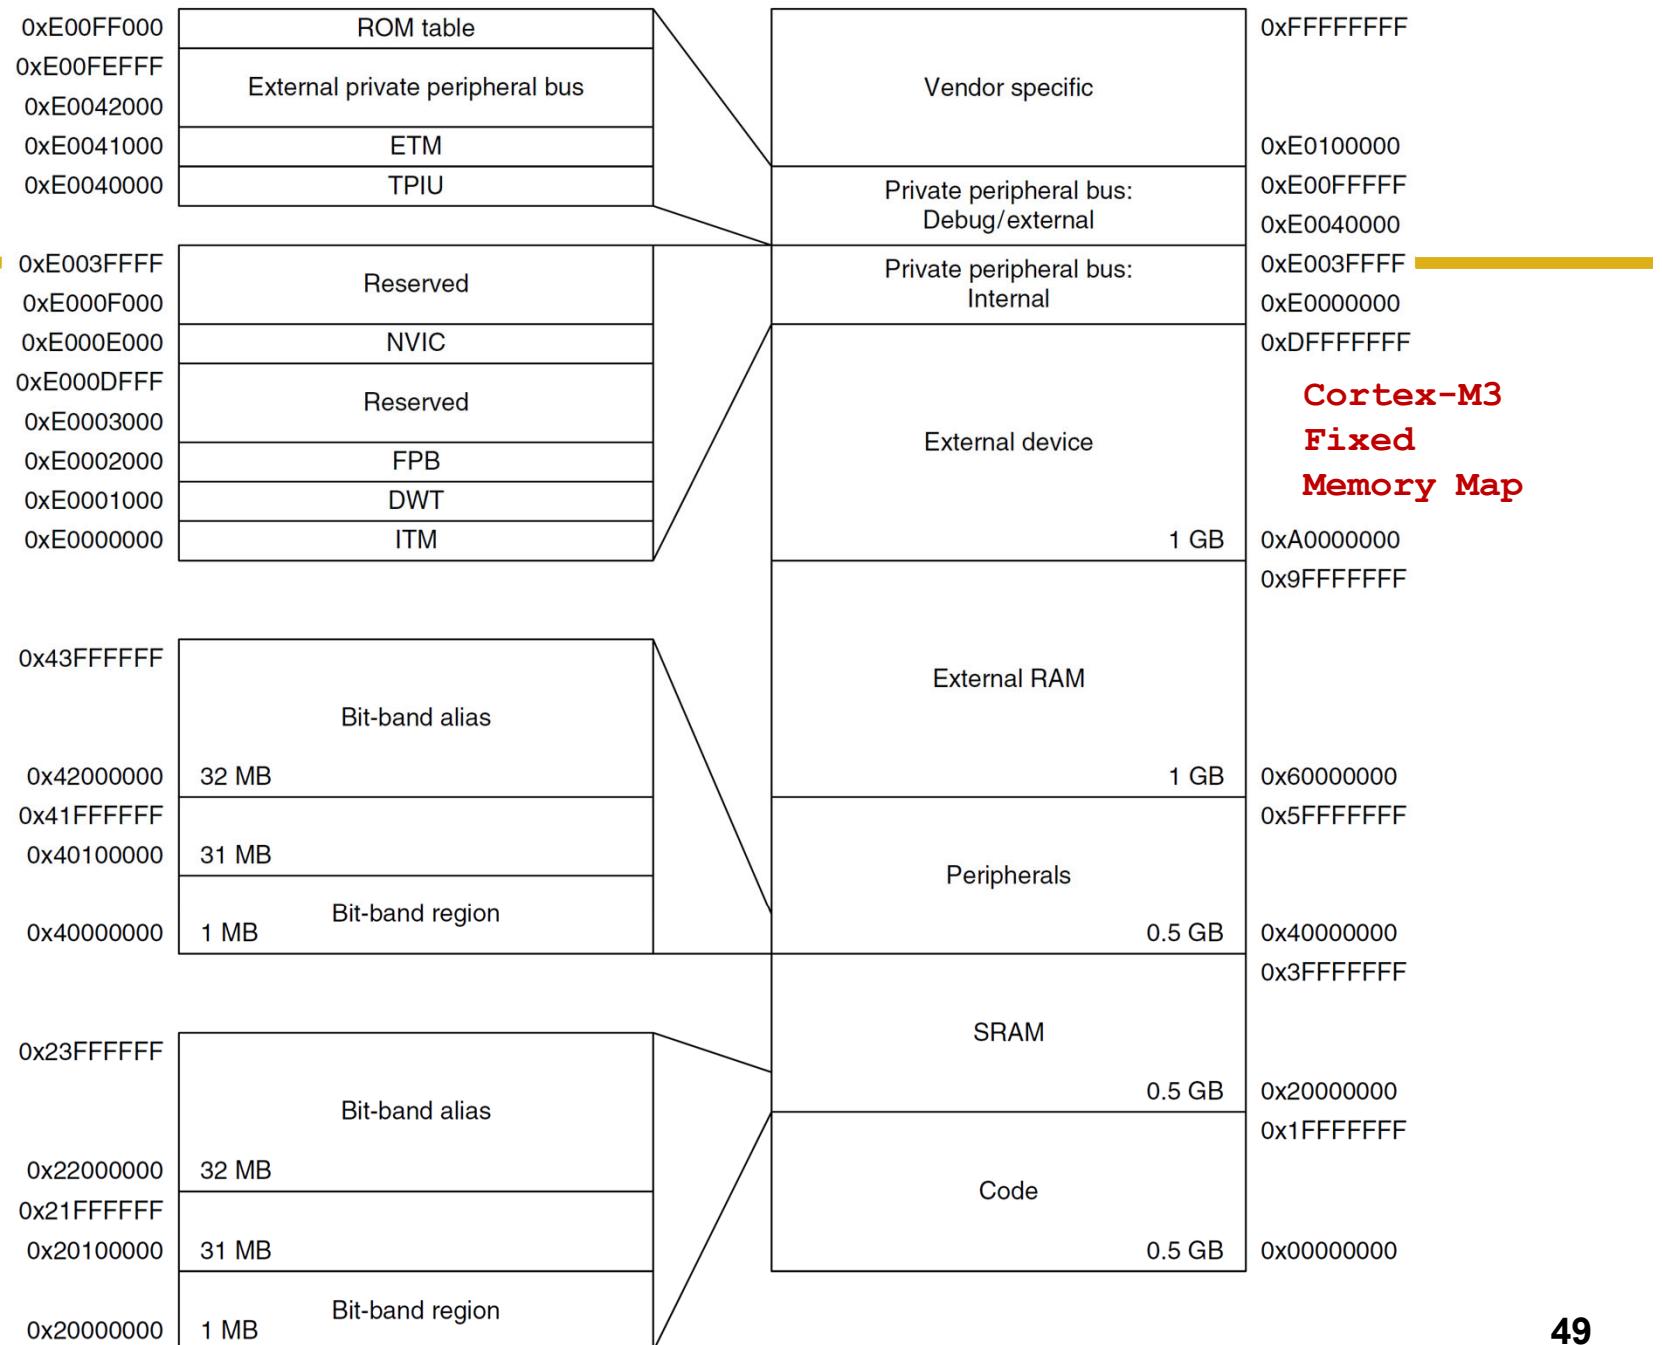
**r1 = -16
r2 = -12
r3 = -8
r7 = -4**

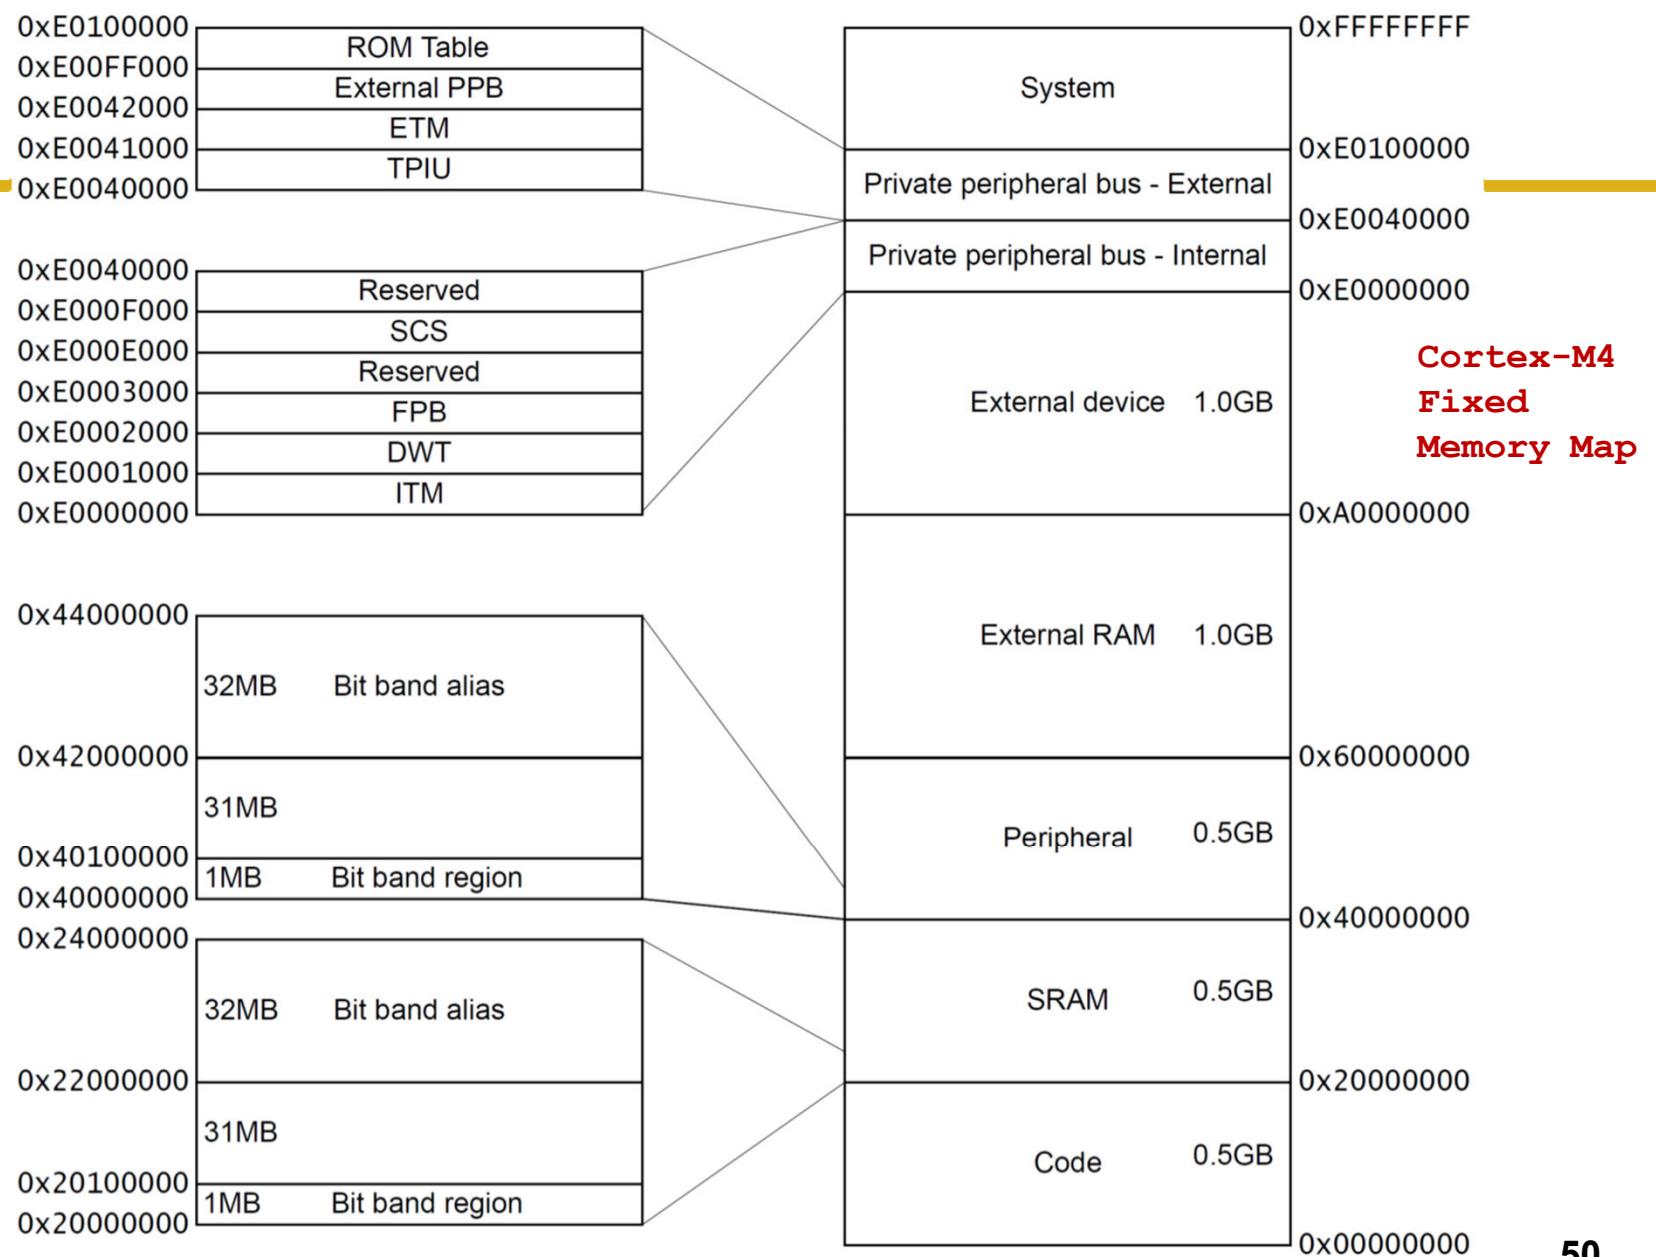


Cortex-M3 & Cortex-M4 Memory Map

- 32-bit Memory Address
- 2^{32} bytes of memory space (4 GB)
- Modified Harvard architecture:
 - Physically separated instruction memory and data memory
 - Separate buses
 - Unified address space









Chapter 6

Flow Control in Assembly



Comparison Instructions

Instruction	Operands	Brief description	Flags
CMP	Rn, Op2	Compare	N,Z,C,V
CMN	Rn, Op2	Compare Negative	N,Z,C,V
TEQ	Rn, Op2	Test Equivalence	N,Z,C
TST	Rn, Op2	Test	N,Z,C

- The only effect of the comparisons is to **update the condition flags**.
 - No need to set S bit.
 - No need to specify Rd.
- Operations are:
 - **CMP** operand1 - operand2, but result not written
 - **CMN** operand1 + operand2, but result not written
 - **TEQ** operand1 ^ operand2, but result not written
 - **TST** operand1 & operand2, but result not written
- Examples:
 - **CMP** r0, r1
 - **TST** r2, #5



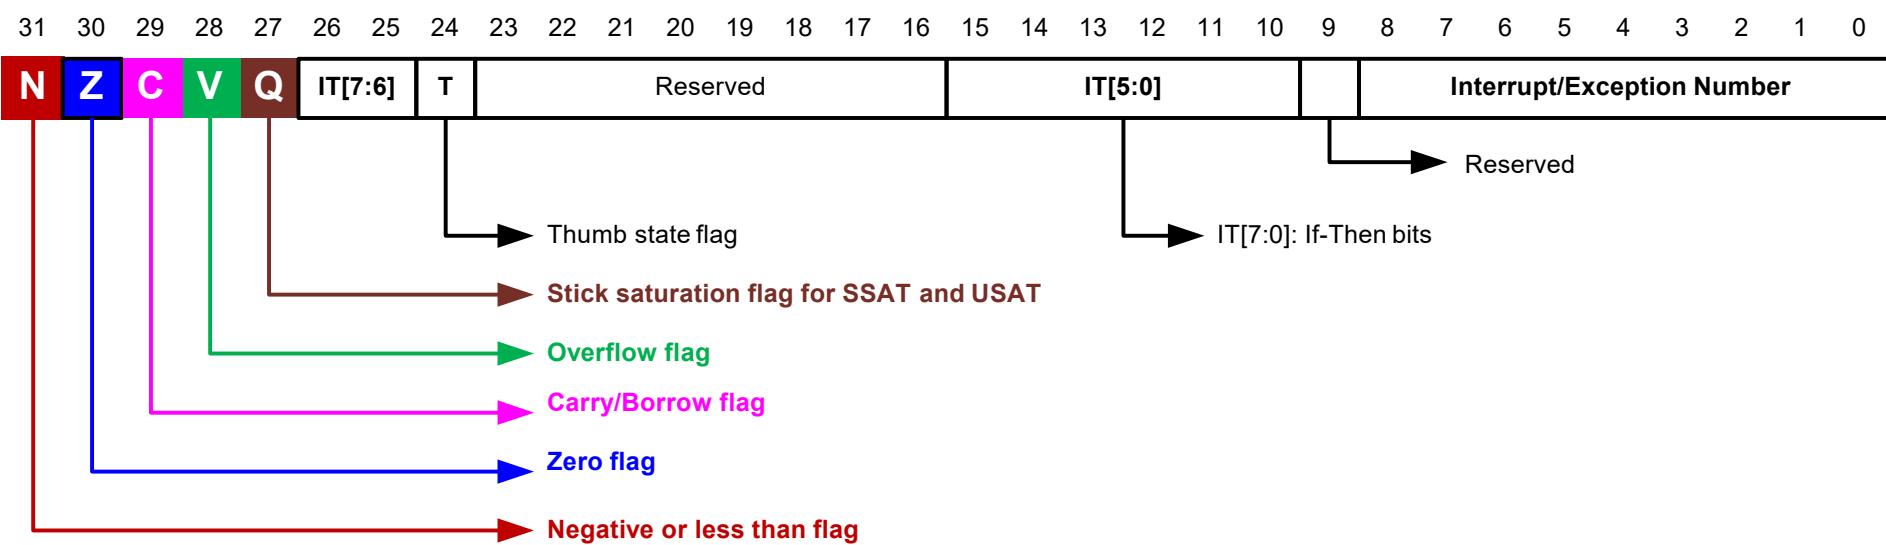
Condition Codes

- Status of a condition code depends on the status of N, Z, C, and V flags.

Suffix	Description	Flags tested
EQ	Equal	Z=1
NE	Not equal	Z=0
CS/HS	Unsigned higher or same	C=1
CC/LO	Unsigned lower	C=0
MI	Negative	N=1
PL	Positive or Zero	N=0
VS	Overflow	V=1
VC	No overflow	V=0
HI	Unsigned higher	C=1 & Z=0
LS	Unsigned lower or same	C=0 or Z=1
GE	Signed Greater or equal	N=V
LT	Signed Less than	N!=V
GT	Signed Greater than	Z=0 & N=V
LE	Signed Less than or equal	Z=1 or N!=V
AL	Always	

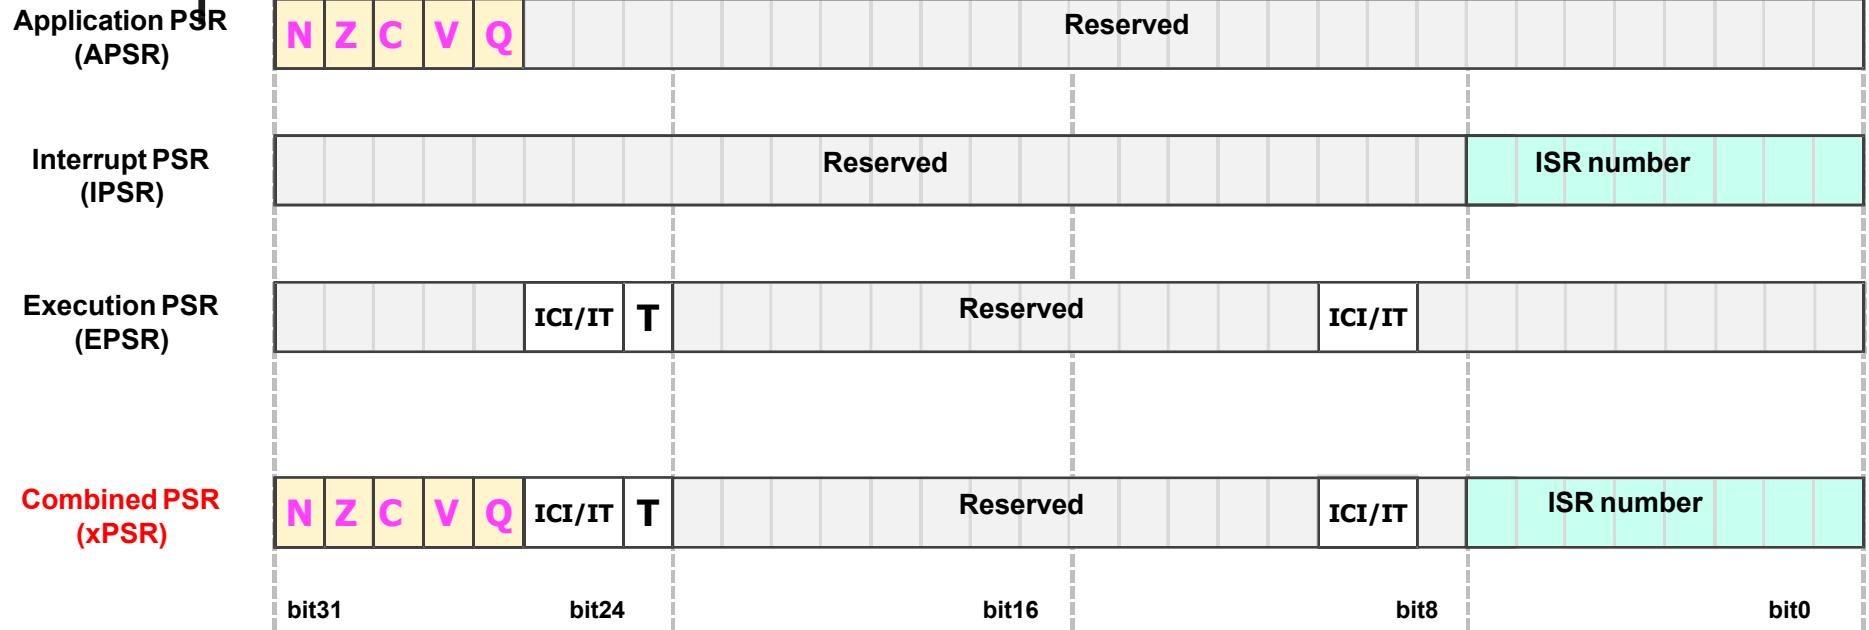


Combined Program Status Registers (xPSR)





Combined Program Status Registers (xPSR)



$$XPSR = APSR + IPSR + EPSR$$



Example: GE - Signed Greater or Equal (N == V)

CMP r0, r1

We in fact perform subtraction $r0 - r1$, without saving the result.

	N = 0	N = 1
V = 0	<ul style="list-style-type: none">• No overflow, implying the result is correct.• The result is non-negative,• Thus $r0 - r1 \geq 0$, i.e., $r0 \geq r1$	<ul style="list-style-type: none">• No overflow, implying the result is correct.• The result is negative.• Thus $r0 - r1 < 0$, i.e., $r0 < r1$
V = 1	<ul style="list-style-type: none">• Overflow occurs, implying the result is incorrect.• The result is mistakenly reported as non-negative and in fact it should be negative.• Thus $r0 - r1 < 0$ in reality, i.e., $r0 < r1$	<ul style="list-style-type: none">• Overflow occurs, implying the result is incorrect.• The result is mistakenly reported as negative and in fact it should be non-negative.• Thus $r0 - r1 \geq 0$ in reality., i.e. $r0 \geq r1$

Conclusions:

- If $N == V$, then it is signed greater or equal (GE).
- Otherwise, it is signed less than (LT)



Branch Instructions

Instruction	Operands	Brief description	Flags
B	label	Branch	-
BL	label	Branch with Link	-
BLX	Rm	Branch indirect with Link	-
BX	Rm	Branch indirect	-

- **B label:** causes a branch to label.
- **BL label:** instruction copies the address of the next instruction into **r14** (**lr**, the link register), and causes a branch to label.
- **BX Rm:** branch to the address held in Rm
- **BLX Rm:** copies the address of the next instruction into **r14** (**lr**, the link register) and branch to the address held in Rm



Branch With Link

- The "Branch with link (BL)" instruction implements a subroutine call by writing $\text{PC} - 4$ into the LR of the current bank.
 - i.e. the address of the next instruction following the branch with link instruction (allowing for the pipeline).
- To return from subroutine, simply need to restore the PC from the LR:
 - `MOV pc, lr`
 - Again, pipeline has to refill before execution continues.
- The "Branch" instruction does not affect LR.



Signed vs. Unsigned

Conditional codes applied to branch instructions

Compare	Signed	Unsigned
==	EQ	EQ
≠	NE	NE
>	GT	HI
≥	GE	HS
<	LT	LO
≤	LE	LS



Compare	Signed	Unsigned
==	BEQ	BEQ
≠	BNE	BNE
>	BGT	BHI
≥	BGE	BHS
<	BLT	BLO
≤	BLE	BLS



Branch Instructions

	Instruction	Description	Flags tested
Unconditional Branch	B label	Branch to label	
Conditional Branch	BEQ label	Branch if EQual	Z = 1
	BNE label	Branch if Not Equal	Z = 0
	BCS/BHS label	Branch if unsigned Higher or Same	C = 1
	BCC/BLO label	Branch if unsigned LOwer	C = 0
	BMI label	Branch if MInus (Negative)	N = 1
	BPL label	Branch if PLus (Positive or Zero)	N = 0
	BVS label	Branch if oVerflow Set	V = 1
	BVC label	Branch if oVerflow Clear	V = 0
	BHI label	Branch if unsigned HIgher	C = 1 & Z = 0
	BLS label	Branch if unsigned Lower or Same	C = 0 or Z = 1
	BGE label	Branch if signed Greater or Equal	N = V
	BLT label	Branch if signed Less Than	N != V
	BGT label	Branch if signed Greater Than	Z = 0 & N = V
	BLE label	Branch if signed Less than or Equal	Z = 1 or N = !V



Which is Greater: 0xFFFFFFFF or 0x00000001?

It's **software's responsibility** to tell computer how to interpret data:

- If written in C, declare the signed vs unsigned variable
- If written in Assembly, use signed vs unsigned branch instructions

```
signed int x, y ;  
x = -1;  
y = 1;  
if (x > y)  
    ...
```

```
MOVS r6, #0xFFFFFFFF ;x  
MOVS r5, #0x00000001 ;y  
CMP r5, r6  
BLE Then_Clause  
...
```

Check
Z=1 or N!=V

BLE: Branch if less than or equal, signed \leq

```
unsigned int x, y ;  
x = 4294967295;  
y = 1;  
if (x > y)  
    ...
```

```
MOVS r6, #0xFFFFFFFF ;x  
MOVS r5, #0x00000001 ;y  
CMP r5, r6  
BLS Then_Clause  
...
```

Check
C=0 or Z=1

BLS: Branch if lower or same, unsigned \leq



Combination

Instruction	Operands	Brief description	Flags
CBZ	Rn, label	Compare and Branch if Zero	-
CBNZ	Rn, label	Compare and Branch if Non Zero	-

- Except that it does not change the condition code flags, **CBZ Rn, label** is equivalent to:
CMP Rn, #0
BEQ label
- Except that it does not change the condition code flags, **CBNZ Rn, label** is equivalent to:
CMP Rn, #0
BNE label



Conditional Execution

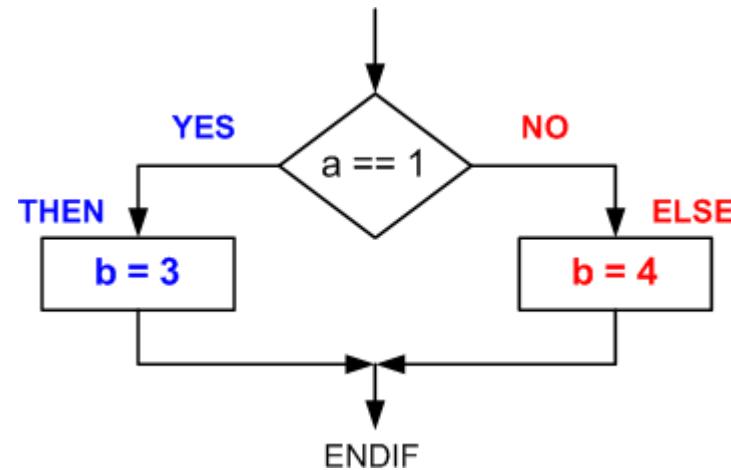
Add instruction	Condition	Flag tested
ADDEQ r3, r2, r1	Add if EQual	Add if Z = 1
ADDNE r3, r2, r1	Add if Not Equal	Add if Z = 0
ADDHS r3, r2, r1	Add if Unsigned Higher or Same	Add if C = 1
ADDLO r3, r2, r1	Add if Unsigned LOwer	Add if C = 0
ADDMI r3, r2, r1	Add if Minus (Negative)	Add if N = 1
ADDPL r3, r2, r1	Add if PLus (Positive or Zero)	Add if N = 0
ADDVS r3, r2, r1	Add if oVerflow Set	Add if V = 1
ADDVC r3, r2, r1	Add if oVerflow Clear	Add if V = 0
ADDHI r3, r2, r1	Add if Unsigned HIgher	Add if C = 1 & Z = 0
ADDLS r3, r2, r1	Add if Unsigned Lower or Same	Add if C = 0 or Z = 1
ADDGE r3, r2, r1	Add if Signed Greater or Equal	Add if N = V
ADDLT r3, r2, r1	Add if Signed Less Than	Add if N != V
ADDGT r3, r2, r1	Add if Signed Greater Than	Add if Z = 0 & N = V
ADDLE r3, r2, r1	Add if Signed Less than or Equal	Add if Z = 1 or N = !V



If-then-else

C Program

```
if (a == 1)
    b = 3;
else
    b = 4;
```



```
; r1 = a, r2 = b
CMP r1, #1      ; compare a and 1
BNE else        ; go to else if a ≠ 1
then  MOV r2, #3  ; b = 3
      B endif     ; go to endif
else  MOV r2, #4  ; b = 4
endif
```

BNE: Branch not equal, Z=0



For Loop

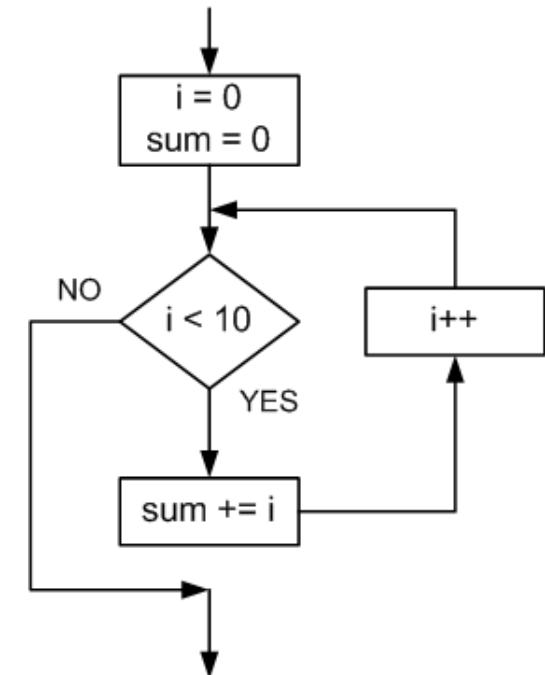
C Program

```
int i;
int sum = 0;
for(i = 0; i < 10; i++){
    sum += i;
}
```

Implementation 1:

```
        MOV r0, #0 ; i
        MOV r1, #0 ; sum

        B check
loop    ADD r1, r1, r0
        ADD r0, r0, #1
check   CMP r0, #10
        BLT loop
endloop
```





Example 1: Greatest Common Divider (GCD)

Euclid's Algorithm

```
uint32_t a, b;  
while (a != b) {  
    if (a > b)  
        a = a - b;  
    else  
        b = b - a;  
}
```



```
gcd  CMP r0, r1  
      SUBHI r0, r0, r1  
      SUBL0 r1, r1, r0  
      BNE gcd
```



```
; suppose r0 = a and r1 = b  
gcd   CMP r0, r1          ; a > b?  
      BEQ end              ; if a = b, done  
      BLO less              ; a < b  
      SUB r0, r0, r1         ; a = a - b  
      B  gcd  
less   SUB r1, r1, r0       ; b = b - a  
      B  gcd  
end
```



Break and Continue

C Program	Assembly Program
<pre>int i; int sum = 0; for(i = 0; i < 10; i++) { if (i == 5) continue; sum += i; }</pre>	<pre>MOVS r0, #0 ; i = 0 MOVS r1, #0 ; sum = 0 loop CMP r0, #10 BGE endloop CMP r0, #5 ADDNE r1, r1, r0 ; sum += i ADD r0, r0, #1 ; i++ B loop endloop</pre>



Switch Statement Example

C Program	Assembly Program
<pre>unsigned int score; char grade; switch(score){ case 10: break; case 9: grade = 'A'; break; case 8: grade = 'B'; break; case 7: grade = 'C'; break; case 6: grade = 'D'; break; default: grade = 'F'; break; }</pre>	<p>; $r0 = \text{numeric score } (0 \leq r0 \leq 10)$; $r1 = \text{letter grade}$</p> <p>CMP r0, #5 BLS default ; branch if $\text{unsigned } r0 \leq 5$ SUBS r2, r0, #6 ; $r2$ holds branch index</p> <p>; $pc = pc + 4 + 2 \times \text{BranchTable}[r2]$ TBB [pc, r2]</p> <p>BranchTable</p> <p>DCB (case_6 - BranchTable)/2 ; index = 0 DCB (case_7 - BranchTable)/2 ; index = 1 DCB (case_8 - BranchTable)/2 ; index = 2 DCB (case_10_9 - BranchTable)/2 ; index = 3 DCB (case_10_9 - BranchTable)/2 ; index = 4</p> <p>ALIGN</p> <p>case_10_9</p> <p>MOV r1, #0x41 ; ASCII 'A' = 0x41 B exit</p> <p>case_8</p> <p>MOV r1, #0x42 ; ASCII 'B' = 0x42 B exit</p>



Switch Statement Example Cont'd

C Program	Assembly Program
<pre>unsigned int score; char grade; switch(score){ case 10: break; case 9: grade = 'A'; break; case 8: grade = 'B'; break; case 7: grade = 'C'; break; case 6: grade = 'D'; break; default: grade = 'F'; break; }</pre>	<pre>case_7 MOV r1, #0x43 ; ASCII 'C' = 0x43 B exit case_6 MOV r1, #0x44 ; ASCII 'D' = 0x44 B exit default MOV r1, #0x46 ; ASCII 'F' = 0x46 B exit</pre>

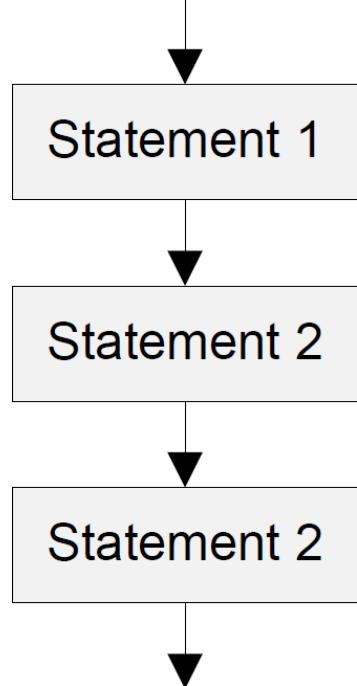


Chapter 7

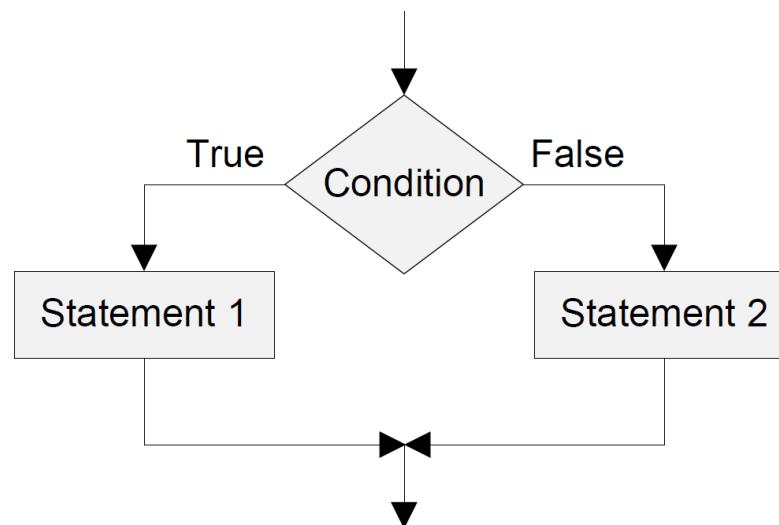
Structured Programming



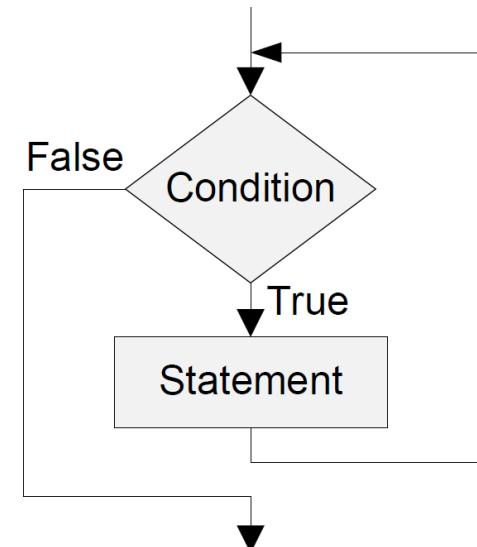
Basic Control Structures



Sequence Structure



Selection Structure

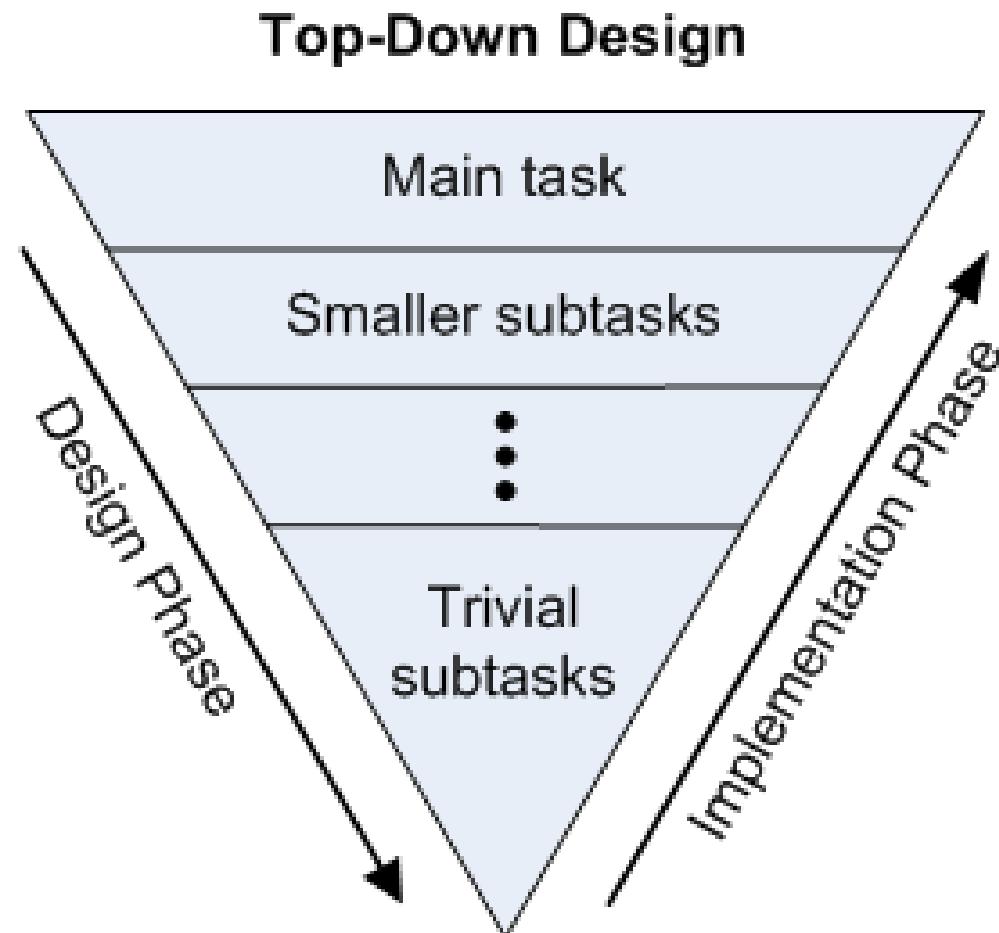


Loop Structure

- Each structure has only 1 entry point and 1 exit point.
- Control structure can be nested or embedded in another structure to form a compound structure.



Top-Down Design





Structured Programming

- Programming using nested control structures as well as subroutines in high level languages is called **structured programming**.

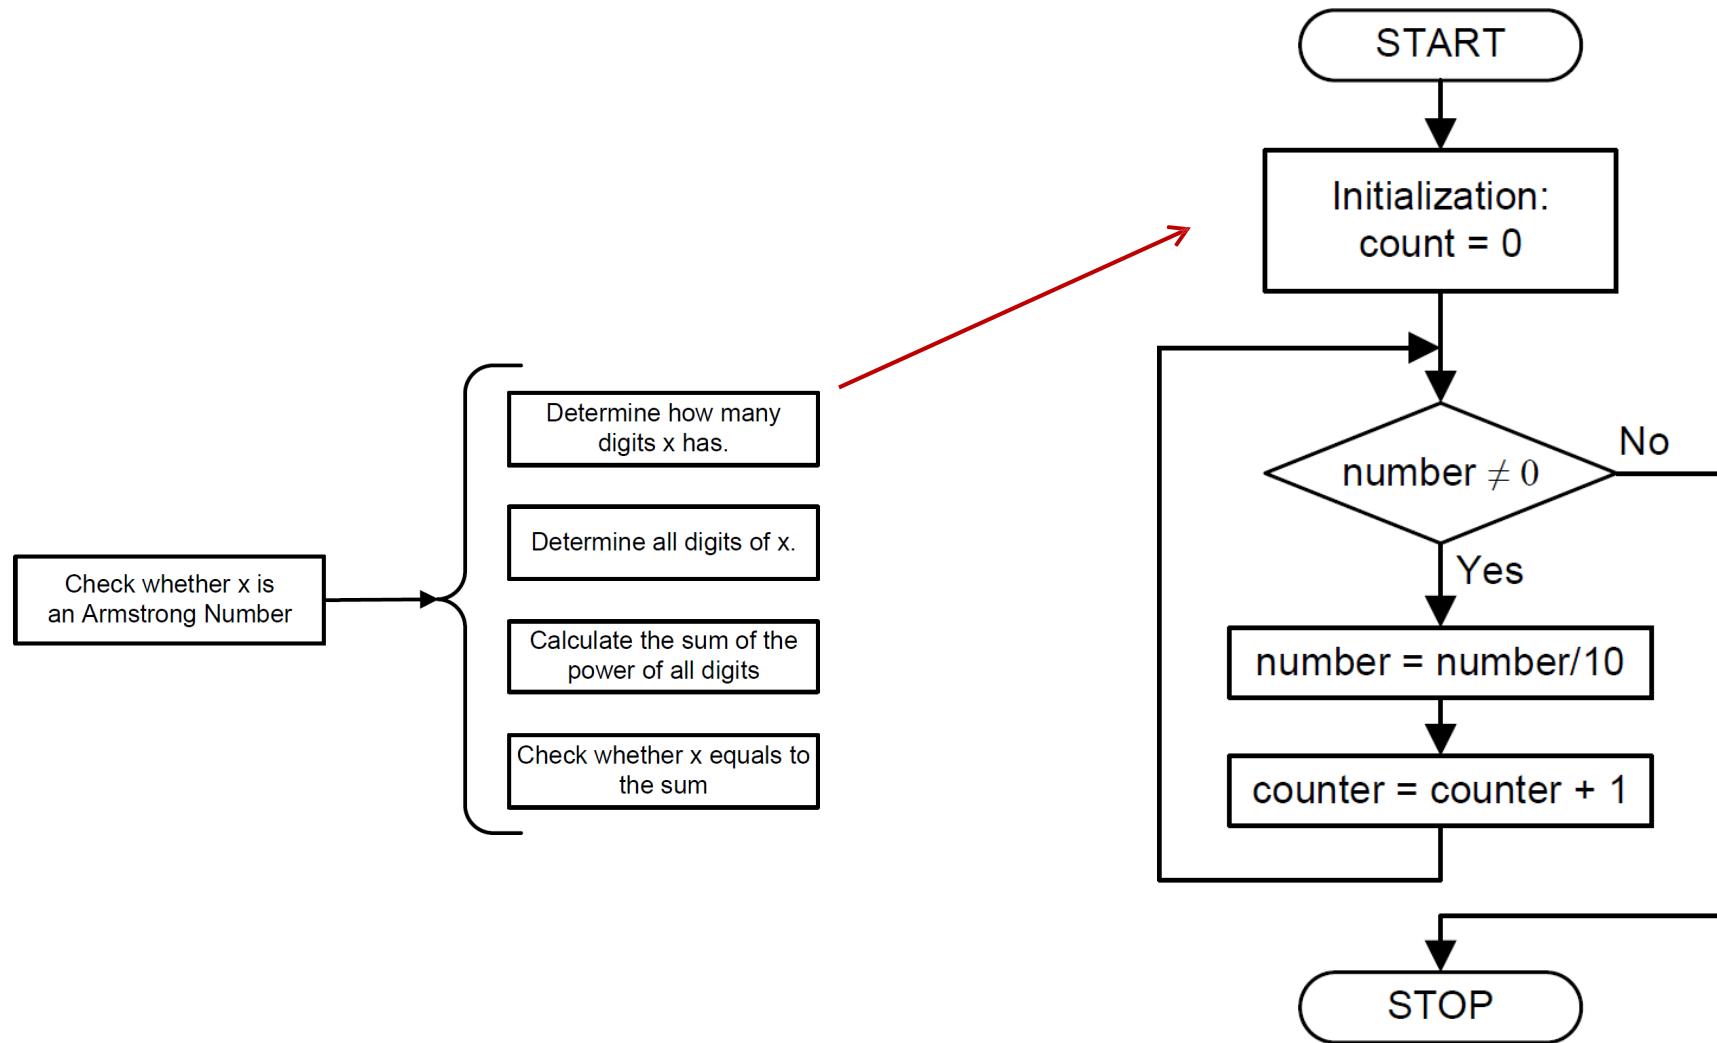


Structured Programming in Assembly

- Structured programming in assembly involves 2 steps:
 1. Top-down Logical Design – Perform stepwise refinement and construct program flow using high level structures.
 2. Implement the high level structures identified in Step 1 in assembly language.



Top-Down Design Example: Counting digits





Top-Down Design Example: Program Implementation

C Program	Assembly Program
int main(void) { int number, sum, r; int flag, t; number = 371; sum = 0; t = number; while(t != 0) { r = t % 10; sum = sum + r*r*r; t = t / 10; } if(number == sum) flag = 1; else flag = 0; while(1); }	AREA Armstrong, CODE, READONLY EXPORT __main ENTRY __main PROC LDR r0, =371 ; number to be checked MOV r4, #0 ; sum = 0 MOV r1, r0 ; save a copy, r1=number loop: MOV r3, #10 CBZ r1, check ; if t = 0, exit Loop SDIV r6, r1, r3 ; r2 = remainder MLS r2, r3, r6, r1 ; r2 = r1 - 10*r6 MUL r3, r2, r2 ; remainder^2 MLA r4, r2, r2, r4 ; r4 = r2*r2 + r4 MOVS r3, #10 SDIV r1, r1, r3 ; t = t/10 CBNZ r1, loop check: CMP r0, r4 yes: MOVEQ r0, #1 ; Armstrong no: MOVNE r0, #0 ; not Armstrong stop: B stop ENDP END



Register Reuse

```
AREA myCode, CODE
EXPORT __main
ENTRY
__main PROC

    LDR r2, =B      } Lifetime of r2
    LDR r3, [r2]    }
    LDR r4, =C      } Lifetime of r4
    LDR r5, [r4]    }
    LDR r6, =D      } Lifetime of r6
    LDR r7, [r6]    }
    ADD r1, r3, r5
    SUB r1, r1, r7
    LDR r0, =A      } Lifetime of r0
    STR r1, [r0]    }

ENDP
```

```
AREA myData, DATA
A DCD 0
B DCD -1
C DCD -2
D DCD 2

END
```

8 registers used

```
AREA myCode, CODE
EXPORT __main
ENTRY
__main PROC

    LDR r2, =B      } Lifetime of r2
    LDR r3, [r2]    }
    LDR r2, =C      } Lifetime of r2
    LDR r5, [r2]    }
    LDR r2, =D      } Lifetime of r2
    LDR r7, [r2]    }
    ADD r3, r3, r5
    SUB r3, r3, r7
    LDR r2, =A      } Lifetime of r2
    STR r3, [r2]    }

ENDP
```

```
AREA myData, DATA
A DCD 0
B DCD -1
C DCD -2
D DCD 2

END
```

4 registers used

```
AREA myCode, CODE
EXPORT __main
ENTRY
__main PROC

    LDR r2, =B
    LDR r3, [r2]
    LDR r2, =C
    LDR r5, [r2]
    LDR r2, =D
    LDR r2, [r2] } Reuse r2
    ADD r3, r3, r5
    SUB r3, r3, r2
    LDR r2, =A
    STR r3, [r2]

ENDP
```

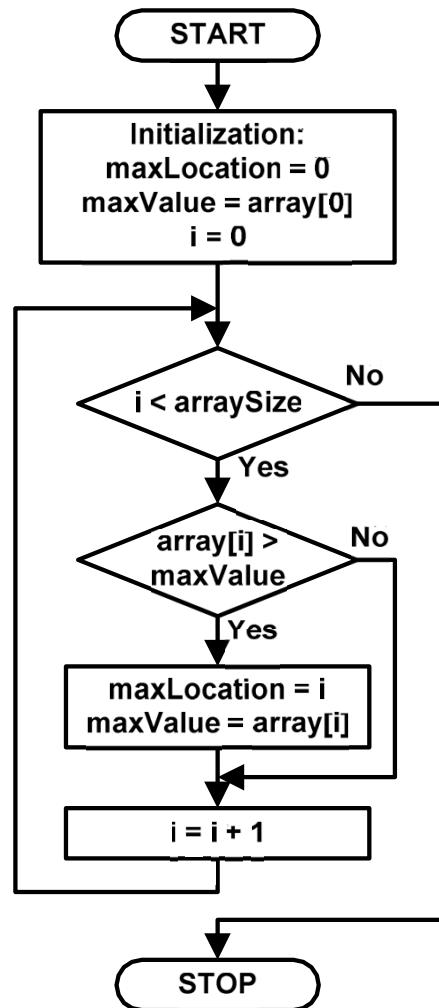
```
AREA myData, DATA
A DCD 0
B DCD -1
C DCD -2
D DCD 2

END
```

3 registers used



Finding Max of an Array





Finding Max of an Array

C Program	Assembly Program
int array[10] = {-1, 5, 3, 8, 10, 23, 6, 5, 2, -10};	AREA myData, DATA ALIGN array DCD -1,5,3,8,10,23,6,5,2,-10 size DCD 10
int size = 10;	
int main(void) {	AREA findMax, CODE EXPORT __main ALIGN ENTRY PROC
int i, maxLocation, maxValue;	; Identify the array size LDR r3, =size LDR r3, [r3] ; array size SUB r3, r3, #1
// Initialize max and location	; Initialize max value and location LDR r4, =array LDR r0, [r4] ; r0 = default max MOV r1, #0 ; r1 = max location
maxLocation = 0;	
maxValue = array[0];	
// Loop through the array	; Loop over the array
for (i = 0; i < size; i++){	MOV r2, #0 ; Loop index i
if (array[i] > maxValue) {	CMP r2, r3 ; compare i & size
maxValue = array[i];	BGE stop ; stop if i ≥ size
maxLocation = i;	LDR r5, [r4,r2,LSL #2] ; array[i]
}	CMP r5, r0 ; compare with max
}	MOVGT r0, r5 ; update max value
	MOVGT r1, r2 ; update location
	ADD r2, r2, #1 ; update index i
	B loop
while(1); //dead Loop	stop B stop ; dead Loop
}	ENDP END



Chapter 8 Subroutines



Introduction

- 2 special considerations when using subroutines in assembly
 - 1. Preserve and recover the caller's environment.
 - Subroutine must push registers at beginning of subroutine.
 - Subroutine must pop registers at end of subroutine.
 - 2. Embedded application binary interface (EABI) has specifications for passing input arguments to a subroutine and returning the result back to the caller.
 - Follow EABI to make assembly subroutines compatible with C programs.



Calling a Subroutine

- Assembly program can use branch and link (BL) instruction to call a subroutine.
- BL instruction performs 2 operations:
 1. Store memory address of instruction after BL instruction in LR ($PC + 4$) → referred to as **return address**
 2. Place memory address of 1st instruction of subroutine in PC.



BL and BX

```
void enable(void) ;  
  
...  
enable() ;  
...
```

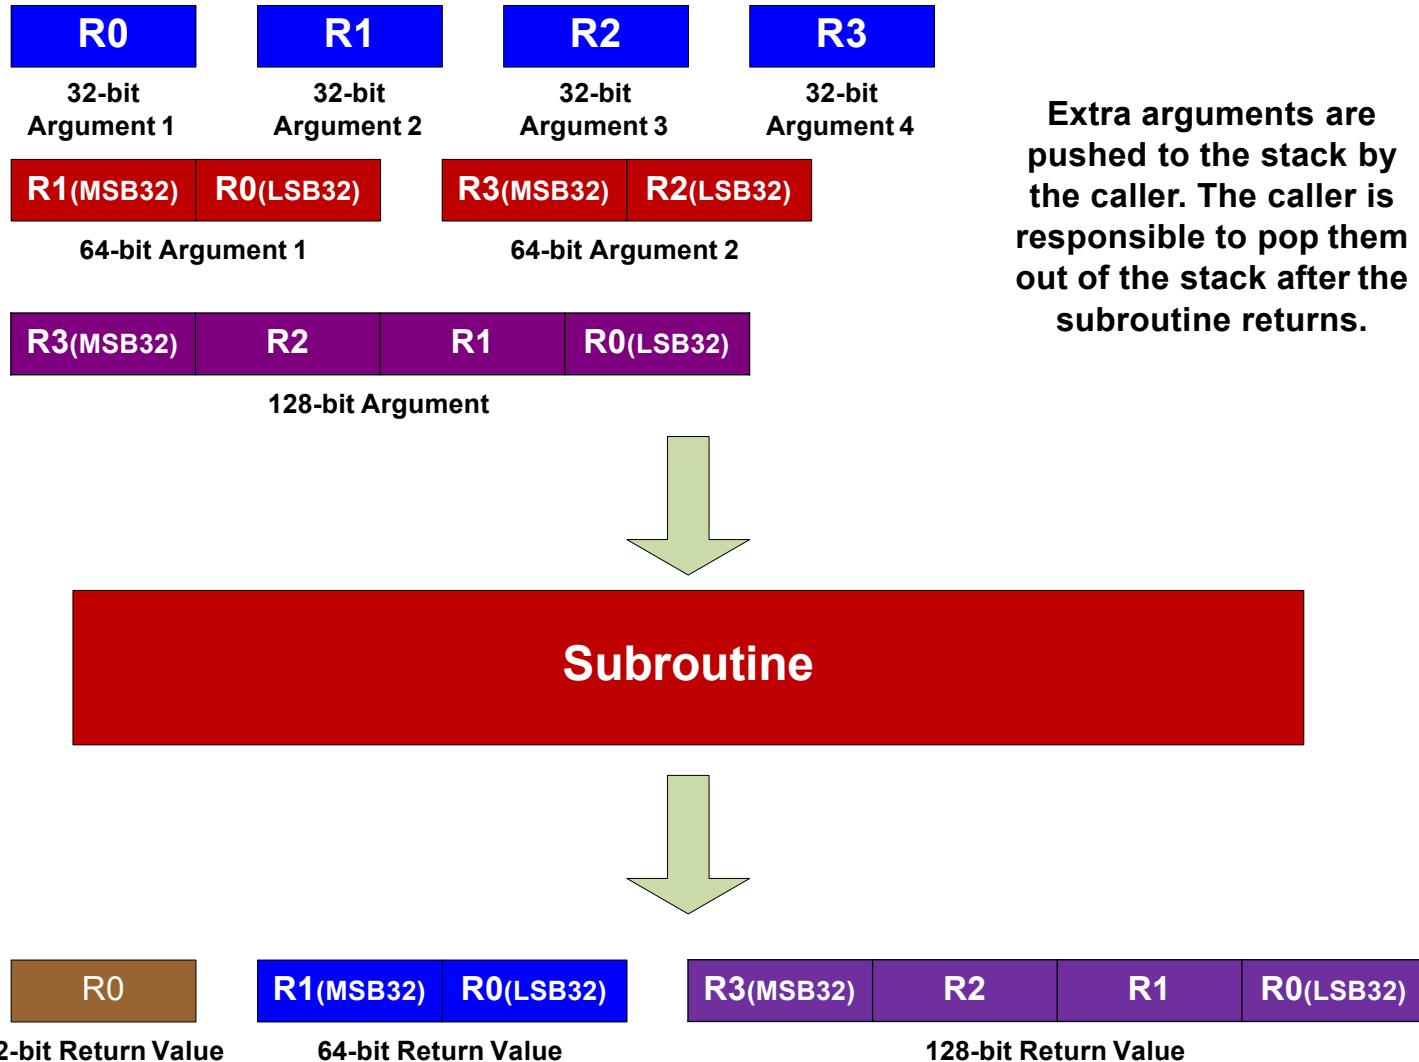


```
...  
BL enable  
...
```

```
export enable  
enable ...  
...  
BX LR
```



Passing Arguments via Registers



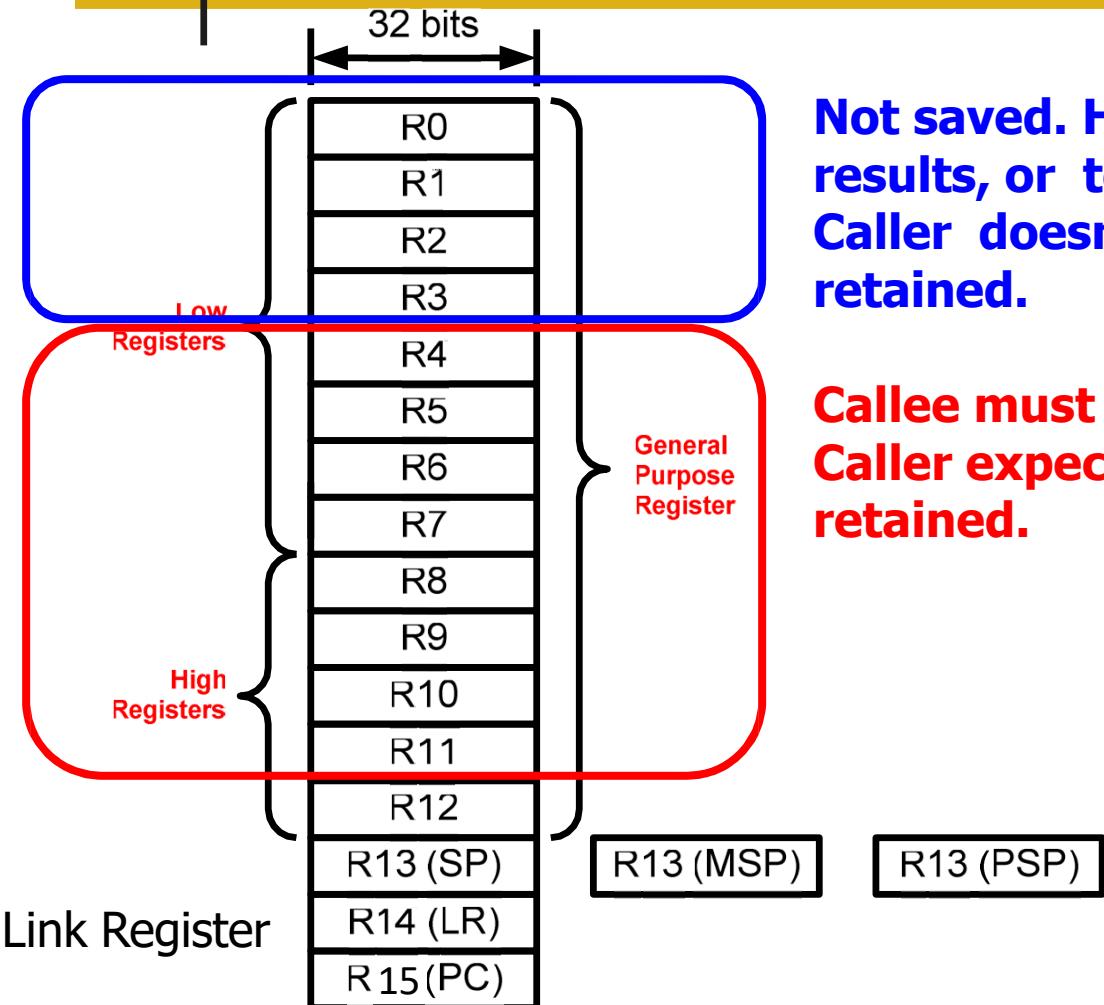


ARM Procedure Call Standard

Register	Usage	Subroutine Preserved	Notes
r0	Argument 1 and return value	No	If return has 64 bits, then r0:r1 hold it. If argument 1 has 64 bits, r0:r1 hold it.
r1	Argument 2	No	
r2	Argument 3	No	If the return has 128 bits, r0-r3 hold it.
r3	Argument 4	No	If more than 4 arguments, use the stack
r4	General-purpose V1	Yes	Variable register 1 holds a local variable.
r5	General-purpose V2	Yes	Variable register 2 holds a local variable.
r6	General-purpose V3	Yes	Variable register 3 holds a local variable.
r7	General-purpose V4	Yes	Variable register 4 holds a local variable.
r8	General-purpose V5	YES	Variable register 5 holds a local variable.
r9	Platform specific/V6	No	Usage is platform-dependent.
r10	General-purpose V7	Yes	Variable register 7 holds a local variable.
r11	General-purpose V8	Yes	Variable register 8 holds a local variable.
r12 (IP)	Intra-procedure-call register	No	It holds intermediate values between a procedure and the sub-procedure it calls.
r13 (SP)	Stack pointer	Yes	SP has to be the same after a subroutine has completed.
r14 (LR)	Link register	No	LR does not have to contain the same value after a subroutine has completed.
r15 (PC)	Program counter	N/A	Do not directly change PC

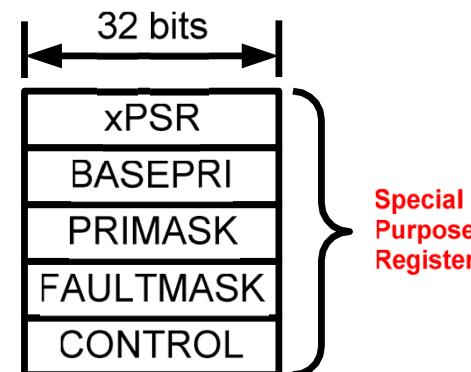


Link Register



Not saved. Hold arguments, results, or temporary values. Caller doesn't expect them to be retained.

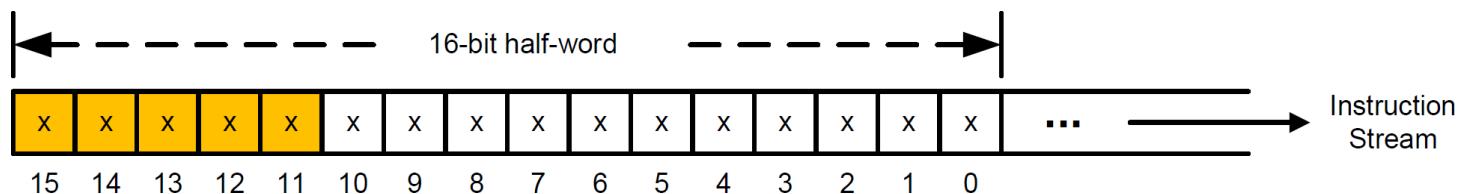
Callee must save them. Caller expects these values are retained.





Realities

- PC is always incremented by **4**.
 - Each time, 4 bytes are fetched from the instruction memory
 - It is either two 16-bit instructions or one 32-bit instruction



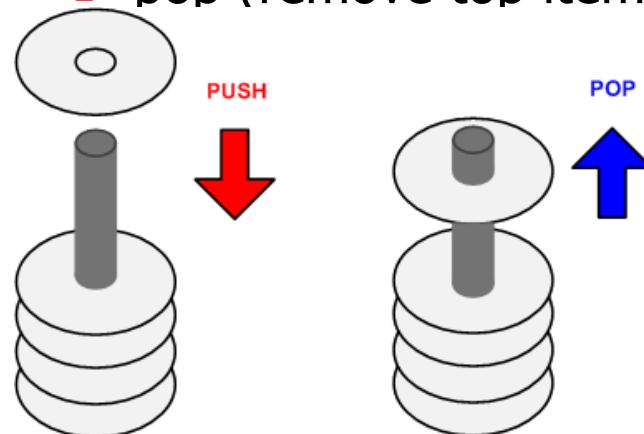
If bit [15-11] = **11101**, **11110**, or **11111**, then, it is the first half-word of a 32-bit instruction. Otherwise, it is a 16-bit instruction.

- The least significant bit of LR is always **1** for ARM Cortex-M
 - This bit is used to control the processor mode:
 - 0 = ARM, 1 = THUMB
 - Cortex-M only supports THUMB.



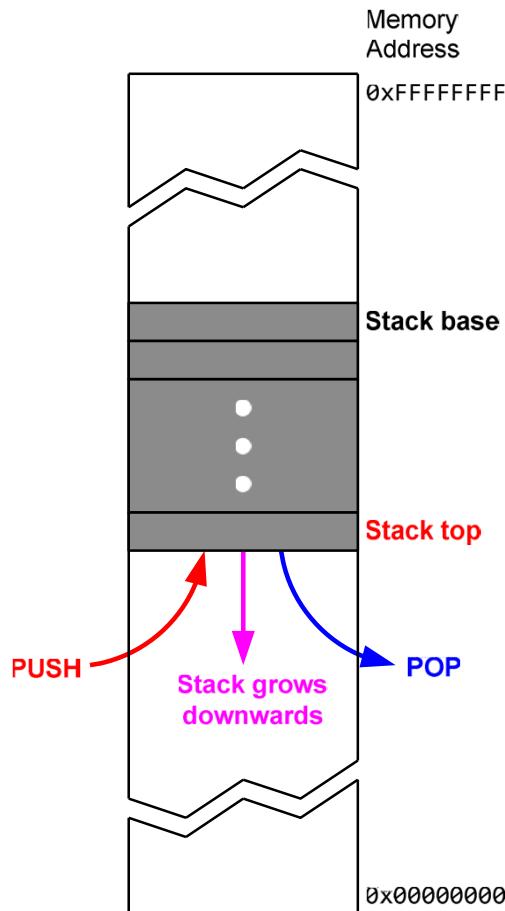
Stack

- A **Last-In-First-Out** data structure
- Only allow to access the most recently added item
 - Also called the top of the stack
- Key operations:
 - push (add item to stack)
 - pop (remove top item from stack)

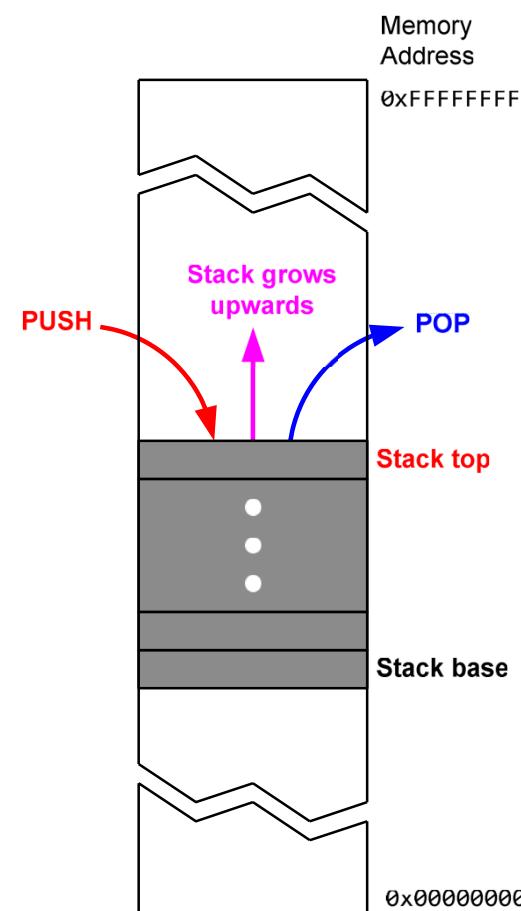




Stack Growth Convention: Ascending vs Descending



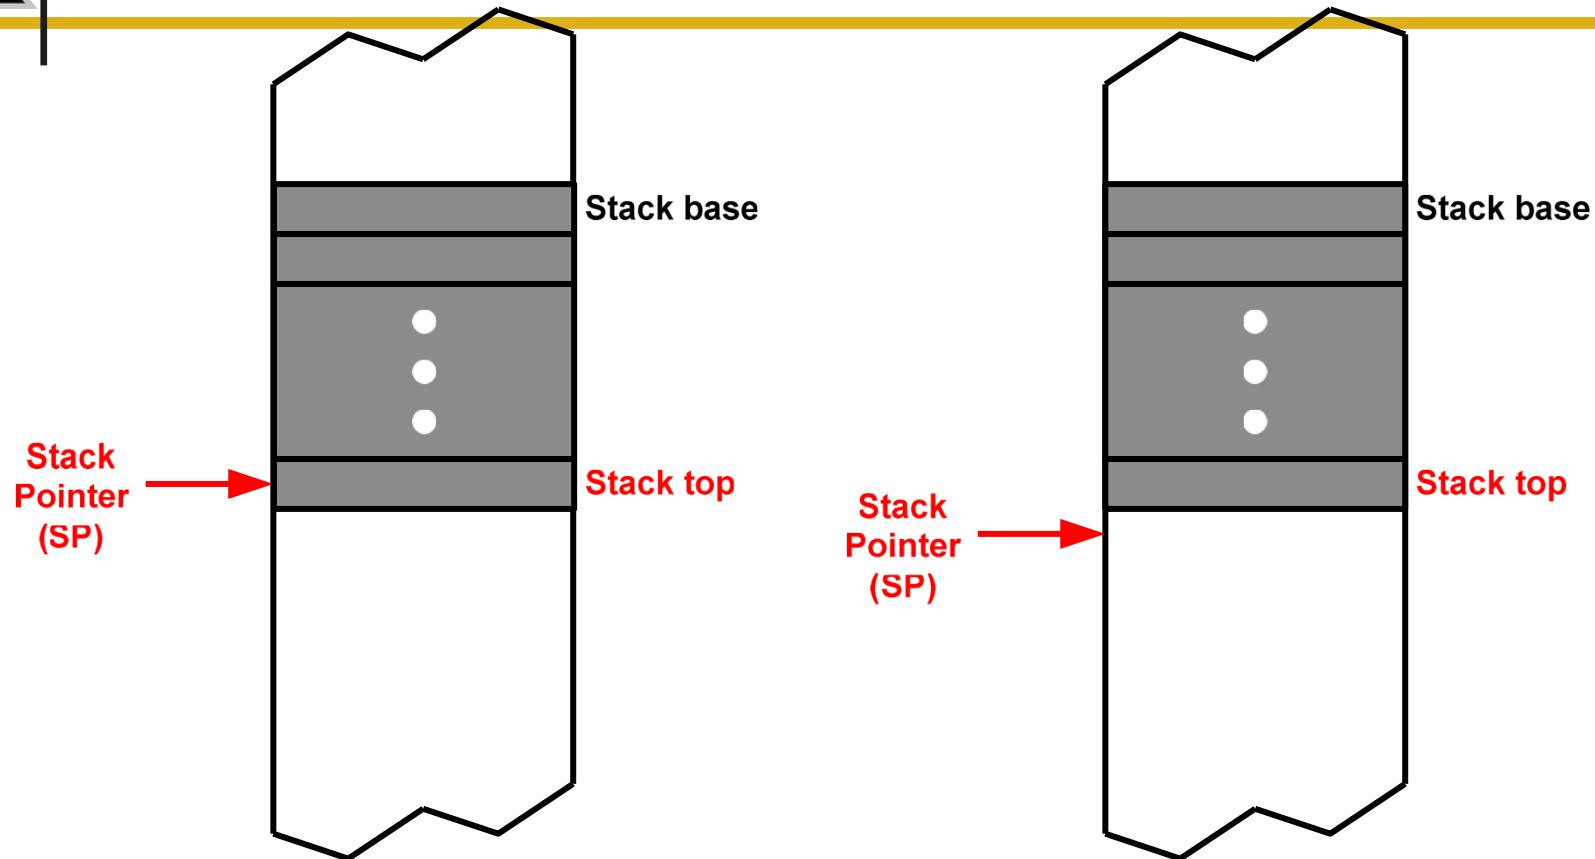
Descending stack: Stack grows towards low memory address



Ascending stack: Stack grows towards high memory address



Stack Growth Convention: Full vs Empty



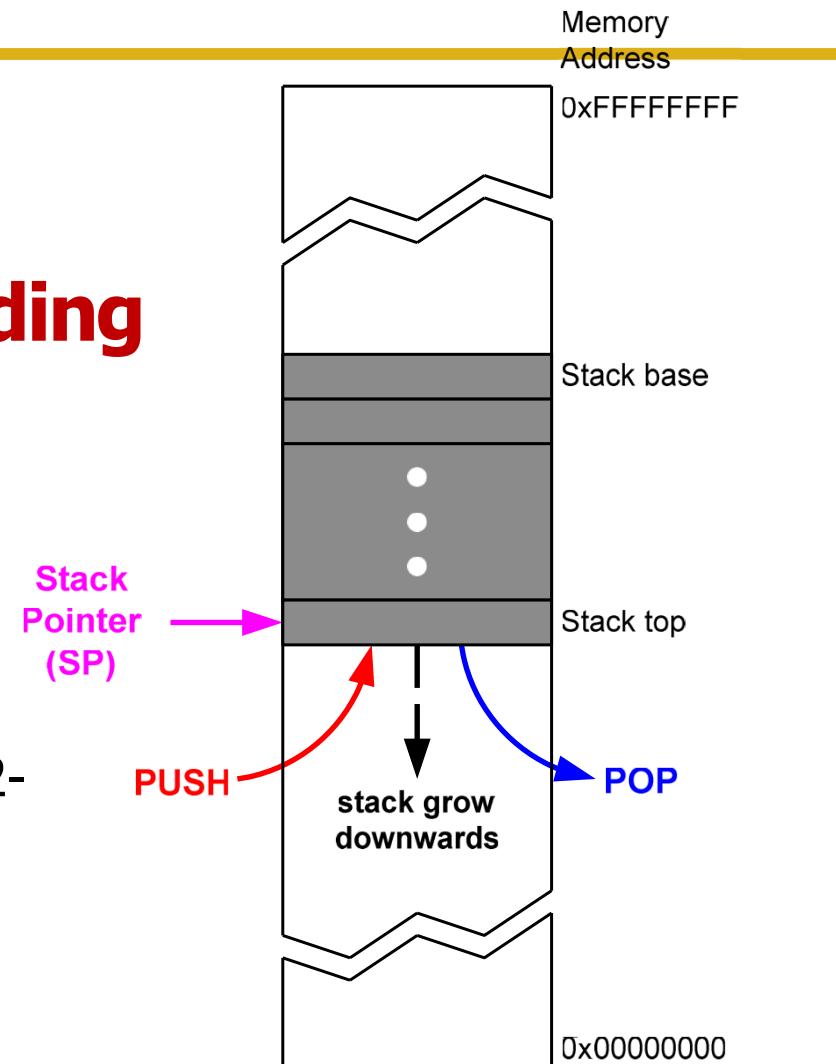
Full stack: SP points to the last item pushed onto the stack

Empty stack: SP points to the next free space on the stack



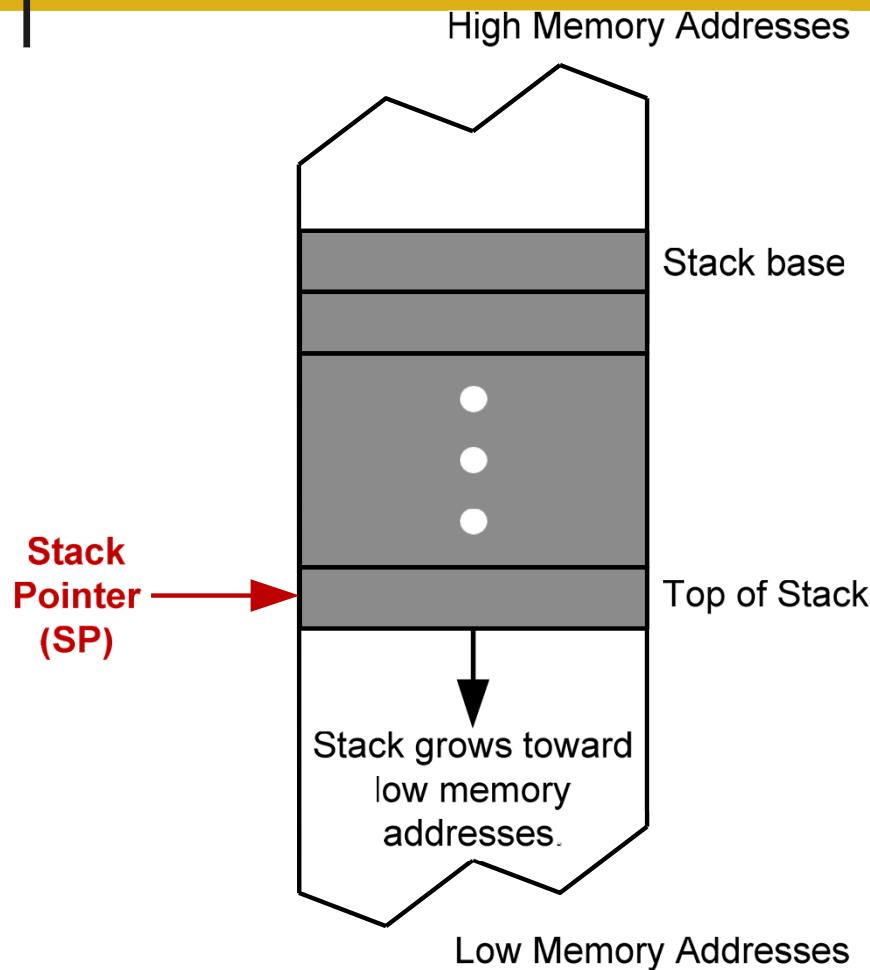
Cortex-M Stack

- Stack pointer (SP) = R13
- Cortex-M uses **full descending stack**
- stack pointer
 - decremented on **PUSH**
 - incremented on **POP**
 - SP starts at **0x20000200** for STM32-Discovery





Full Descending Stack



PUSH {register_list}
equivalent to:
STMDB SP!, {register_list}

DB: Decrement Before

POP {register_list}
equivalent to:
LDMIA SP!, {register_list}

IA: Increment After



Stack Implementation via STM and LDM

Stock Name	Push		Pop	
	Equivalent	Alternative	Equivalent	Alternative
Full Descending(FD)	STMFD SP!,list	STMDB SP!,list	LDMFD SP!,list	LDMIA SP!,list
Empty Descending(ED)	STMED SP!,list	STMDA SP!,list	LDMED SP!,list	LDMIB SP!,list
Full Ascending(FA)	STMFA SP!,list	STMIB SP!,list	LDMFA SP!,list	LDMDA SP!,list
Empty Ascending(EA)	STMEA SP!,list	STMIA SP!,list	LDMEA SP!,list	LDMDB SP!,list

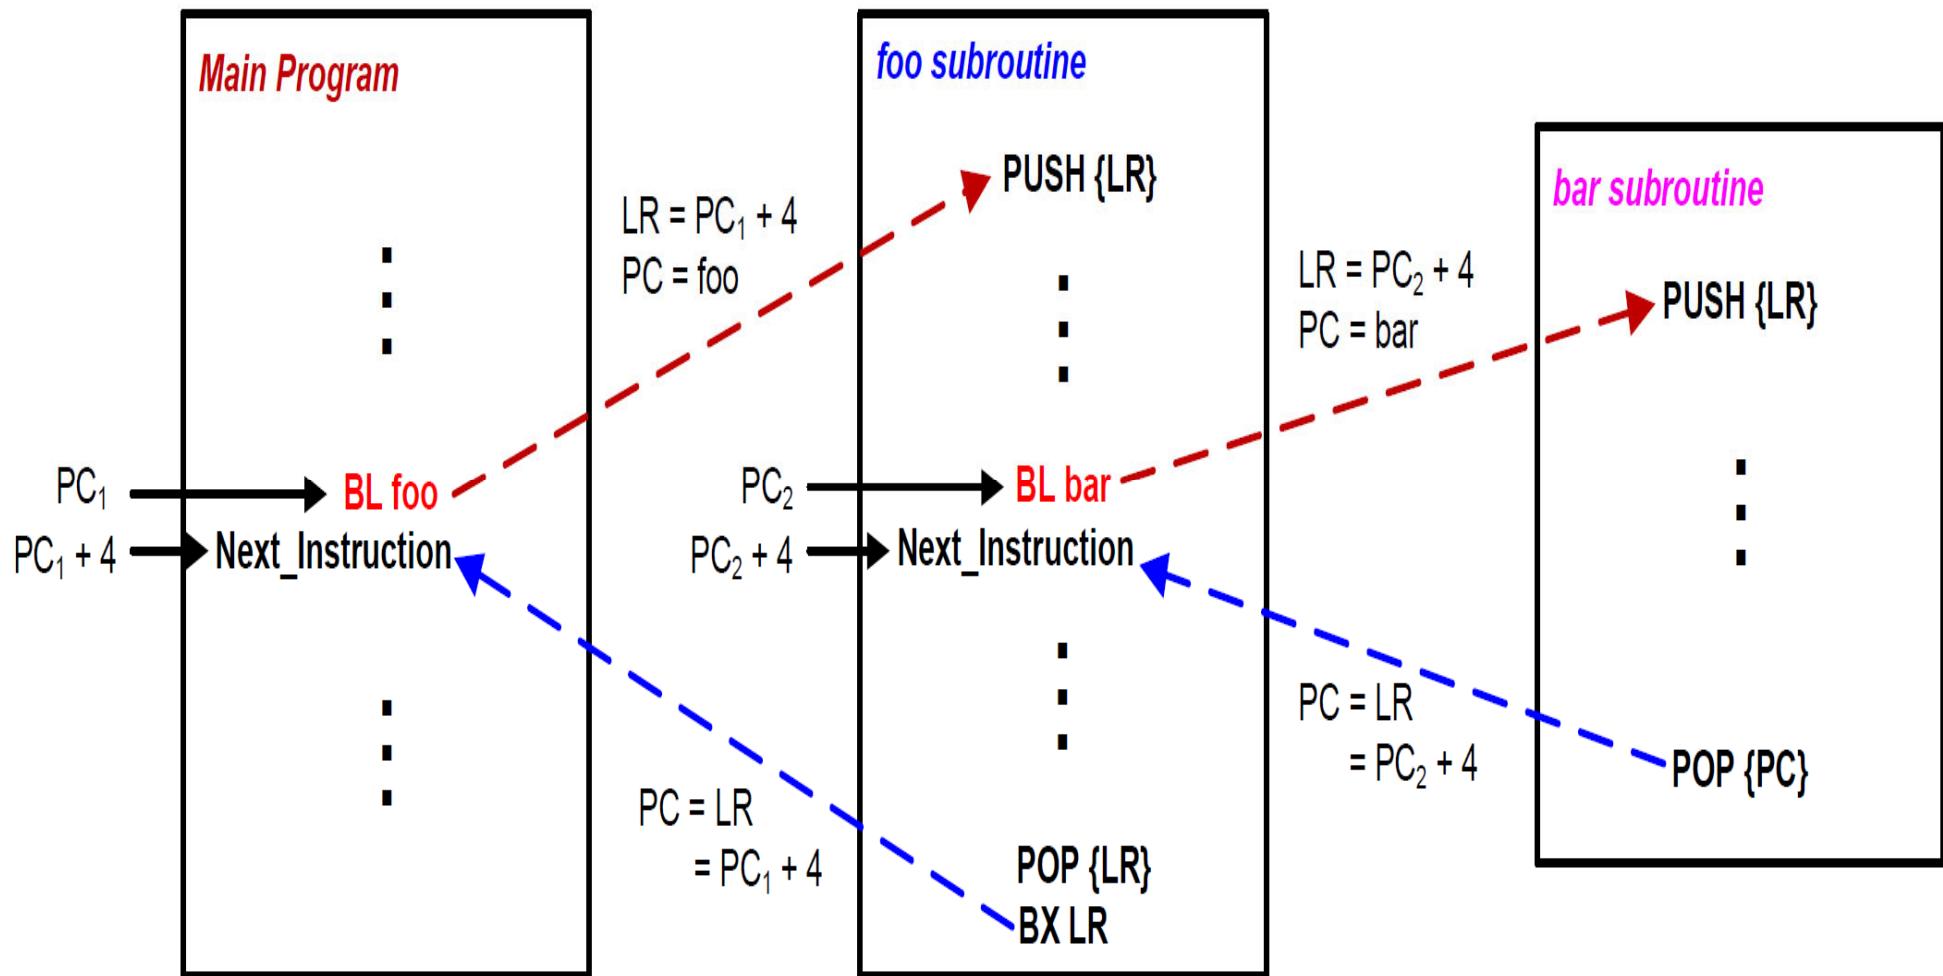


Stacks and Subroutines

- Stack can also be used to return from a subroutine.
- This is done by using PUSH {LR} at beginning of subroutine followed by POP {PC} at end of subroutine.
- This approach is required when a subroutine calls another subroutine.
- Under these circumstances, the LR will be overwritten in the caller subroutine.



Subroutine Calling Another Subroutine





Subroutine Calling Another Subroutine

MAIN PROC

```
    MOV R0 ,#2  
    BL QUAD  
    ENDL  
    ...  
    ENDP
```

QUAD

PROC

```
    PUSH {LR}  
    BL SQ  
    BL SQ  
    POP {LR}  
    BX LR  
    ENDP
```

SQ

PROC
 MUL R0 ,R0
 BX LR
 ENDP

Function MAIN

Function SQ

Function QUAD



2 Ways to Effectively Return Within Nested Subroutines

Method 1

PUSH {LR}

...

POP {LR}

BX LR

Method 2

PUSH {LR}

...

POP {PC}



Initializing the stack pointer (SP)

- Before using the stack, software has to define stack space and initialize the stack pointer (SP).
- The assembly file **startup.s** defines stack space and initialize SP.



Recursive Functions

- A **recursive function** is one that solves its task by **calling itself** on smaller pieces of data.
- An effective tactic is to
 - divide a problem into sub-problems of the same type as the original,
 - solve those sub-problems, and
 - combine the results



Recursion vs Iteration

- Any problem that can be solved **recursively** can also be solved **iteratively** (using loop).
- **Recursive functions vs Iterative functions**
 - Cons:
 - Recursive functions are slow
 - Recursive function take more memory
 - Pros
 - Recursive functions resembles the problem more naturally
 - Recursive function are easier to program, and debug.



Recursive Functions

- **push LR** (& working registers) onto stack before nested call
- **pop LR** (& working registers) off stack after nested return